

st125457_ulugbek_assignment

September 6, 2024

1 Exploratory data analysis

More resources for understanding data visualization

- <https://builtin.com/data-science/boxplot>
- <https://www.reddit.com/r/dataisbeautiful/>
- <https://mode.com/blog/violin-plot-examples/>
- <https://www.knowledgehut.com/blog/data-science/data-visualization-in-data-science>
- <https://www.analyticssteps.com/blogs/data-visualization-techniques>

```
[ ]: from IPython.display import Image
from IPython.core.display import HTML
Image(url= "https://reach2020.eu/wp-content/uploads/2019/05/
↳REACH-data-analytics-sophistication.png")
```

```
[ ]: <IPython.core.display.Image object>
```

1.1 Step1: fruits data

1.1) Load library and data

```
[ ]: # Include necessary library
import pandas as pd
import matplotlib.pyplot as plt

# make the plot appear without explicitly show()
%matplotlib inline
```

```
[ ]: fruits = pd.read_csv('http://msds.science.swu.ac.th/csv/fruit_data.csv ')
fruits.head()
```

```
[ ]: 
```

	Area	Perimeter	Length	Width	Bounding Box	Eccentricity	\
0	7293	322.350288	102.772555	90.655051	9100	1.133666	
1	6916	310.249783	99.093942	89.173834	8633	1.111245	
2	7057	316.007143	101.747826	88.561243	8900	1.148898	
3	6897	313.735065	99.056651	88.982012	8633	1.113221	
4	6883	309.421356	98.948844	88.909916	8633	1.112911	

```
Roundness Convexity Average Color Red Average Color Green \
```

0	0.641405	1.172637	178.950413	102.710744
---	----------	----------	------------	------------

1	0.641755	1.186141	177.669421	104.322314
2	0.633996	1.183518	178.628099	106.057851
3	0.639992	1.172964	178.851240	107.793388
4	0.638693	1.189317	179.628099	109.760331

	Average Color Blue	Target
0	65.719008	Apple
1	61.289256	Apple
2	61.933884	Apple
3	62.743802	Apple
4	64.462810	Apple

1.2) Explore data

```
[ ]: # data dimension
fruits.shape
```

```
[ ]: (80, 12)
```

```
[ ]: # show the name of data columns (they are candidates for the model's features)
fruits.columns
```

```
[ ]: Index(['Area', 'Perimeter', 'Length', 'Width', 'Bounding Box', 'Eccentricity',
          'Roundness', 'Convexity', 'Average Color Red', 'Average Color Green',
          'Average Color Blue', 'Target'],
          dtype='object')
```

```
[ ]: # check if any feature has null or empty value
fruits.isnull().sum()
```

```
[ ]: Area                0
     Perimeter           0
     Length              0
     Width               0
     Bounding Box        0
     Eccentricity         0
     Roundness            0
     Convexity            0
     Average Color Red    0
     Average Color Green  0
     Average Color Blue   0
     Target               0
     dtype: int64
```

```
[ ]: # Data stat
fruits.describe()
```

```
[ ]:
```

	Area	Perimeter	Length	Width	Bounding Box \
count	80.000000	80.000000	80.000000	80.000000	80.000000
mean	6274.46250	312.445633	101.136797	82.003298	8547.337500
std	1433.14043	14.467298	4.347729	13.551783	897.119327
min	3445.00000	276.007143	93.877625	53.408523	6300.000000
25%	5915.50000	303.178716	99.001695	78.172658	8175.000000
50%	6652.50000	309.300036	100.846689	84.732940	8376.000000
75%	7367.75000	326.967298	102.413086	91.990114	9528.000000
max	7861.00000	345.462987	119.525558	97.574934	9800.000000

	Eccentricity	Roundness	Convexity	Average Color Red \
count	80.000000	80.000000	80.000000	80.000000
mean	1.272982	0.580589	1.173267	166.620455
std	0.253911	0.101963	0.018514	44.815759
min	1.036767	0.368820	1.105078	84.041322
25%	1.099497	0.574367	1.168206	151.415289
50%	1.174686	0.634364	1.180198	180.157025
75%	1.292117	0.644497	1.184057	190.506198
max	1.863729	0.653344	1.196027	252.231405

	Average Color Green	Average Color Blue
count	80.000000	80.000000
mean	130.669008	73.239669
std	38.577439	59.671417
min	100.504132	17.834711
25%	102.710744	27.995868
50%	122.628099	65.702479
75%	133.733471	89.654959
max	252.008264	250.264463

```
[ ]: #fruits.groupby(['Target']).size()
```

```
[ ]: # Count number of samples for each target
fruits.Target.value_counts()
```

```
[ ]: Target
Apple      20
Banana     20
Mango      20
Orange     20
Name: count, dtype: int64
```

1.2: First visualization of data

```
[ ]: # load seaborn library
import seaborn as sns
```

```
[ ]: # create columns name as list and remove target column from the list
allcols = list(fruits.columns)
features = allcols[0:11]
print(features)

['Area', 'Perimeter', 'Length', 'Width', 'Bounding Box', 'Eccentricity',
'Roundness', 'Convexity', 'Average Color Red', 'Average Color Green', 'Average
Color Blue']

[ ]: fruits.columns

[ ]: Index(['Area', 'Perimeter', 'Length', 'Width', 'Bounding Box', 'Eccentricity',
'Roundness', 'Convexity', 'Average Color Red', 'Average Color Green',
'Average Color Blue', 'Target'],
dtype='object')


[ ]: ## Heatmap

# Select only the numerical columns for the heatmap
numerical_data = fruits.select_dtypes(include=['float64', 'int64'])

# Compute the correlation matrix for numerical columns
corr_matrix_numerical = numerical_data.corr()

# Plotting the heatmap for numerical columns
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix_numerical, annot=True, cmap='coolwarm', fmt=".2f",
↪linewidths=0.5)
plt.title('Correlation Matrix of Numerical Variables')
plt.show()

# Features: "Area", "Length", "Eccentricity", "Convexity", "Average Color Red",
↪"Average Color Green", "Average Color Blue"
```



box plots are useful for

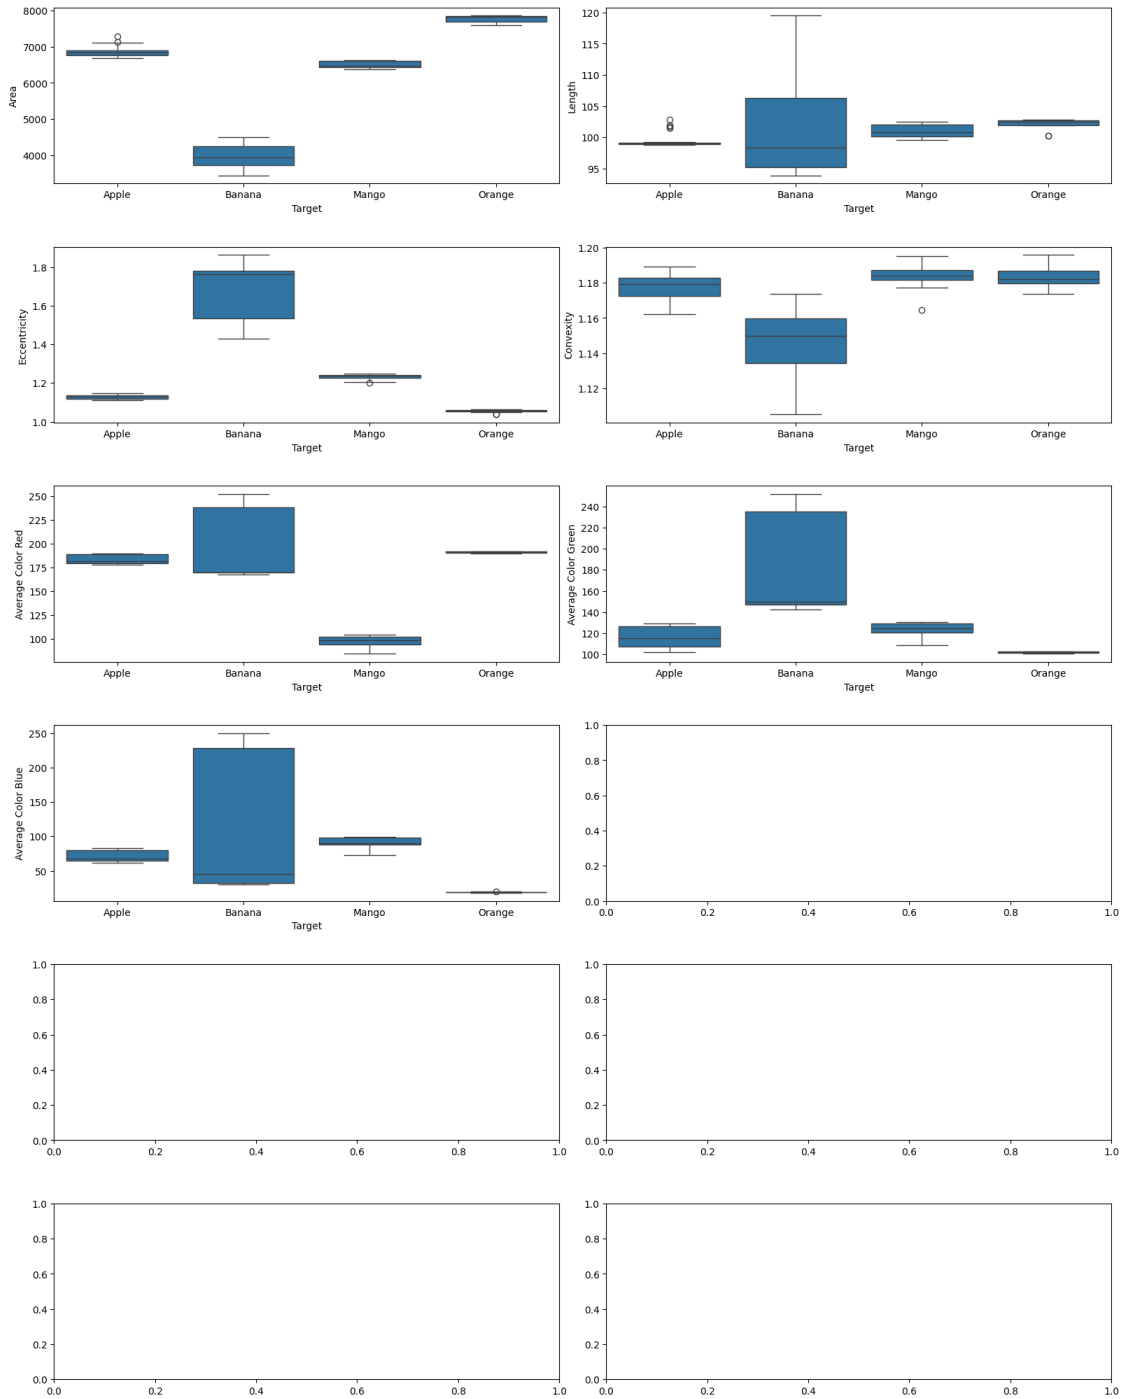
- tell the values of outliers
- Identify if data is symmetrical
- Check Skewness

Nature of data: Qualitative: (Nominal & Ordinal) and Quantitative: (Discrete and Continuous)

```
[ ]: # plot box plot of all candidate features for each target
candidate_features = ["Area", "Length", "Eccentricity", "Convexity", "Average_
↪Color Red", "Average Color Green", "Average Color Blue"]

fig, axs = plt.subplots(ncols=2, nrows=6, figsize=(16, 20))
index = 0
axs = axs.flatten()
```

```
for i, k in enumerate(candidate_features):
    sns.boxplot(x='Target', y=k, data=fruits, ax=axes[i])
plt.tight_layout(pad=0.5, w_pad=0.5, h_pad=3)
```



```
[ ]: # plot distribution of all candidate features
fig, axs = plt.subplots(ncols=3, nrows=4, figsize=(16, 10))
index = 0
axs = axs.flatten()

# sns.distplot(fruits[k].values, bins=20, ax=axs[index])
# for k, v in fruits.items():
for i, k in enumerate(candidate_features):
    sns.distplot(fruits[k].values, bins=20, ax=axs[i])

plt.tight_layout(pad=0.5, w_pad=0.5, h_pad=3)
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2473474893.py:9: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(fruits[k].values, bins=20, ax=axs[i])
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2473474893.py:9: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(fruits[k].values, bins=20, ax=axs[i])
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2473474893.py:9: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(fruits[k].values, bins=20, ax=axs[i])
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2473474893.py:9: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(fruits[k].values, bins=20, ax=axes[i])
C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2473474893.py:9: UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(fruits[k].values, bins=20, ax=axes[i])
C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2473474893.py:9: UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

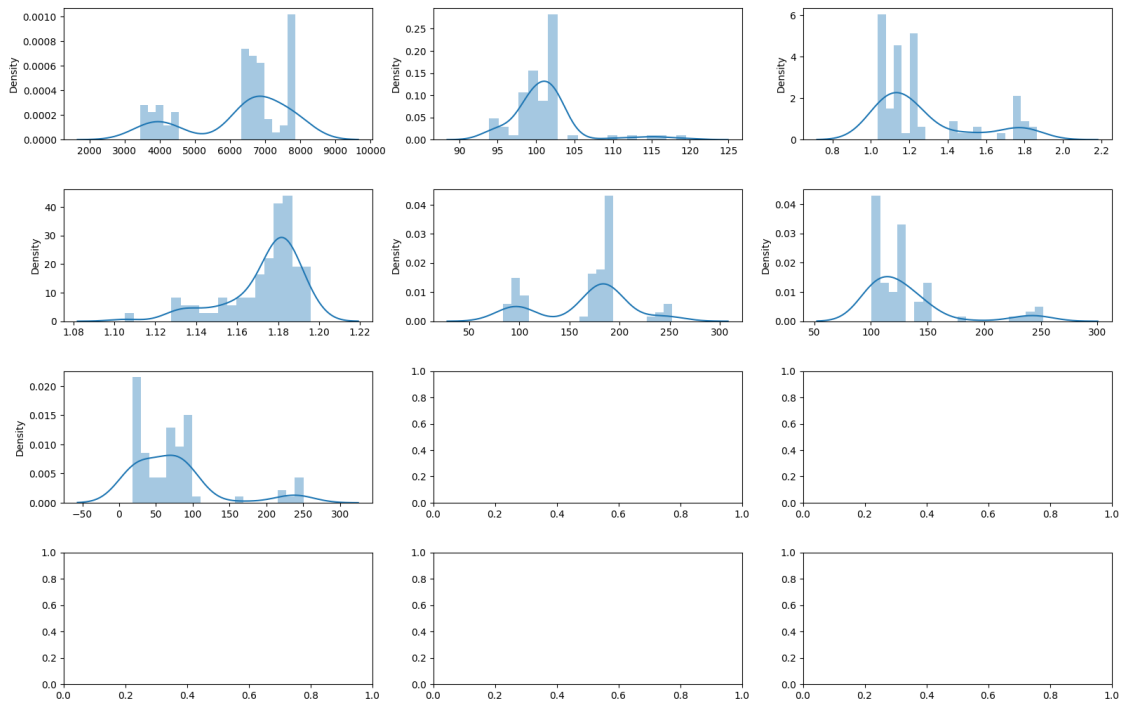
```
sns.distplot(fruits[k].values, bins=20, ax=axes[i])
C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2473474893.py:9: UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(fruits[k].values, bins=20, ax=axes[i])
```

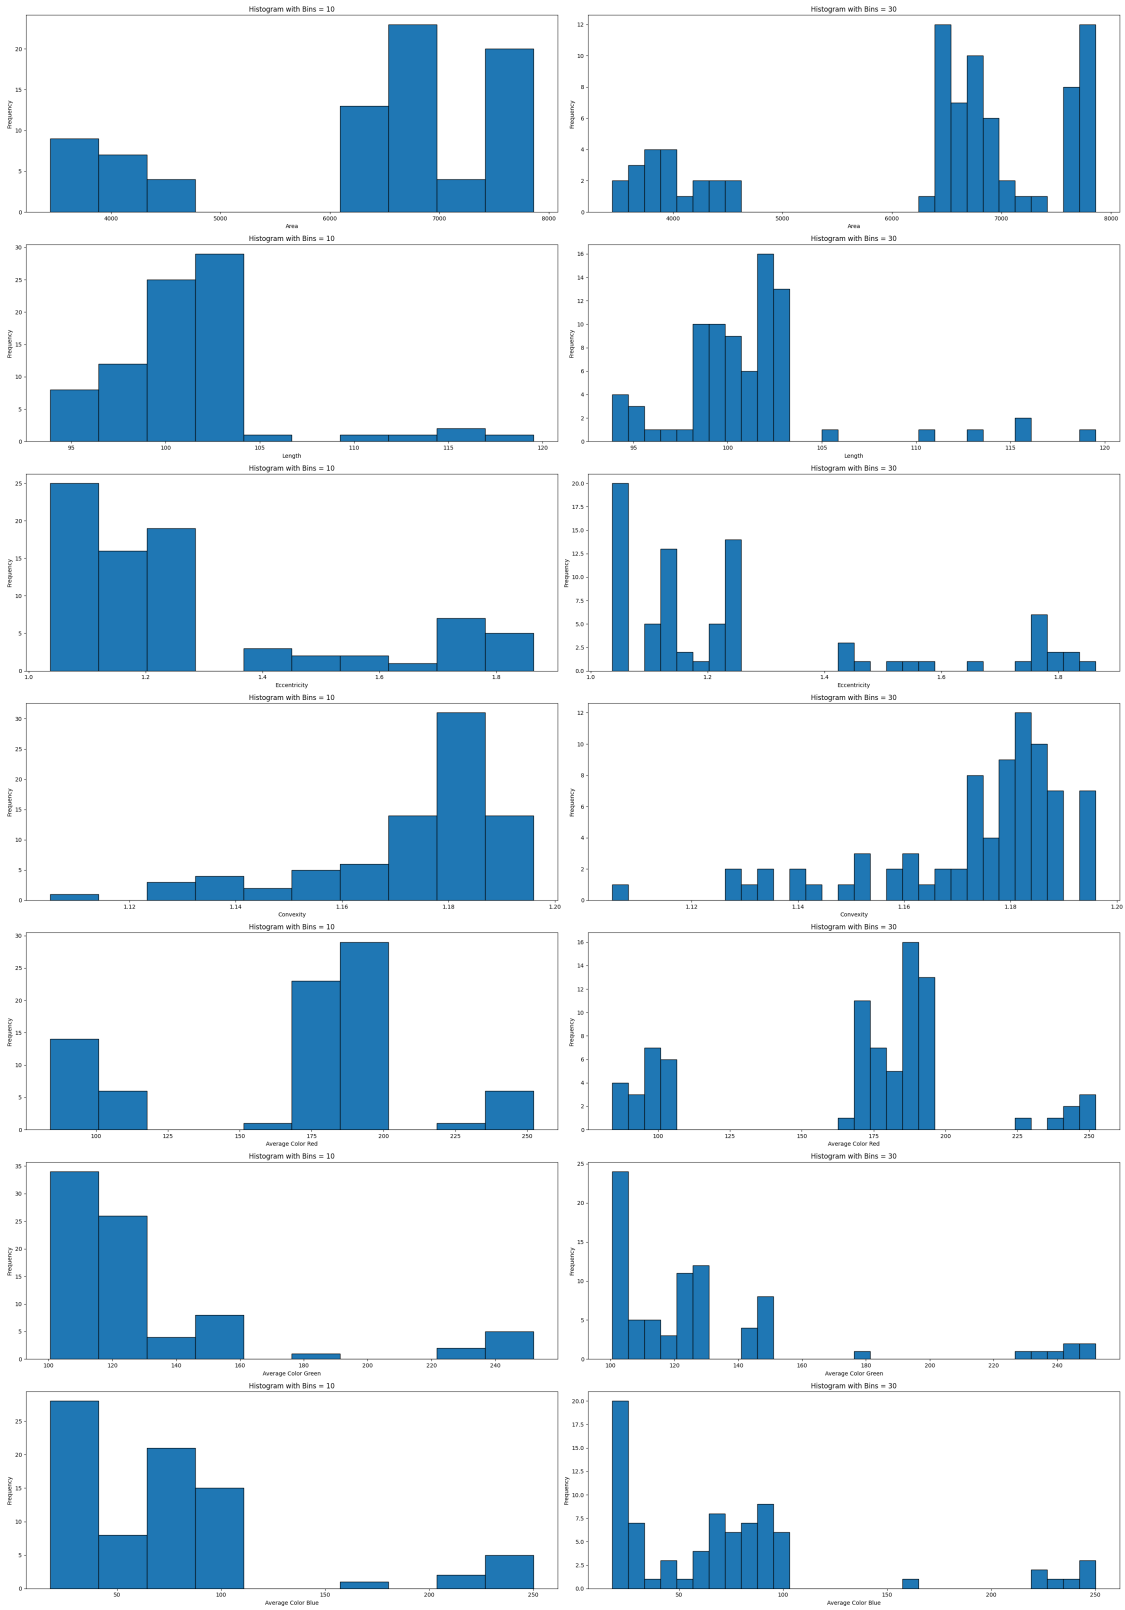
```
[ ]: # exercise1 : create histogram with bins = 10 and 30
numerical_data = fruits.select_dtypes(include=['float64', 'int64'])

fig, axs = plt.subplots(nrows=7, ncols=2, figsize=(28, 40))

for i, k in enumerate(candidate_features):
    axs[i, 0].hist(fruits[candidate_features[i]], bins=10, edgecolor='black')
    axs[i, 0].set_title('Histogram with Bins = 10')
    axs[i, 0].set_xlabel(candidate_features[i])
    axs[i, 0].set_ylabel('Frequency')

    axs[i, 1].hist(fruits[candidate_features[i]], bins=30, edgecolor='black')
    axs[i, 1].set_title('Histogram with Bins = 30')
    axs[i, 1].set_xlabel(candidate_features[i])
    axs[i, 1].set_ylabel('Frequency')

plt.tight_layout()
plt.show()
```



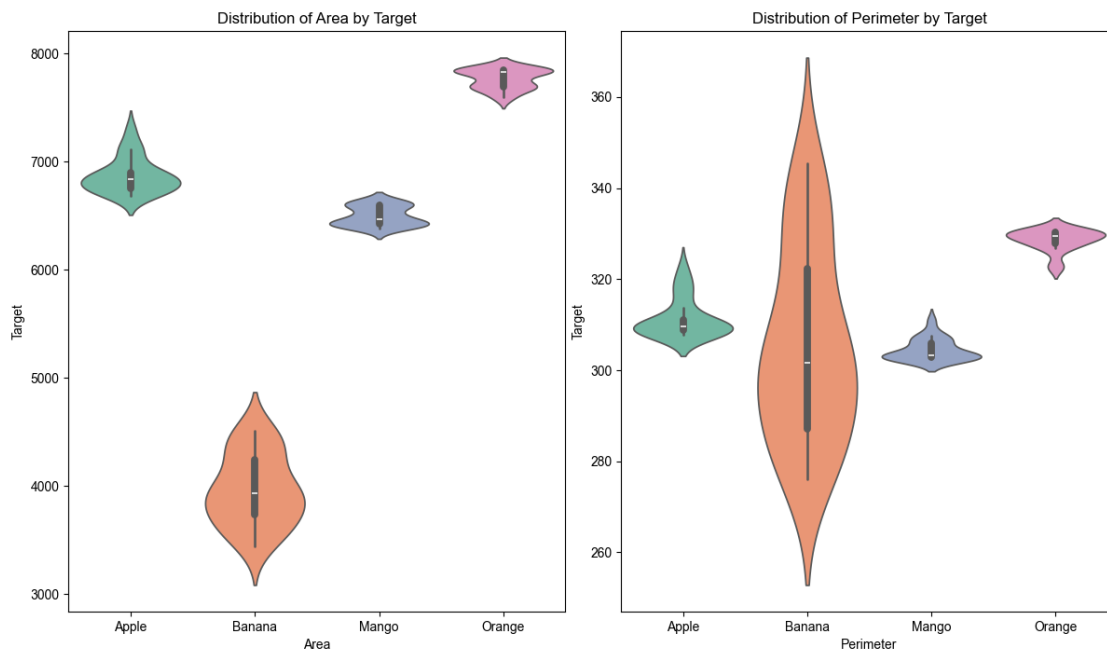
How to select number of bins: http://en.wikipedia.org/wiki/Histogram#Number_of_bins_and_w

violin plots combines box plots and kde

```
[ ]: # plot distribution of all candidate features for each target

fig, axs = plt.subplots(ncols=2, nrows=1, figsize=(12, 7))
index = 0
axs = axs.flatten()
sns.set_style('whitegrid')

for i, feature in enumerate(features[:2]):
    sns.violinplot(x='Target', hue='Target', data=fruits, ax=axs[i], y=feature,
                  palette='Set2')
    axs[i].set_title(f'Distribution of {feature} by Target')
    axs[i].set_xlabel(feature)
    axs[i].set_ylabel('Target')
plt.tight_layout(pad=0.5, w_pad=0.5, h_pad=3)
```




```
[ ]: sns.__version__
```

```
[ ]: '0.13.2'
```

```
[ ]: # see the distribution of individual feature by target value using kdeplot
sns.displot(data=fruits, x="Area", hue="Target", multiple="stack", kind="kde")
```

```
[ ]: <seaborn.axisgrid.FacetGrid at 0x20aff71a400>
```



```
[ ]: # deprecated after the seaborn version 0.14
sns.distplot(fruits[fruits['Target']=='Apple']['Area'], bins=20,color='r')
sns.distplot(fruits[fruits['Target']=='Banana']['Area'], bins=20,color='g')
sns.distplot(fruits[fruits['Target']=='Mango']['Area'], bins=20,color='b')
sns.distplot(fruits[fruits['Target']=='Orange']['Area'], bins=20,color='m')
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2696945370.py:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(fruits[fruits['Target']=='Apple']['Area'], bins=20,color='r')
C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2696945370.py:3: UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(fruits[fruits['Target']=='Banana']['Area'], bins=20,color='g')
C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2696945370.py:4: UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(fruits[fruits['Target']=='Mango']['Area'], bins=20,color='b')
C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2696945370.py:5: UserWarning:
```

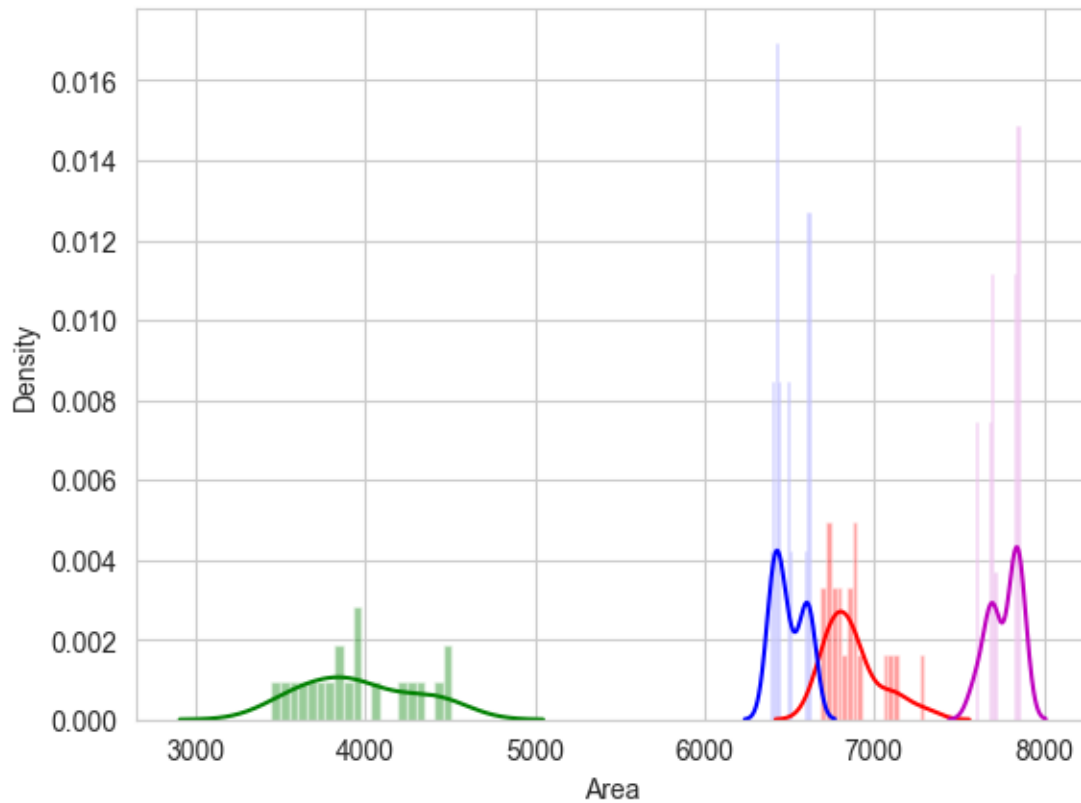
``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(fruits[fruits['Target']=='Orange']['Area'], bins=20,color='m')
```

```
[ ]: <Axes: xlabel='Area', ylabel='Density'>
```



```
[ ]: ## plot distribution of all candidate features for each target

## fig, axs = plt.subplots(ncols=2, nrows=8, figsize=(12, 7))
# index = 0
# axs = axs.flatten()
sns.set_style('whitegrid')
sns.set(style="ticks", color_codes=True)

sns.pairplot(data=fruits[candidate_features + ['Target']], hue='Target')
plt.tight_layout(pad=0.5, w_pad=0.5, h_pad=3)
```



```
[ ]: fruits.columns
```


```
[ ]: Index(['Area', 'Perimeter', 'Length', 'Width', 'Bounding Box', 'Eccentricity',
          'Roundness', 'Convexity', 'Average Color Red', 'Average Color Green',
          'Average Color Blue', 'Target'],
          dtype='object')
```

```
[ ]: correlation = fruits[candidate_features].corr()

# plot the cross relationship between each pair using heatmap
plt.figure(figsize=(16, 10))


# Plot the heatmap
sns.heatmap(correlation, annot=True, cmap='coolwarm', vmin=-1, vmax=1, fmt='.
↪2f', linewidths=0.5)
```

```
plt.title('Heatmap of Feature Correlations')
plt.show()
```




```
[ ]: import numpy as np
# exercise2: find correlation with positive and negative (hint:np.sign)
plt.figure(figsize=(16, 10))

# Plot the heatmap
sns.heatmap(np.sign(correlation), annot=True, cmap='coolwarm', vmin=-1, vmax=1,
            fmt='.2f', linewidths=0.5)
plt.title('Heatmap of Feature Correlations')
plt.show()
```





```
[ ]: # See the relationship between the area and width since it has 0.99
plt.figure(figsize=(10,6))


sns.scatterplot(data=fruits, x='Area', y='Width')
plt.show()
```




```
[ ]: # exercise3: select any 2 features that may be linearly dependent and plot ↵  
      ↪ regplot  
plt.figure(figsize=(10,6))  
sns.regplot(data=fruits, x='Area', y='Width')  
plt.show()
```




```
[ ]: # see the relationship between Perimeter and Length for each target  
# color = r,g,b,m  
  
fig, ax = plt.subplots()
```




```
[ ]: # see the relationship between Width and Length for each target  
fig, ax = plt.subplots()
```



```
[ ]: # Exercise: see the relationship between Roundness and Eccentricity for each ↵  
      ↪target  
fig, ax = plt.subplots()
```



```
[ ]: # Exercise: see the relationship between Area and "Average Color Red" for each ↵  
      ↪target  
fig, ax = plt.subplots()
```



```
[ ]: # exercise4: choose other pair of feature and scatter plot them.
```

1.2: Scale data

```
[ ]: # Load library for scaling  
from sklearn import preprocessing  
scaler = preprocessing.MinMaxScaler()  
#scaler = preprocessing.StandardScaler()
```

```
[ ]: # function to scale each column  
def scaleColumns(df, cols_to_scale):  
    pass
```

```
[ ]: # scale data in each column  
scaled_fruits = scaleColumns(fruits, features)
```

```
[ ]: # see the data after scaled  
scaled_fruits.head()
```

```
[ ]: # distribution after scaled  
fig, axs = plt.subplots(ncols=3, nrows=4, figsize=(16, 10))  
index = 0
```

```

axs = axs.flatten()

# sns.distplot(fruits[k].values, bins=20, ax=axs[index])
# for k,v in scaled_fruits.items():
for k in features:
    scaled_fruits[k].hist(bins=20, ax=axs[index])
    axs[index].set_title(k)
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.5, h_pad=3)

```

Dataset taken from <https://www.kaggle.com/datasets/harishkumardatalab/data-science-salary-2021-2023>

```
[ ]: salary_data = pd.read_csv('data_science_salary_2021_2023.csv')
```

```
[ ]: salary_data.shape
```

```
[ ]: (3761, 9)
```

```
[ ]: salary_data.head()
```

```
[ ]:
work_year  experience_level  employment_type  job_title \
0         2023              EN             FT    Applied Scientist
1         2023              EN             FT    Applied Scientist
2         2023              EN             FT  Data Quality Analyst
3         2023              EN             FT  Compliance Data Analyst
4         2023              EN             FT    Applied Scientist

salary  salary_currency  salary_in_usd  company_location  company_size
0   213660             USD         213660             US             L
1   130760             USD         130760             US             L
2   100000             USD         100000             NG             L
3    30000             USD          30000             NG             L
4   204620             USD         204620             US             L

```

```
[ ]: salary_data.columns
```

```
[ ]: Index(['work_year', 'experience_level', 'employment_type', 'job_title',
          'salary', 'salary_currency', 'salary_in_usd', 'company_location',
          'company_size'],
          dtype='object')
```

```
[ ]: salary_data['experience_level'].unique()
```

```
[ ]: array(['EN', 'EX', 'MI', 'SE'], dtype=object)
```

```
[ ]: salary_data['company_location'].unique()
```



```
[ ]: array(['US', 'NG', 'IN', 'CA', 'ES', 'GH', 'DE', 'CH', 'AU', 'SE', 'BR',
          'GB', 'VN', 'BA', 'GR', 'HK', 'NL', 'FI', 'IE', 'SG', 'SI', 'MX',
          'FR', 'HR', 'AM', 'KE', 'RO', 'TH', 'CF', 'UA', 'IL', 'CO', 'PT',
          'EE', 'LV', 'MK', 'PK', 'IT', 'MA', 'AR', 'CR', 'IR', 'HU', 'AS',
          'BE', 'AT', 'ID', 'LU', 'MY', 'CZ', 'DZ', 'RU', 'PL', 'LT', 'TR',
          'BO', 'EG', 'AL', 'SK', 'PR', 'AE', 'DK', 'IQ', 'CN', 'BS', 'JP',
          'CL', 'MD', 'MT', 'PH', 'HN', 'NZ'], dtype=object)
```

```
[ ]: salary_data['job_title'].value_counts()
```

```
[ ]: job_title
Data Engineer          1040
Data Scientist          840
Data Analyst            614
Machine Learning Engineer 291
Analytics Engineer      103

...
Compliance Data Analyst    1
BI Data Engineer           1
Deep Learning Researcher   1
Head of Machine Learning   1
Staff Data Analyst         1
Name: count, Length: 93, dtype: int64
```

```
[ ]: salary_data['work_year'].unique()
```

```
[ ]: array([2023, 2022, 2021, 2020], dtype=int64)
```

```
[ ]: salary_data['job_title'].unique()
```

```
[ ]: array(['Applied Scientist', 'Data Quality Analyst',
          'Compliance Data Analyst', 'Machine Learning Engineer',
          'Research Scientist', 'Data Engineer', 'Data Analyst',
          'Data Scientist', 'BI Data Engineer', 'Research Engineer',
          'Business Data Analyst', 'Autonomous Vehicle Technician',
          'Applied Machine Learning Scientist', 'AI Programmer',
          'AI Developer', 'Computer Vision Engineer', 'BI Developer',
          'Big Data Engineer', 'Deep Learning Engineer', 'Head of Data',
          'Analytics Engineer', 'Data Architect', 'Director of Data Science',
          'Head of Data Science', 'Data Analytics Manager', 'ML Engineer',
          'Applied Machine Learning Engineer', 'Applied Data Scientist',
          'ETL Engineer', 'Data Specialist', 'Lead Data Analyst',
          'Data Manager', 'Machine Learning Scientist', 'MLOps Engineer',
          'Financial Data Analyst', 'Software Data Engineer',
          'Data Science Manager', 'Data Science Consultant',
          'Machine Learning Infrastructure Engineer', 'Insight Analyst',
          'AI Scientist', 'Data Infrastructure Engineer',
          'Data Science Lead', 'Product Data Analyst', 'Data Analytics Lead',
```

```
'Machine Learning Research Engineer', 'Principal Data Scientist',
'Data Modeler', 'Business Intelligence Engineer',
'Data Strategist', 'Data DevOps Engineer',
'Machine Learning Researcher', 'Cloud Database Engineer',
'Data Operations Engineer', 'Deep Learning Researcher',
'BI Analyst', 'Data Analytics Specialist', 'BI Data Analyst',
'Machine Learning Software Engineer', 'Big Data Architect',
'Computer Vision Software Engineer', 'Azure Data Engineer',
'Data Lead', 'Data Science Engineer', 'NLP Engineer',
'Machine Learning Developer', 'Data Analytics Engineer',
'Data Analytics Consultant', '3D Computer Vision Researcher',
'Lead Data Engineer', 'Head of Machine Learning',
'Data Scientist Lead', 'Lead Data Scientist', 'ETL Developer',
'Principal Data Analyst', 'Manager Data Management',
'Principal Machine Learning Engineer',
'Data Management Specialist', 'Data Science Tech Lead',
'Cloud Data Engineer', 'Data Operations Analyst',
'Marketing Data Analyst', 'Product Data Scientist',
'Principal Data Architect', 'Machine Learning Manager',
'Lead Machine Learning Engineer', 'Marketing Data Engineer',
'Power BI Developer', 'Cloud Data Architect',
'Principal Data Engineer', 'Staff Data Scientist',
'Finance Data Analyst', 'Staff Data Analyst'], dtype=object)
```

```
[ ]: salary_data.salary
```

```
[ ]: 0      213660
      1      130760
      2     100000
      3      30000
      4     204620
      ...
     3756    130000
     3757     80000
     3758    190200
     3759     40000
     3760    412000
      Name: salary, Length: 3761, dtype: int64
```

Salary vs Experience | Use proper visualization technique to show the relationship between salary vs experience (2 marks)

```
[ ]: plt.figure(figsize=(12, 8))

      sns.boxplot(x='experience_level', y='salary_in_usd', data=salary_data,
                  palette='viridis')

      plt.title('Salary vs Experience')
```


```
plt.xlabel('Experience')
plt.ylabel('Salary in USD')
plt.show()
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\3867932903.py:3:

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='experience_level', y='salary_in_usd', data=salary_data,
palette='viridis')
```



```
[ ]: plt.figure(figsize=(12, 8))


sns.violinplot(x='experience_level', y='salary_in_usd', data=salary_data,
palette='viridis')

plt.title('Salary vs Experience')
plt.xlabel('Experience')
plt.ylabel('Salary in USD')
plt.show()
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\771222664.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.violinplot(x='experience_level', y='salary_in_usd', data=salary_data,
palette='viridis')
```



```
[ ]: plt.figure(figsize=(12, 8))


# Plots individual data points for each category, showing the spread and
↳clustering of values
sns.stripplot(x='experience_level', y='salary_in_usd', data=salary_data,↳
↳palette='viridis')

plt.title('Salary vs Experience')
plt.xlabel('Experience')
plt.ylabel('Salary in USD')
plt.show()
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2105702522.py:3:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.stripplot(x='experience_level', y='salary_in_usd', data=salary_data,
palette='viridis')
```



Salary vs employment_type | Use proper visualization technique to show the relationship between salary vs employment_type (2 marks)

```
[ ]: salary_data.employment_type.unique()
```

```
[ ]: array(['FT', 'PT', 'CT', 'FL'], dtype=object)
```

```
[ ]: plt.figure(figsize=(12, 8))

sns.boxplot(x='employment_type', y='salary_in_usd', data=salary_data,
            palette='viridis')

plt.title("Salary vs Employment Type")
plt.xlabel("Employment Type")
plt.ylabel("Salary in USD")
```

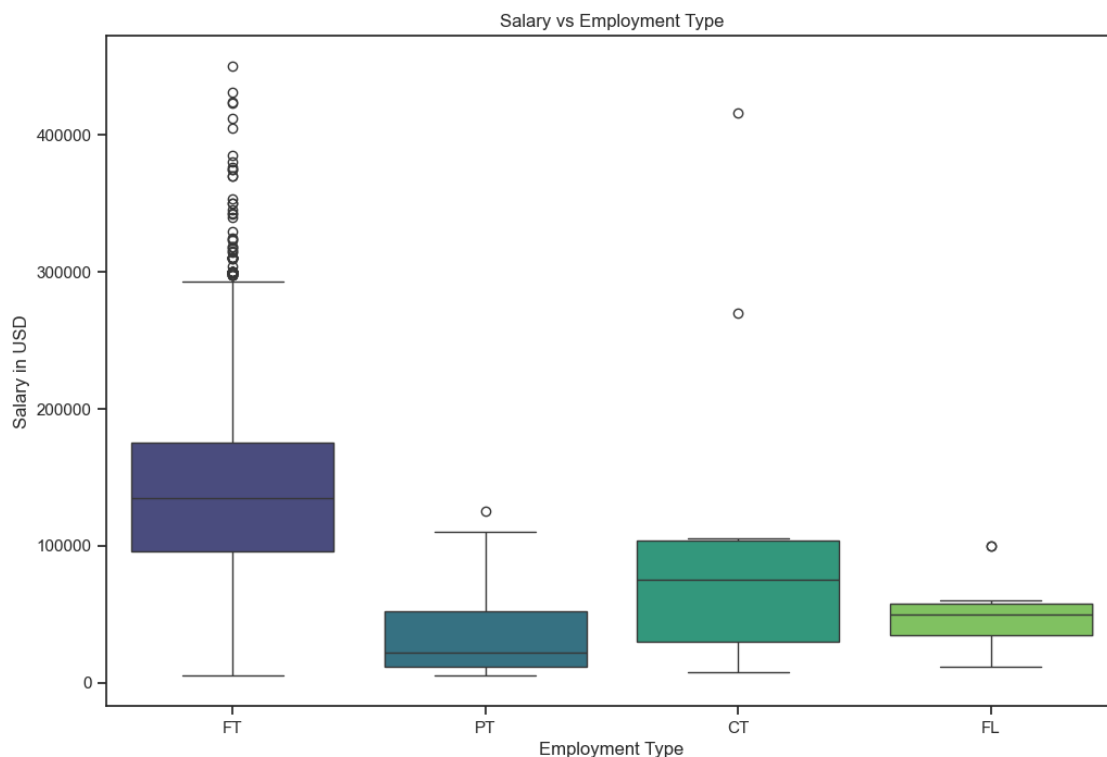
```
plt.show()
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\3496671623.py:3:

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='employment_type', y='salary_in_usd', data=salary_data,
palette='viridis')
```



```
[ ]: plt.figure(figsize=(12, 8))

sns.violinplot(x='employment_type', y='salary_in_usd', data=salary_data,
palette='viridis')

plt.title("Salary vs Employment Type")
plt.xlabel("Employment Type")
plt.ylabel("Salary in USD")

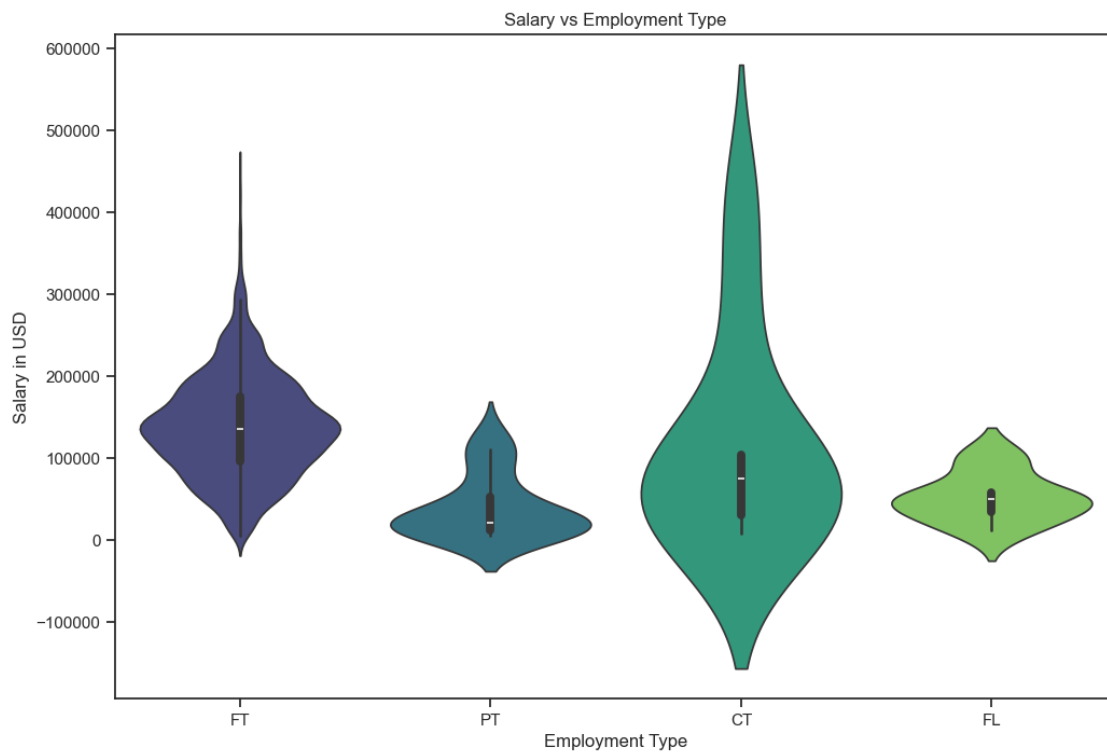
plt.show()
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\1881712386.py:3:

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.violinplot(x='employment_type', y='salary_in_usd', data=salary_data,
palette='viridis')
```



```
[ ]: plt.figure(figsize=(12, 8))

sns.stripplot(x='employment_type', y='salary_in_usd', data=salary_data,
palette='viridis')

plt.title("Salary vs Employment Type")
plt.xlabel("Employment Type")
plt.ylabel("Salary in USD")

plt.show()
```

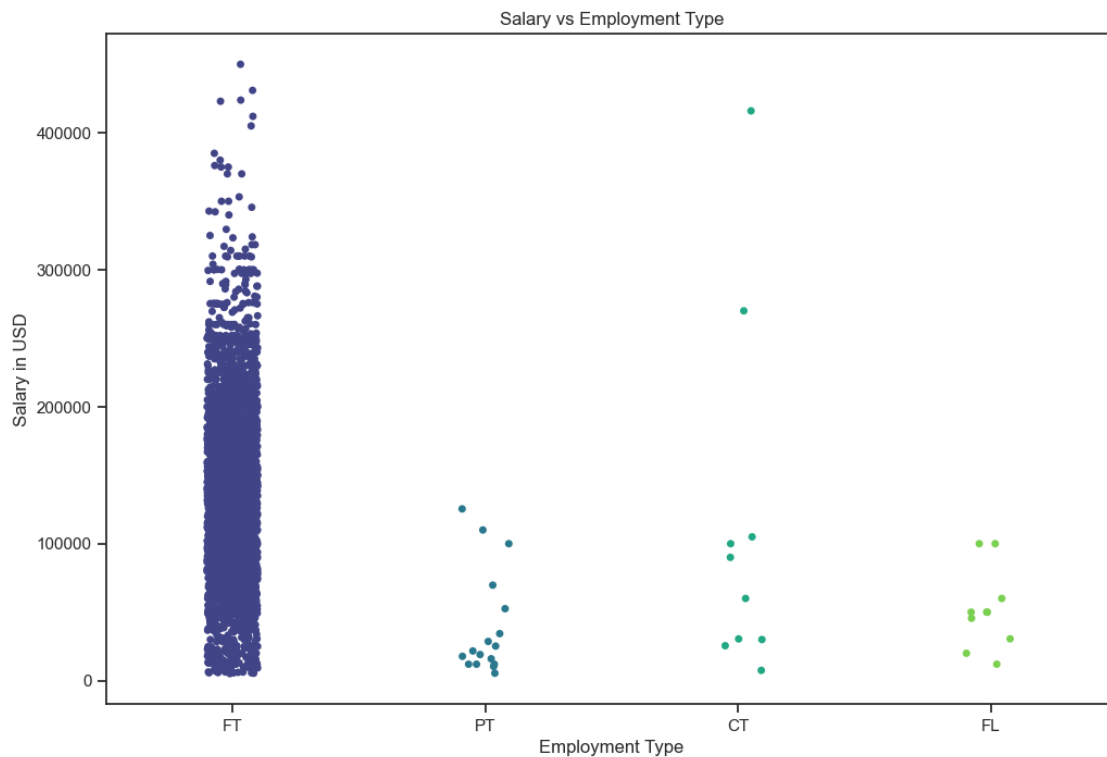
C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\3358226939.py:3:

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in

v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.stripplot(x='employment_type', y='salary_in_usd', data=salary_data,
palette='viridis')
```



Salary vs company_size | Use proper visualization technique to show the relationship between salary vs company_size (2 marks)

```
[ ]: salary_data.company_size.unique()
```

```
[ ]: array(['L', 'M', 'S'], dtype=object)
```

```
[ ]: plt.figure(figsize=(12, 8))
```

```
sns.boxplot(x='company_size', y='salary_in_usd', data=salary_data,
palette='viridis')
```

```
plt.title('Salary vs Company Size')
```

```
plt.xlabel('Company Size')
```


```
plt.ylabel('Salary')
```

```
plt.show()
```


C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\879199461.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='company_size', y='salary_in_usd', data=salary_data,
palette='viridis')
```



```
[ ]: plt.figure(figsize=(12, 8))

sns.violinplot(x='company_size', y='salary_in_usd', data=salary_data,
palette='viridis')

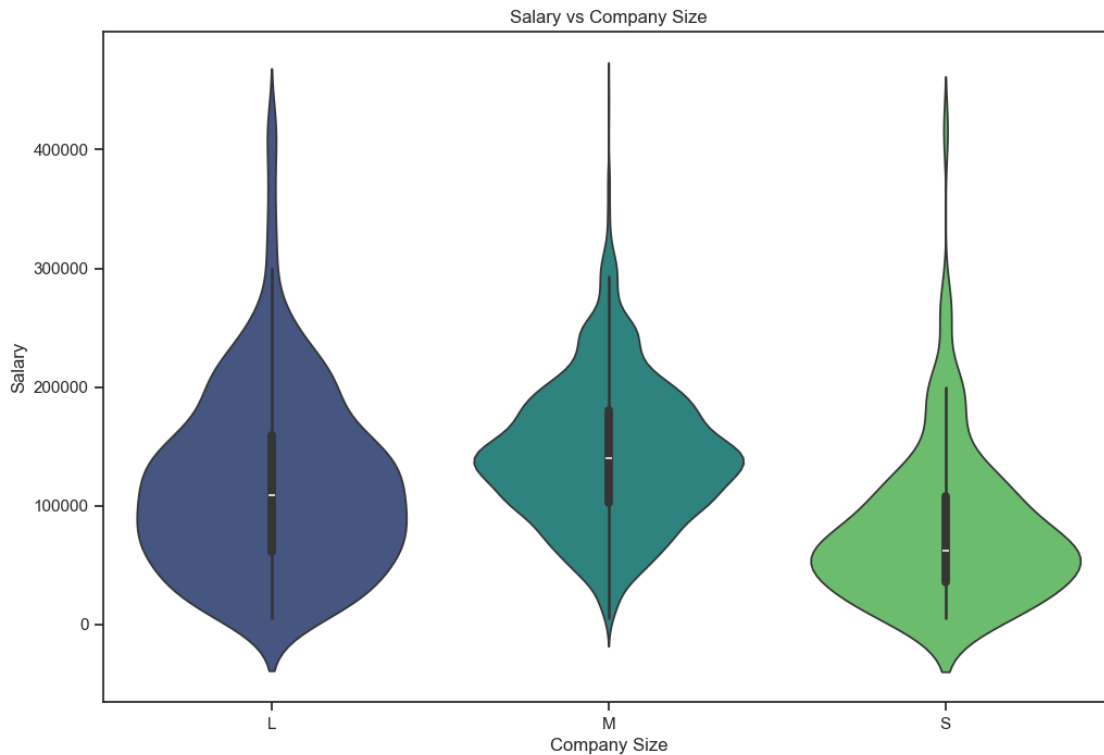
plt.title('Salary vs Company Size')
plt.xlabel('Company Size')
plt.ylabel('Salary')

plt.show()
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\3889054652.py:3:
FutureWarning:

Passing ``palette`` without assigning ``hue`` is deprecated and will be removed in v0.14.0. Assign the ``x`` variable to ``hue`` and set ``legend=False`` for the same effect.

```
sns.violinplot(x='company_size', y='salary_in_usd', data=salary_data,  
palette='viridis')
```




```
[ ]: plt.figure(figsize=(12, 8))  
  
sns.stripplot(x='company_size', y='salary_in_usd', data=salary_data,  
             palette='viridis')  
  
plt.title('Salary vs Company Size')  
plt.xlabel('Company Size')  
plt.ylabel('Salary')  
  
plt.show()
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2185805539.py:3:
FutureWarning:

Passing ``palette`` without assigning ``hue`` is deprecated and will be removed in v0.14.0. Assign the ``x`` variable to ``hue`` and set ``legend=False`` for the same

effect.

```
sns.stripplot(x='company_size', y='salary_in_usd', data=salary_data,
palette='viridis')
```



Salary vs Location | Get average salary for all the location and write out the top 3 and bottom 3 location according to salaries (5 marks)

```
[ ]: salary_per_loc = salary_data.groupby('company_location').salary_in_usd.mean().
      ↪reset_index().sort_values(by='salary_in_usd', ascending=False)
top_3_locations = salary_per_loc.head(3)
bottom_3_locations = salary_per_loc.tail(3)

top_3_locations, bottom_3_locations
```

```
[ ]: (  company_location  salary_in_usd
37          IL  271446.500000
59          PR  167500.000000
70          US  151801.053859,
      company_location  salary_in_usd
47          MA    10000.0
9           BO     7500.0
49          MK     6304.0)
```

```
[ ]: plt.figure(figsize=(20, 18))

sns.boxplot(x='company_location', y='salary_in_usd', data=salary_data,
           palette='viridis')

plt.title("Salary vs Location")
plt.xlabel("Location")
plt.ylabel("Salary")


plt.tight_layout()
```

C:\Users\eraco\AppData\Local\Temp\ipykernel_14268\2941615401.py:3:

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='company_location', y='salary_in_usd', data=salary_data,
palette='viridis')
```



Plot average salaries over the year(line chart). What is the trend of salaries/Is it increasing or decreasing? Explain. (5 marks)


```
[ ]: avg_salary_year = salary_data.groupby('work_year').salary_in_usd.mean().
    ↪reset_index()

plt.figure(figsize=(12, 8))

sns.lineplot(x='work_year', y='salary_in_usd', data=avg_salary_year,
    ↪marker='o', color='b')

plt.title("Average salary over Years")
plt.xlabel("Years")
plt.ylabel("Salary in USD")

plt.show()
```



It is clear that the salary is increasing over the years worker work in the same job. It is natural - the more experience you have, the more you competitive in the market, the more your salary.

Count values for each job title. Select the jobs that have less than 50 entries. Replace those values with Others. Create a pie chart. (10 marks)

```
[ ]: job_counts = salary_data['job_title'].value_counts()

salary_data['job_title'] = salary_data['job_title'].apply(lambda x: x if
    ↪ job_counts[x] >= 50 else 'Others')

updated_job_counts = salary_data['job_title'].value_counts()

plt.figure(figsize=(16, 14))
plt.pie(updated_job_counts, labels=updated_job_counts.index, autopct='%1.1f%%',
    ↪ startangle=90, colors=sns.color_palette('pastel'))

plt.title('Job Title')
plt.axis('equal')

plt.tight_layout()
```

