

# Copy\_of\_Lab\_8\_Data\_Pre\_processing\_Homework

November 4, 2024

```
[ ]: #Run this command before importing the dataset.
```

```
!pip install ucimlrepo
```

Requirement already satisfied: ucimlrepo in /usr/local/lib/python3.10/dist-packages (0.0.7)

Requirement already satisfied: pandas>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from ucimlrepo) (2.2.2)

Requirement already satisfied: certifi>=2020.12.5 in /usr/local/lib/python3.10/dist-packages (from ucimlrepo) (2024.8.30)

Requirement already satisfied: numpy>=1.22.4 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0->ucimlrepo) (1.26.4)

Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0->ucimlrepo) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0->ucimlrepo) (2024.2)

Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0->ucimlrepo) (2024.2)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas>=1.0.0->ucimlrepo) (1.16.0)

Load the dataset from the repository <https://archive.ics.uci.edu/dataset/2/adult> (4 marks)

```
[ ]: from ucimlrepo import fetch_ucirepo
```

```
# fetch dataset
```

```
adult = fetch_ucirepo(id=2)
```

```
# data (as pandas dataframes)
```

```
X = adult.data.features
```

```
y = adult.data.targets
```

```
[ ]: X.head()
```

```
[ ]:
   age  workclass  fnlwgt  education  education-num  \
0   39   State-gov   77516   Bachelors             13
1   50  Self-emp-not-inc   83311   Bachelors             13
2   38    Private  215646   HS-grad              9
3   53    Private  234721    11th              7
```

4    28                    Private   338409   Bachelors                    13

	marital-status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	capital-gain	capital-loss	hours-per-week	native-country
0	2174	0	40	United-States
1	0	0	13	United-States
2	0	0	40	United-States
3	0	0	40	United-States
4	0	0	40	Cuba

```
[ ]: y.head()
```

```
[ ]: income
0  <=50K
1  <=50K
2  <=50K
3  <=50K
4  <=50K
```

```
[ ]: X.shape, y.shape
```

```
[ ]: ((48842, 14), (48842, 1))
```

```
[ ]: X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   48842 non-null  int64
1   workclass             47879 non-null  object
2   fnlwgt               48842 non-null  int64
3   education             48842 non-null  object
4   education-num         48842 non-null  int64
5   marital-status       48842 non-null  object
6   occupation            47876 non-null  object
7   relationship          48842 non-null  object
8   race                 48842 non-null  object
9   sex                  48842 non-null  object
10  capital-gain          48842 non-null  int64
11  capital-loss          48842 non-null  int64
```

```

12 hours-per-week 48842 non-null int64
13 native-country 48568 non-null object
dtypes: int64(6), object(8)
memory usage: 5.2+ MB

```

```
[ ]: X.describe()
```

```
[ ]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss \
count	48842.000000	4.884200e+04	48842.000000	48842.000000	48842.000000
mean	38.643585	1.896641e+05	10.078089	1079.067626	87.502314
std	13.710510	1.056040e+05	2.570973	7452.019058	403.004552
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000
25%	28.000000	1.175505e+05	9.000000	0.000000	0.000000
50%	37.000000	1.781445e+05	10.000000	0.000000	0.000000
75%	48.000000	2.376420e+05	12.000000	0.000000	0.000000
max	90.000000	1.490400e+06	16.000000	99999.000000	4356.000000

  

	hours-per-week
count	48842.000000
mean	40.422382
std	12.391444
min	1.000000
25%	40.000000
50%	40.000000
75%	45.000000
max	99.000000

```
[ ]: X.isna().sum()
```

```
[ ]: age                0
workclass             963
fnlwgt                0
education              0
education-num          0
marital-status         0
occupation            966
relationship           0
race                  0
sex                   0
capital-gain           0
capital-loss           0
hours-per-week         0
native-country         274
dtype: int64
```

```
[ ]: y.isna().sum()
```

```
[ ]: income      0
      dtype: int64
```

```
[ ]: for cat in X.select_dtypes(include='object'):
      print(cat, len(X[cat].unique()), X[cat].unique())
```

```
workclass 10 ['State-gov' 'Self-emp-not-inc' 'Private' 'Federal-gov' 'Local-gov'
'?'
'Self-emp-inc' 'Without-pay' 'Never-worked' nan]
education 16 ['Bachelors' 'HS-grad' '11th' 'Masters' '9th' 'Some-college'
'Assoc-acdm'
'Assoc-voc' '7th-8th' 'Doctorate' 'Prof-school' '5th-6th' '10th'
'1st-4th' 'Preschool' '12th']
marital-status 7 ['Never-married' 'Married-civ-spouse' 'Divorced' 'Married-
spouse-absent'
'Separated' 'Married-AF-spouse' 'Widowed']
occupation 16 ['Adm-clerical' 'Exec-managerial' 'Handlers-cleaners' 'Prof-
specialty'
'Other-service' 'Sales' 'Craft-repair' 'Transport-moving'
'Farming-fishing' 'Machine-op-inspct' 'Tech-support' '?'
'Protective-serv' 'Armed-Forces' 'Priv-house-serv' nan]
relationship 6 ['Not-in-family' 'Husband' 'Wife' 'Own-child' 'Unmarried' 'Other-
relative']
race 5 ['White' 'Black' 'Asian-Pac-Islander' 'Amer-Indian-Eskimo' 'Other']
sex 2 ['Male' 'Female']
native-country 43 ['United-States' 'Cuba' 'Jamaica' 'India' '?' 'Mexico' 'South'
'Puerto-Rico' 'Honduras' 'England' 'Canada' 'Germany' 'Iran'
'Philippines' 'Italy' 'Poland' 'Columbia' 'Cambodia' 'Thailand' 'Ecuador'
'Laos' 'Taiwan' 'Haiti' 'Portugal' 'Dominican-Republic' 'El-Salvador'
'France' 'Guatemala' 'China' 'Japan' 'Yugoslavia' 'Peru'
'Outlying-US(Guam-USVI-etc)' 'Scotland' 'Trinidad&Tobago' 'Greece'
'Nicaragua' 'Vietnam' 'Hong' 'Ireland' 'Hungary' 'Holand-Netherlands' nan]
```

```
[ ]: y.value_counts()
```

```
[ ]: income
<=50K      24720
<=50K.     12435
>50K        7841
>50K.       3846
Name: count, dtype: int64
```

```
[ ]: y = y.income.str.replace('<=50K.', '<=50K')
      y = y.str.replace('>50K.', '>50K')
      y.value_counts()
```

```
[ ]: income
<=50K      37155
```

```
>50K      11687
Name: count, dtype: int64
```

**Split the dataset into Train and Test Dataset in 80:20 ratio (1 marks)**

```
[ ]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=52)
```

```
[ ]: for cat in X_train.select_dtypes(include='object'):
    print(cat, len(X_train[cat].unique()), X_train[cat].unique())
    print(len(X_test[cat].unique()), X_test[cat].unique())
```

```
workclass 10 ['?' 'Self-emp-not-inc' 'Private' 'Self-emp-inc' 'Federal-gov'
'Local-gov'
'State-gov' nan 'Never-worked' 'Without-pay']
10 ['Self-emp-not-inc' 'Self-emp-inc' 'Private' 'State-gov' 'Local-gov'
'Federal-gov' '?' nan 'Without-pay' 'Never-worked']
education 16 ['5th-6th' 'Assoc-voc' 'Assoc-acdm' 'HS-grad' 'Some-college' '9th'
'Bachelors' 'Doctorate' '7th-8th' 'Masters' '12th' '10th' '11th'
'Prof-school' '1st-4th' 'Preschool']
16 ['Masters' 'Some-college' 'Assoc-voc' '11th' '9th' '12th' 'HS-grad'
'Assoc-acdm' '10th' 'Prof-school' '7th-8th' 'Bachelors' 'Doctorate'
'Preschool' '1st-4th' '5th-6th']
marital-status 7 ['Never-married' 'Married-civ-spouse' 'Divorced' 'Widowed'
'Married-spouse-absent' 'Separated' 'Married-AF-spouse']
7 ['Married-civ-spouse' 'Never-married' 'Divorced' 'Separated' 'Widowed'
'Married-spouse-absent' 'Married-AF-spouse']
occupation 16 ['?' 'Farming-fishing' 'Craft-repair' 'Tech-support' 'Sales'
'Adm-clerical' 'Exec-managerial' 'Other-service' 'Prof-specialty'
'Machine-op-inspct' 'Transport-moving' 'Priv-house-serv' nan
'Protective-serv' 'Handlers-cleaners' 'Armed-Forces']
16 ['Sales' 'Exec-managerial' 'Adm-clerical' 'Handlers-cleaners'
'Craft-repair' 'Prof-specialty' 'Other-service' 'Tech-support'
'Machine-op-inspct' 'Farming-fishing' 'Transport-moving'
'Protective-serv' '?' nan 'Priv-house-serv' 'Armed-Forces']
relationship 6 ['Unmarried' 'Husband' 'Not-in-family' 'Wife' 'Other-relative'
'Own-child']
6 ['Husband' 'Own-child' 'Wife' 'Unmarried' 'Not-in-family' 'Other-relative']
race 5 ['White' 'Black' 'Asian-Pac-Islander' 'Other' 'Amer-Indian-Eskimo']
5 ['Asian-Pac-Islander' 'White' 'Amer-Indian-Eskimo' 'Black' 'Other']
sex 2 ['Female' 'Male']
2 ['Male' 'Female']
native-country 43 ['El-Salvador' 'United-States' 'Canada' 'Mexico' 'Cuba' '?'
'Italy'
'Philippines' 'Guatemala' 'Puerto-Rico' 'Cambodia' 'Taiwan' 'Vietnam'
'Dominican-Republic' 'China' 'India' 'Ireland' nan 'Columbia' 'England']
```

```

'Japan' 'South' 'Haiti' 'Laos' 'Jamaica' 'Peru' 'Germany' 'Poland'
'Yugoslavia' 'Hong' 'Greece' 'Outlying-US(Guam-USVI-etc)' 'Scotland'
'France' 'Iran' 'Portugal' 'Nicaragua' 'Thailand' 'Ecuador' 'Honduras'
'Hungary' 'Trinidad&Tobago' 'Holand-Netherlands']
42 ['India' 'United-States' nan 'China' 'Haiti' 'Germany' 'Mexico'
'Puerto-Rico' 'Cambodia' 'El-Salvador' 'Philippines' 'Columbia' 'Japan'
'Honduras' '?' 'Dominican-Republic' 'Portugal' 'Ecuador' 'Cuba' 'Greece'
'Vietnam' 'South' 'Guatemala' 'Hungary' 'Taiwan' 'France' 'Yugoslavia'
'Iran' 'Italy' 'Canada' 'England' 'Nicaragua' 'Poland' 'Peru'
'Trinidad&Tobago' 'Jamaica' 'Ireland' 'Scotland' 'Hong' 'Thailand'
'Outlying-US(Guam-USVI-etc)' 'Laos']

```

```
[ ]: X_train.shape, X_test.shape
```

```
[ ]: ((39073, 14), (9769, 14))
```

Create a data pipeline that does:

1. Imputation (5 marks)
2. Standardization and Scaling (5 marks)
3. Discretization (5 marks)
4. Encoding (5 marks)
5. Prediction (5 marks)

(Total 25 marks)

```

[ ]: from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder,
    OrdinalEncoder, MinMaxScaler
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.metrics import classification_report, accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import GridSearchCV

import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

```

```

[ ]: age_transformer = Pipeline([
    ('scaler', StandardScaler())
])

workclass_transformer = Pipeline([
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('encoder', OneHotEncoder(handle_unknown='ignore'))
])

fnlwgt_transformer = Pipeline([
    ('scaler', MinMaxScaler())
])

education_transformer = Pipeline([
    ('encoder', OrdinalEncoder())
])

education_num_transformer = Pipeline([
    ('scaler', StandardScaler())
])

marital_status_transformer = Pipeline([
    ('encoder', OrdinalEncoder())
])

occupation_transformer = Pipeline([
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('encoder', OrdinalEncoder())
])

relationship_transformer = Pipeline([
    ('encoder', OrdinalEncoder())
])

race_transformer = Pipeline([
    ('encoder', OrdinalEncoder())
])

sex_transformer = Pipeline([
    ('encoder', OrdinalEncoder())
])

capital_gain_transformer = Pipeline([
    ('scaler', StandardScaler())
])

capital_loss_transformer = Pipeline([

```

```

        ('scaler', StandardScaler())
    ])

hours_per_week_transformer = Pipeline([
    ('scaler', StandardScaler())
])

native_country_transformer = Pipeline([
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('encoder', OneHotEncoder(handle_unknown='ignore'))
])

preprocessor = ColumnTransformer(
    transformers=[
        ('age_transformer', age_transformer, ['age']),
        ('workclass_transformer', workclass_transformer, ['workclass']),
        ('fnlwgt_transformer', fnlwgt_transformer, ['fnlwgt']),
        ('education_transformer', education_transformer, ['education']),
        ('education_num_transformer', education_num_transformer,
↪ ['education-num']),
        ('marital_status_transformer', marital_status_transformer,
↪ ['marital-status']),
        ('occupation_transformer', occupation_transformer, ['occupation']),
        ('relationship_transformer', relationship_transformer,
↪ ['relationship']),
        ('race_transformer', race_transformer, ['race']),
        ('sex_transformer', sex_transformer, ['sex']),
        ('capital_gain_transformer', capital_gain_transformer,
↪ ['capital-gain']),
        ('capital_loss_transformer', capital_loss_transformer,
↪ ['capital-loss']),
        ('hours_per_week_transformer', hours_per_week_transformer,
↪ ['hours-per-week']),
        ('native_country_transformer', native_country_transformer,
↪ ['native-country']),
        ], remainder='passthrough',)

```

```

[ ]: le = LabelEncoder()
y_train = le.fit_transform(y_train)
y_test = le.transform(y_test)

y_train, y_test

```

```

[ ]: (array([0, 0, 0, ..., 0, 0, 0]), array([0, 1, 0, ..., 0, 0, 0]))

```





```

StandardScaler()))],
SimpleImputer(strategy='most_frequent')),
OneHotEncoder(handle_unknown='ignore'))]),
('classifier',
RandomForestClassifier(max_depth=20, n_estimators=200))))
('fnlwtg_transformer',
Pipeline(
['capital-loss']),
('hours_per_week_transformer',
Pipeline(steps=[('scaler',
['hours-per-week']),
('native_country_transformer',
Pipeline(steps=[('imputer',
('encoder',
['native-country']])))),

```

```

[ ]: print(classification_report(y_test, y_pred))
# print("Accuracy: ", accuracy_score(y_test, predictions))
print(confusion_matrix(y_test, y_pred))

cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```

	precision	recall	f1-score	support
0	0.89	0.95	0.92	7479
1	0.78	0.60	0.68	2290
accuracy			0.87	9769
macro avg	0.83	0.77	0.80	9769
weighted avg	0.86	0.87	0.86	9769

```

[[7088 391]
 [ 915 1375]]

```

