

Copy_of_House_Price_Prediction

October 19, 2024

0.1 House Price Prediction dataset taken from
<https://www.kaggle.com/competitions/house-prices-advanced-regression-techniques/d>

```
[ ]: # !unzip house-prices-advanced-regression-techniques.zip
```

Archive: house-prices-advanced-regression-techniques.zip
replace data_description.txt? [y]es, [n]o, [A]ll, [N]one, [r]ename:

```
[ ]: import pandas as pd
```

```
[ ]: house_df = pd.read_csv('train.csv')
```

```
[ ]: house_df.columns
```

```
[ ]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',  
        'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',  
        'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',  
        'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',  
        'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',  
        'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',  
        'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',  
        'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',  
        'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',  
        'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',  
        'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',  
        'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',  
        'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',  
        'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',  
        'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',  
        'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',  
        'SaleCondition', 'SalePrice'],  
        dtype='object')
```

```
[ ]: house_df.isna().sum()
```

```
[ ]: Id                0  
     MSSubClass        0  
     MSZoning          0  
     LotFrontage      259
```

```

LotArea          0
...
MoSold           0
YrSold           0
SaleType         0
SaleCondition    0
SalePrice        0
Length: 81, dtype: int64

```

```

[ ]: import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
### I will just use the columns at random but in the lab work please see what
### relationship there is between target variable and other attribute
numeric_df = house_df.select_dtypes(['int64', 'float64'])


corr = numeric_df.corr()
mask = np.triu(np.ones_like(corr, dtype=bool))
f, ax = plt.subplots(figsize=(11, 9))
cmap = sns.diverging_palette(230, 20, as_cmap=True)
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})

```

```

[ ]: <Axes: >

```



```
[ ]: plt.figure(figsize=(9, 8))
sns.distplot(house_df['SalePrice'], color='g', bins=100, hist_kws={'alpha': 0.
↪4});
```


<ipython-input-73-f6ab0e3193d7>:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(house_df['SalePrice'], color='g', bins=100, hist_kws={'alpha':
0.4});
```



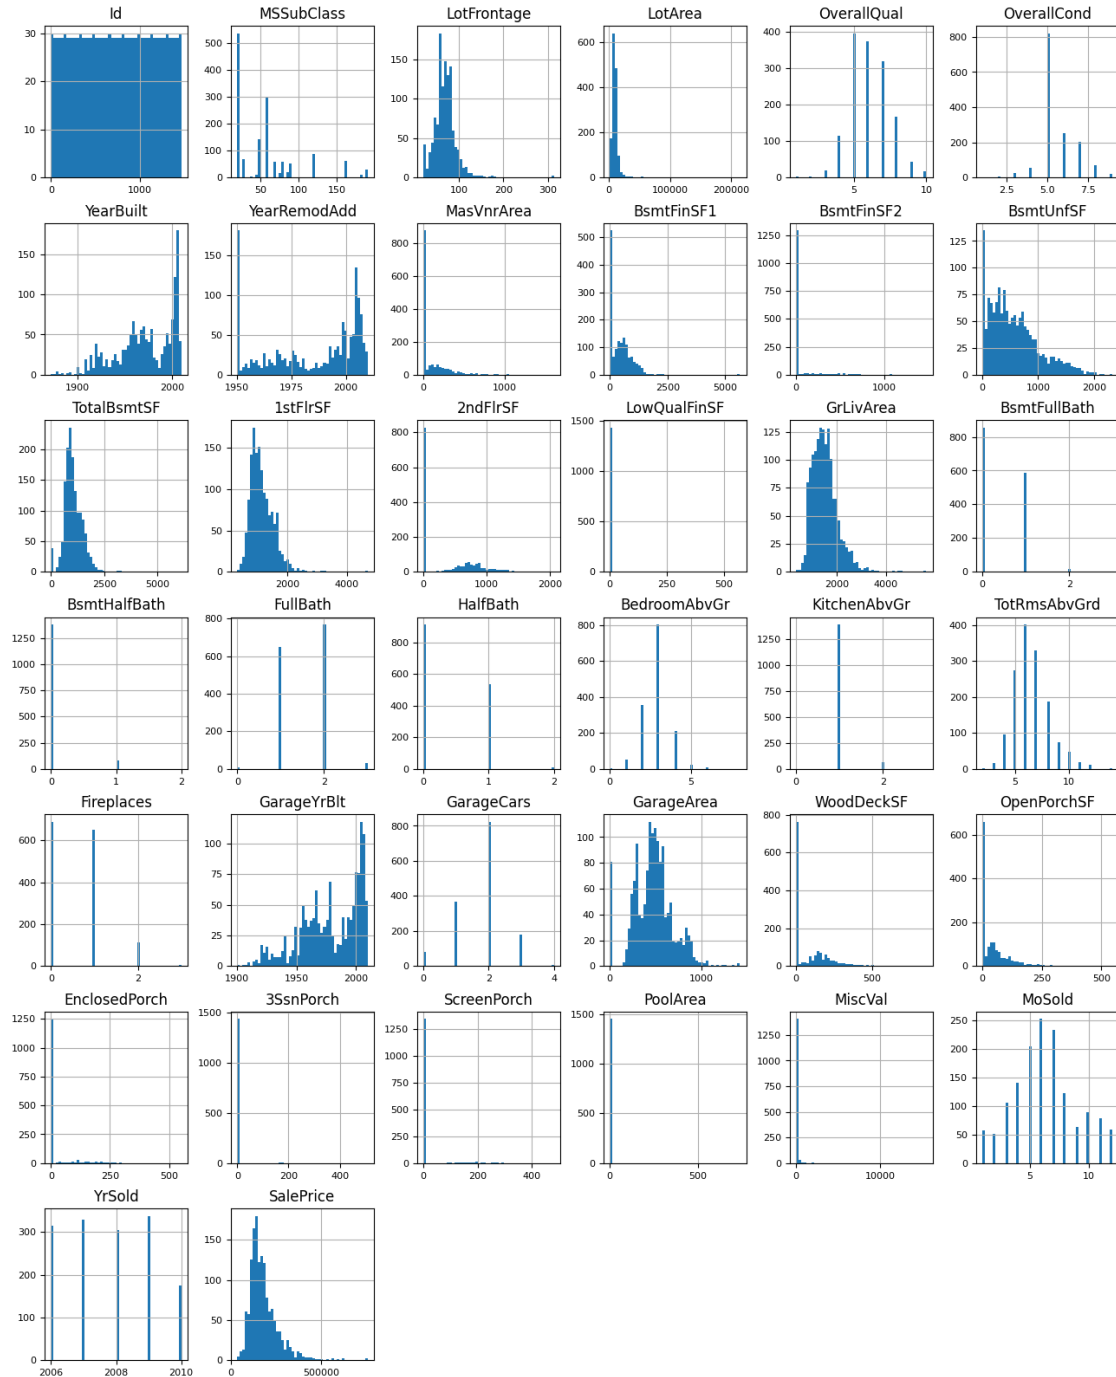
```
[ ]: numeric_df.hist(figsize=(16, 20), bins=50, xlabelsize=8, ylabelsize=8)
```

```
[ ]: array([[<Axes: title={'center': 'Id'}>,
<Axes: title={'center': 'MSSubClass'}>,
<Axes: title={'center': 'LotFrontage'}>,
<Axes: title={'center': 'LotArea'}>,
<Axes: title={'center': 'OverallQual'}>,
<Axes: title={'center': 'OverallCond'}>],
[<Axes: title={'center': 'YearBuilt'}>,
<Axes: title={'center': 'YearRemodAdd'}>,
<Axes: title={'center': 'MasVnrArea'}>,
<Axes: title={'center': 'BsmtFinSF1'}>,
<Axes: title={'center': 'BsmtFinSF2'}>,
<Axes: title={'center': 'BsmtUnfSF'}>],
[<Axes: title={'center': 'TotalBsmtSF'}>],
```

```

<Axes: title={'center': '1stFlrSF'}>,
<Axes: title={'center': '2ndFlrSF'}>,
<Axes: title={'center': 'LowQualFinSF'}>,
<Axes: title={'center': 'GrLivArea'}>,
<Axes: title={'center': 'BsmtFullBath'}>],
[<Axes: title={'center': 'BsmtHalfBath'}>,
<Axes: title={'center': 'FullBath'}>,
<Axes: title={'center': 'HalfBath'}>,
<Axes: title={'center': 'BedroomAbvGr'}>,
<Axes: title={'center': 'KitchenAbvGr'}>,
<Axes: title={'center': 'TotRmsAbvGrd'}>],
[<Axes: title={'center': 'Fireplaces'}>,
<Axes: title={'center': 'GarageYrBlt'}>,
<Axes: title={'center': 'GarageCars'}>,
<Axes: title={'center': 'GarageArea'}>,
<Axes: title={'center': 'WoodDeckSF'}>,
<Axes: title={'center': 'OpenPorchSF'}>],
[<Axes: title={'center': 'EnclosedPorch'}>,
<Axes: title={'center': '3SsnPorch'}>,
<Axes: title={'center': 'ScreenPorch'}>,
<Axes: title={'center': 'PoolArea'}>,
<Axes: title={'center': 'MiscVal'}>,
<Axes: title={'center': 'MoSold'}>],
[<Axes: title={'center': 'YrSold'}>,
<Axes: title={'center': 'SalePrice'}>, <Axes: >, <Axes: >,
<Axes: >, <Axes: >]], dtype=object)

```



```
[ ]: house_df.head()
```

```
[ ]:   Id  MSSubClass MSZoning  LotFrontage  LotArea  Street  Alley  LotShape  \
0    1           60      RL         65.0     8450   Pave   NaN     Reg
1    2           20      RL         80.0     9600   Pave   NaN     Reg
```

2	3	60	RL	68.0	11250	Pave	NaN	IR1
3	4	70	RL	60.0	9550	Pave	NaN	IR1
4	5	60	RL	84.0	14260	Pave	NaN	IR1

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	\
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	5	
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0	9	
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0	12	

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

```
[ ]: columns = ['MSSubClass', 'MSZoning', 'LotArea', 'LotFrontage', 'LotShape',
↳ 'Street', 'SaleCondition', 'OverallQual', 'YearBuilt', 'YearRemodAdd',
↳ 'MasVnrArea', 'BsmtFinSF1', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF',
↳ 'GrLivArea', 'FullBath', 'HalfBath', 'GarageYrBlt', 'GarageCars',
↳ 'GarageArea', 'WoodDeckSF', 'OpenPorchSF', 'SalePrice']
```

```
[ ]: house_df[columns].head()
```

	MSSubClass	MSZoning	LotArea	LotFrontage	LotShape	Street	SaleCondition	\
0	60	RL	8450	65.0	Reg	Pave	Normal	
1	20	RL	9600	80.0	Reg	Pave	Normal	
2	60	RL	11250	68.0	IR1	Pave	Normal	
3	70	RL	9550	60.0	IR1	Pave	Abnorml	
4	60	RL	14260	84.0	IR1	Pave	Normal	

	OverallQual	YearBuilt	YearRemodAdd	...	2ndFlrSF	GrLivArea	FullBath	\
0	7	2003	2003	...	854	1710	2	
1	6	1976	1976	...	0	1262	2	
2	7	2001	2002	...	866	1786	2	
3	7	1915	1970	...	756	1717	1	
4	8	2000	2000	...	1053	2198	2	

	HalfBath	GarageYrBlt	GarageCars	GarageArea	WoodDeckSF	OpenPorchSF	\
0	1	2003.0	2	548	0	61	
1	0	1976.0	2	460	298	0	
2	1	2001.0	2	608	0	42	
3	0	1998.0	3	642	0	35	

4	1	2000.0	3	836	192	84
---	---	--------	---	-----	-----	----

	SalePrice
0	208500
1	181500
2	223500
3	140000
4	250000

[5 rows x 24 columns]

```
[ ]: house_df[columns].isna().sum()
```

```
[ ]: MSSubClass      0
     MSZoning        0
     LotArea         0
     LotFrontage    259
     LotShape        0
     Street          0
     SaleCondition   0
     OverallQual     0
     YearBuilt       0
     YearRemodAdd    0
     MasVnrArea      8
     BsmtFinSF1      0
     TotalBsmtSF     0
     1stFlrSF        0
     2ndFlrSF        0
     GrLivArea       0
     FullBath        0
     HalfBath        0
     GarageYrBlt     81
     GarageCars      0
     GarageArea      0
     WoodDeckSF      0
     OpenPorchSF     0
     SalePrice       0
     dtype: int64
```

```
[ ]: null_cols = ['LotFrontage', 'GarageYrBlt', 'MasVnrArea']

     for col in null_cols:
         house_df[col].fillna(house_df[col].mean(), inplace=True)
```

<ipython-input-79-caa3c8083d9b>:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work

because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing `'df[col].method(value, inplace=True)'`, try using `'df.method({col: value}, inplace=True)'` or `df[col] = df[col].method(value)` instead, to perform the operation inplace on the original object.

```
house_df[col].fillna(house_df[col].mean(), inplace=True)
```

```
[ ]: house_df[columns].isna().sum()
```

```
[ ]: MSSubClass      0
      MSZoning       0
      LotArea        0
      LotFrontage    0
      LotShape       0
      Street         0
      SaleCondition   0
      OverallQual     0
      YearBuilt       0
      YearRemodAdd    0
      MasVnrArea      0
      BsmtFinSF1      0
      TotalBsmtSF     0
      1stFlrSF        0
      2ndFlrSF        0
      GrLivArea       0
      FullBath        0
      HalfBath        0
      GarageYrBlt     0
      GarageCars      0
      GarageArea      0
      WoodDeckSF      0
      OpenPorchSF     0
      SalePrice       0
      dtype: int64
```

```
[ ]: X = house_df[['MSSubClass', 'MSZoning', 'LotArea', 'LotFrontage', 'LotShape',
↳ 'Street', 'SaleCondition', 'OverallQual', 'YearBuilt', 'YearRemodAdd',
↳ 'MasVnrArea', 'BsmtFinSF1', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF',
↳ 'GrLivArea', 'FullBath', 'HalfBath', 'GarageYrBlt', 'GarageCars',
↳ 'GarageArea', 'WoodDeckSF', 'OpenPorchSF']]
      Y = house_df['SalePrice']
```

```
[ ]: X
```

[]:	MSSubClass	MSZoning	LotArea	LotFrontage	LotShape	Street	SaleCondition	\
0	60	RL	8450	65.0	Reg	Pave	Normal	
1	20	RL	9600	80.0	Reg	Pave	Normal	
2	60	RL	11250	68.0	IR1	Pave	Normal	
3	70	RL	9550	60.0	IR1	Pave	Abnorml	
4	60	RL	14260	84.0	IR1	Pave	Normal	
...	
1455	60	RL	7917	62.0	Reg	Pave	Normal	
1456	20	RL	13175	85.0	Reg	Pave	Normal	
1457	70	RL	9042	66.0	Reg	Pave	Normal	
1458	20	RL	9717	68.0	Reg	Pave	Normal	
1459	20	RL	9937	75.0	Reg	Pave	Normal	

	OverallQual	YearBuilt	YearRemodAdd	...	1stFlrSF	2ndFlrSF	\
0	7	2003	2003	...	856	854	
1	6	1976	1976	...	1262	0	
2	7	2001	2002	...	920	866	
3	7	1915	1970	...	961	756	
4	8	2000	2000	...	1145	1053	
...	
1455	6	1999	2000	...	953	694	
1456	6	1978	1988	...	2073	0	
1457	7	1941	2006	...	1188	1152	
1458	5	1950	1996	...	1078	0	
1459	5	1965	1965	...	1256	0	

	GrLivArea	FullBath	HalfBath	GarageYrBlt	GarageCars	GarageArea	\
0	1710	2	1	2003.0	2	548	
1	1262	2	0	1976.0	2	460	
2	1786	2	1	2001.0	2	608	
3	1717	1	0	1998.0	3	642	
4	2198	2	1	2000.0	3	836	
...	
1455	1647	2	1	1999.0	2	460	
1456	2073	2	0	1978.0	2	500	
1457	2340	2	0	1941.0	1	252	
1458	1078	1	0	1950.0	1	240	
1459	1256	1	1	1965.0	1	276	

	WoodDeckSF	OpenPorchSF
0	0	61
1	298	0
2	0	42
3	0	35
4	192	84
...
1455	0	40

1456	349	0
1457	0	60
1458	366	0
1459	736	68

[1460 rows x 23 columns]

```
[ ]: X.isna().sum()
```

```
[ ]: MSSubClass      0
      MSZoning       0
      LotArea        0
      LotFrontage    0
      LotShape       0
      Street         0
      SaleCondition  0
      OverallQual    0
      YearBuilt      0
      YearRemodAdd   0
      MasVnrArea     0
      BsmtFinSF1     0
      TotalBsmtSF    0
      1stFlrSF       0
      2ndFlrSF       0
      GrLivArea      0
      FullBath       0
      HalfBath       0
      GarageYrBlt    0
      GarageCars     0
      GarageArea     0
      WoodDeckSF     0
      OpenPorchSF    0
      dtype: int64
```

```
[ ]: import matplotlib.pyplot as plt
```

```
[ ]: import seaborn as sns
```

```
[ ]: sns.distplot(X['LotFrontage']) #only works for old version of seaborn use ↵
      ↪ different method histplot/kdeplot
```

<ipython-input-86-e9dcc26ed7ac>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(X['LotFrontage']) #only works for old version of seaborn use  
different method histplot/kdeplot
```

```
[ ]: <Axes: xlabel='LotFrontage', ylabel='Density'>
```



```
[ ]: median_lotfrontage = X['LotFrontage'].median()
```

```
[ ]: median_lotfrontage
```

```
[ ]: 70.04995836802665
```

```
[ ]: X['LotFrontage'].fillna(median_lotfrontage,inplace=True)
```

<ipython-input-89-ca9726a0c5bd>:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing `'df[col].method(value, inplace=True)'`, try using `'df.method({col: value}, inplace=True)'` or `df[col] = df[col].method(value)` instead, to perform the operation inplace on the original object.

```
X['LotFrontage'].fillna(median_lotfrontage,inplace=True)
<ipython-input-89-ca9726a0c5bd>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
X['LotFrontage'].fillna(median_lotfrontage,inplace=True)
```

```
[ ]: X.isna().sum()
```

```
[ ]: MSSubClass      0
     MSZoning        0
     LotArea         0
     LotFrontage     0
     LotShape        0
     Street          0
     SaleCondition   0
     OverallQual     0
     YearBuilt       0
     YearRemodAdd    0
     MasVnrArea      0
     BsmtFinSF1      0
     TotalBsmtSF     0
     1stFlrSF        0
     2ndFlrSF        0
     GrLivArea       0
     FullBath        0
     HalfBath        0
     GarageYrBlt     0
     GarageCars      0
     GarageArea      0
     WoodDeckSF      0
     OpenPorchSF     0
     dtype: int64
```

0.1.1 In categorical values that needs to be converted to numerical form and

0.1.2 In numerical values scale it down and split it

```
[ ]: from sklearn.model_selection import train_test_split
     from sklearn.preprocessing import StandardScaler
```

```
[ ]: X.dtypes
```

```
[ ]: MSSubClass      int64
     MSZoning        object
     LotArea         int64
     LotFrontage     float64
     LotShape        object
     Street          object
     SaleCondition   object
     OverallQual     int64
     YearBuilt       int64
     YearRemodAdd    int64
     MasVnrArea      float64
     BsmtFinSF1      int64
     TotalBsmtSF     int64
     1stFlrSF        int64
     2ndFlrSF        int64
     GrLivArea       int64
     FullBath        int64
     HalfBath        int64
     GarageYrBlt     float64
     GarageCars      int64
     GarageArea      int64
     WoodDeckSF      int64
     OpenPorchSF     int64
     dtype: object
```

```
[ ]: scaler = StandardScaler()
```

```
[ ]: numerical_columns = ['MSSubClass', 'LotArea', 'LotFrontage', 'OverallQual',
    ↪ 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'TotalBsmtSF',
    ↪ '1stFlrSF', '2ndFlrSF', 'GrLivArea', 'FullBath', 'HalfBath', 'GarageYrBlt',
    ↪ 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF']
     categorical_columns = ['MSZoning', 'LotShape', 'Street', 'SaleCondition']
```

```
[ ]: X[numerical_columns] = scaler.fit_transform(X[numerical_columns])
```

<ipython-input-95-cb7cdb47f1f9>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
X[numerical_columns] = scaler.fit_transform(X[numerical_columns])

```
[ ]: X['MSZoning'].unique()
```

```
[ ]: array(['RL', 'RM', 'C (all)', 'FV', 'RH'], dtype=object)
```

```
[ ]: X['LotShape'].unique()
[ ]: array(['Reg', 'IR1', 'IR2', 'IR3'], dtype=object)
[ ]: X['Street'].unique()
[ ]: array(['Pave', 'Grvl'], dtype=object)
[ ]: X['SaleCondition'].unique()
[ ]: array(['Normal', 'Abnorml', 'Partial', 'AdjLand', 'Alloca', 'Family'],
          dtype=object)
[ ]: encoded_value = pd.get_dummies(X[categorical_columns], sparse=False).
      ↪astype(int) #Alternative way to one hot encode categorical values
[ ]: X = pd.concat([X, encoded_value],axis=1)
[ ]: X.drop(columns=categorical_columns,inplace=True)
[ ]: X_train, X_test, y_train, y_test = train_test_split(X,Y,train_size=0.8)
[ ]: X_train.shape, X_test.shape, X.shape
[ ]: ((1168, 36), (292, 36), (1460, 36))
```

0.2 Define model for evaluation

```
[ ]: from sklearn.linear_model import LinearRegression
     from sklearn.ensemble import RandomForestRegressor
[ ]: linear = LinearRegression()
     random_forest = RandomForestRegressor()
[ ]: from sklearn.model_selection import KFold, cross_val_score
[ ]: algorithms = [linear, random_forest]
[ ]: algorithm_names = ['linear_regression', 'random forest']
[ ]: kfold = KFold(n_splits=5, shuffle=True)
[ ]: for i,model in enumerate(algorithms):
      scores = cross_val_score(model, X_train, y_train, cv=kfold,
      ↪scoring='neg_mean_squared_error')
      print(f"{algorithm_names[i]} - Score:{scores} - Mean Score{scores.mean()}")

linear_regression - Score: [-1.02688362e+09 -1.47332768e+09 -2.85422088e+09
-9.47078572e+08
```

```
-1.50245785e+09] - Mean Score-1560793720.2089233
random forest - Score:[-6.34254723e+08 -1.49475153e+09 -5.53717701e+08
-1.25223134e+09
-8.05098016e+08] - Mean Score-948010663.7559055
```

```
[ ]: from sklearn.model_selection import GridSearchCV
```

```
[ ]: param_grid = {
    'max_depth': [5, 10, 15],
    'n_estimators': [5, 6, 7, 8, 9, 10, 11, 12, 13, 15]
}
```

```
[ ]: grid = GridSearchCV(estimator=random_forest, param_grid=param_grid, cv=kfold,
    ↪scoring='neg_mean_squared_error')
```

```
[ ]: grid.fit(X_train, y_train)
```

```
/usr/local/lib/python3.10/dist-packages/numpy/ma/core.py:2820: RuntimeWarning:
invalid value encountered in cast
```

```
    _data = np.array(data, dtype=dtype, copy=copy,
```

```
[ ]: GridSearchCV(cv=KFold(n_splits=5, random_state=None, shuffle=True),
    estimator=RandomForestRegressor(),
    param_grid={'max_depth': [5, 10, 15],
                'n_estimators': [5, 6, 7, 8, 9, 10, 11, 12, 13, 15]},
    scoring='neg_mean_squared_error')
```

```
[ ]: grid.best_params_
```

```
[ ]: {'max_depth': 10, 'n_estimators': 11}
```

```
[ ]: grid.best_score_
```

```
[ ]: -952854163.2083702
```

```
[ ]: yhat = grid.predict(X_test)
```

```
[ ]: from sklearn.metrics import mean_squared_error
```

```
[ ]: mean_squared_error(y_test, yhat)
```

```
[ ]: 674499445.7271372
```

```
[ ]: rf = grid.best_estimator_
```

```
[ ]: rf.feature_importances_
```

```
[ ]: array([2.23487517e-03, 1.63126416e-02, 7.86446822e-03, 5.65746555e-01,
    1.10634195e-02, 1.39303711e-02, 5.68856760e-03, 2.83416752e-02,
```



```


3.84127177e-02, 2.46239340e-02, 5.42631552e-02, 1.09831676e-01,
1.29365578e-03, 1.02920292e-03, 6.54701874e-03, 2.67570357e-02,
5.95508356e-02, 8.16524694e-03, 8.34737969e-03, 4.72864434e-04,
2.88807749e-04, 4.97098318e-05, 1.18677436e-03, 2.61996939e-03,
3.86548875e-04, 1.61504808e-04, 7.48374244e-06, 1.08429694e-03,
1.38070688e-05, 1.20365826e-06, 6.12056197e-04, 0.00000000e+00,
2.26699342e-05, 1.12559084e-03, 1.17657771e-03, 7.85702989e-04])

```

```
[ ]: sorted_idx = rf.feature_importances_.argsort()
```

```
[ ]: plt.barh(X.columns[sorted_idx], rf.feature_importances_[sorted_idx])
```

```
[ ]: <BarContainer object of 36 artists>
```



1 Replicate this in LeaveOneOut

```
[ ]: # Code here
from sklearn.model_selection import LeaveOneOut

loo = LeaveOneOut()

random_forest = RandomForestRegressor()

param_grid = {

```

```

    'max_depth': [5, 10],
    'n_estimators': [5, 10, 15]
}

grid = GridSearchCV(estimator=random_forest, param_grid=param_grid, cv=loo,
                    scoring='neg_mean_squared_error')

```

```
[ ]: grid.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(cv=LeaveOneOut(), estimator=RandomForestRegressor(),
                 param_grid={'max_depth': [5, 10], 'n_estimators': [5, 10, 15]},
                 scoring='neg_mean_squared_error')
```

```
[ ]: print(grid.best_params_)
     print(grid.best_score_)

     yhat = grid.predict(X_test)
     mean_squared_error(y_test, yhat)
```

```

{'max_depth': 10, 'n_estimators': 10}
-993319189.5942743

```

```
[ ]: 789141213.2074965
```

2 Analyze

Analyze here the difference between the two and possible reasons why they are different:

The K-Fold took a minute with gridsearchcv to find best model parameters, while LeaveOneOut was taking a lot of time, that's why I reduced the number of param_grid parameters

LOO is an extreme case of KFold where k equals the number of samples in the dataset. Each sample is used as the test set exactly once, while the rest of the dataset (all other samples) is used for training.

KFold often has a balance between bias and variance, especially with a reasonably chosen value for k (like 5 or 10), while LOO gives a very low-bias estimate (since the model is almost always trained on the entire dataset), can suffer from high variance, meaning that test errors can fluctuate significantly between iterations, especially on noisy data.

It is recommended to use LOO when the dataset is small, since leaving out one sample at a time doesn't lose too much information, and the higher variance can be tolerable.

Experiment with types of K-Fold cross-validations.

```
[ ]: #Code here
     from sklearn.model_selection import GridSearchCV, RepeatedKFold
     from sklearn.ensemble import RandomForestRegressor

     repeated_kfold = RepeatedKFold(n_splits=5, n_repeats=3, random_state=42)
```

```

param_grid = {
    'max_depth': [5, 10, 15],
    'n_estimators': [50, 100, 200]
}

model = RandomForestRegressor()

grid = GridSearchCV(estimator=model, param_grid=param_grid, cv=repeated_kfold,
    ↪scoring='neg_mean_squared_error')

grid.fit(X_train, y_train)

print("Best parameters:", grid.best_params_)

print("Best score:", grid.best_score_)
yhat = grid.predict(X_test)
mean_squared_error(y_test, yhat)

```

Best parameters: {'max_depth': 15, 'n_estimators': 200}

Best score: -1021803092.167171

[]: 650606371.9831071

```

[ ]: from sklearn.model_selection import GridSearchCV, TimeSeriesSplit
from sklearn.ensemble import RandomForestRegressor

tscv = TimeSeriesSplit(n_splits=5)

param_grid = {
    'max_depth': [5, 10, 15],
    'n_estimators': [50, 100, 200]
}

model = RandomForestRegressor()

grid = GridSearchCV(estimator=model, param_grid=param_grid, cv=tscv,
    ↪scoring='neg_mean_squared_error')

grid.fit(X_train, y_train)

print("Best parameters:", grid.best_params_)
print(grid.best_score_)

yhat = grid.predict(X_test)
mean_squared_error(y_test, yhat)

```

Best parameters: {'max_depth': 10, 'n_estimators': 100}

-1169338831.673592

[]: 665601982.8576424

2.1 Model this in Linear Regression (Conduct Grid Search for Ridge, Lasso)

```
[ ]: #Code here
from sklearn.linear_model import Ridge, Lasso
from sklearn.datasets import make_regression

ridge_param_grid = {
    'alpha': np.logspace(-3, 3, 7) # Regularization strength
}

ridge_model = Ridge()

kfold = KFold(n_splits=5, shuffle=True, random_state=42)

ridge_grid = GridSearchCV(estimator=ridge_model, param_grid=ridge_param_grid,
                           cv=kfold, scoring='neg_mean_squared_error')

ridge_grid.fit(X_train, y_train)

print("Best parameters for Ridge:", ridge_grid.best_params_)
print("Best score for Ridge:", ridge_grid.best_score_)

print("SCORE: ", ridge_grid.score(X_test, y_test))
```

Best parameters for Ridge: {'alpha': 100.0}

Best score for Ridge: -1606732817.4639132

SCORE: -938248170.5369987

```
[ ]: lasso_param_grid = {
    'alpha': np.logspace(-3, 3, 7) # Regularization strength
}

lasso_model = Lasso()

lasso_grid = GridSearchCV(estimator=lasso_model, param_grid=lasso_param_grid,
                           cv=kfold, scoring='neg_mean_squared_error')

lasso_grid.fit(X_train, y_train)

print("Best parameters for Lasso:", lasso_grid.best_params_)
print("Best score for Lasso:", lasso_grid.best_score_)
print("SCORE: ", lasso_grid.score(X_test, y_test))
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 5.551e+11, tolerance: 5.429e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 5.065e+11, tolerance: 6.293e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 5.509e+11, tolerance: 5.749e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 5.798e+11, tolerance: 6.364e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 6.053e+11, tolerance: 6.045e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 5.177e+11, tolerance: 5.429e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.929e+11, tolerance: 6.293e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.626e+11, tolerance: 5.749e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:

```

Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 4.512e+11, tolerance: 6.364e+08

```
model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 5.741e+11, tolerance: 6.045e+08
```

```
model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.020e+11, tolerance: 5.429e+08
```

```
model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 4.292e+10, tolerance: 5.749e+08
```

```
model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.519e+09, tolerance: 6.364e+08
```

```
model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.269e+11, tolerance: 6.045e+08
```

```
model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 6.952e+08, tolerance: 5.749e+08
```

```
model = cd_fast.enet_coordinate_descent(
```

Best parameters for Lasso: {'alpha': 1000.0}

Best score for Lasso: -1599765745.3049054

SCORE: -966231868.7686061

Analyze and state the difference with the two (if any)

Analyze here: Ridge:

- Best suited for scenarios where all features are believed to contribute to the outcome, and

there is multicollinearity among the features. For example, in cases like image processing or when dealing with polynomial regression.

Lasso:

- More effective in situations where feature selection is needed. It is beneficial when the model needs to be interpretable, or when you suspect that many features are irrelevant

Ridge and Lasso regression are valuable tools in the machine learning toolbox, each with its unique strengths. The choice between them often depends on the specific characteristics of the dataset and the goals of the analysis:

- Use Ridge when we believe all features should contribute to the prediction, especially in the presence of multicollinearity.
- Use Lasso when we want to reduce the number of features by eliminating less important ones, leading to a simpler and more interpretable model.

2.2 Try Lasso/ Ridge with a polynomial Feature

```
[ ]: from sklearn.preprocessing import PolynomialFeatures

poly = PolynomialFeatures(degree=2)
X_train_poly = poly.fit_transform(X_train)
X_test_poly = poly.transform(X_test)

ridge_param_grid = {
    'alpha': np.logspace(-3, 3, 7) # Regularization strength
}

ridge_model = Ridge()

kfold = KFold(n_splits=5, shuffle=True, random_state=42)

ridge_grid = GridSearchCV(estimator=ridge_model, param_grid=ridge_param_grid,
                           cv=kfold, scoring='neg_mean_squared_error')

ridge_grid.fit(X_train_poly, y_train)

print("Best parameters for Ridge:", ridge_grid.best_params_)
print("Best score for Ridge:", ridge_grid.best_score_)

print("SCORE: ", ridge_grid.score(X_test_poly, y_test))
```

Best parameters for Ridge: {'alpha': 1000.0}

Best score for Ridge: -1270653031.4742155

SCORE: -640853884.7209431

```
[ ]: lasso_param_grid = {
    'alpha': np.logspace(-3, 3, 7) # Regularization strength
}
```

```

lasso_model = Lasso()

lasso_grid = GridSearchCV(estimator=lasso_model, param_grid=lasso_param_grid,
                           cv=kfold, scoring='neg_mean_squared_error')

lasso_grid.fit(X_train_poly, y_train)

print("Best parameters for Lasso:", lasso_grid.best_params_)
print("Best score for Lasso:", lasso_grid.best_score_)
print("SCORE: ", lasso_grid.score(X_test_poly, y_test))

```

```

/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.295e+11, tolerance: 5.429e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.364e+11, tolerance: 6.293e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.267e+11, tolerance: 5.749e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.142e+11, tolerance: 6.364e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.356e+11, tolerance: 6.045e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.295e+11, tolerance: 5.429e+08
    model = cd_fast.enet_coordinate_descent(

```



```

/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.365e+11, tolerance: 6.293e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.267e+11, tolerance: 5.749e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.143e+11, tolerance: 6.364e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.356e+11, tolerance: 6.045e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.299e+11, tolerance: 5.429e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.368e+11, tolerance: 6.293e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.266e+11, tolerance: 5.749e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.149e+11, tolerance: 6.364e+08
    model = cd_fast.enet_coordinate_descent(

```

```

/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.357e+11, tolerance: 6.045e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.311e+11, tolerance: 5.429e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.393e+11, tolerance: 6.293e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.287e+11, tolerance: 5.749e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.183e+11, tolerance: 6.364e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.369e+11, tolerance: 6.045e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 9.984e+10, tolerance: 5.429e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.456e+11, tolerance: 6.293e+08
    model = cd_fast.enet_coordinate_descent(

```

```

/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.193e+11, tolerance: 5.749e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.289e+11, tolerance: 6.364e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.406e+11, tolerance: 6.045e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.052e+10, tolerance: 5.429e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.150e+10, tolerance: 6.293e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.884e+10, tolerance: 5.749e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.207e+10, tolerance: 6.364e+08
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.990e+10, tolerance: 6.045e+08
    model = cd_fast.enet_coordinate_descent(

```

```
Best parameters for Lasso: {'alpha': 1000.0}
Best score for Lasso: -1212745280.7858984
SCORE:  -648673265.3379784
```

Analyze and state the difference with the model with/out polynomial feature

:Analyze here..

Incorporating polynomial features into regression models can significantly improve their ability to capture complex relationships within the data. However, it also introduces challenges related to overfitting, model complexity, and interpretability.

Choosing the Right Model: - The decision to use polynomial features should be based on a thorough understanding of the data, its relationships, and the specific goals of the modeling task. - Cross-validation and evaluation metrics should guide the decision, ensuring the selected model generalizes well to unseen data.