

Term Project Proposal

# Towards Intelligent Financial Fraud Detection: Exploring Modern Techniques for Existing Threats

Professor: Chantri Polprasert

Presented by:  
Team Duo

st124997 - Suryansh Srivastava  
st125457 - Ulugbek Shernazarov



# Overview

**01** Introduction

**02** Problem Statement

**03** Related Work

**04** Methodology

**05** Model Evaluation

**06** Discussions

**07** Conclusion

**4.1** EDA

**4.2** Data Preprocessing

**5.1** Pipeline

**5.2** Machine Learning Model



# ML Pipeline

```
# Modeling
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)
X_train.shape, X_test.shape
```

```
((5090096, 6), (1272524, 6))
```

Train test split

SMOTE

Pipeline design

BaseLine Model

# ML Pipeline

```
categorical_transformer = Pipeline(steps=[
    ('encoder', OneHotEncoder()),
])

# No need for imputation - no missing values
numerical_transformer = Pipeline(steps=[
    ('scaler', StandardScaler())
])

preprocessor = ColumnTransformer(transformers=[
    ('num', numerical_transformer, num_columns),
    ('cat', categorical_transformer, cat_columns),
])

X_train_preprocessed = preprocessor.fit_transform(X_train)
X_test_preprocessed = preprocessor.transform(X_test)
```

Train test split

Pipeline design

SMOTE

BaseLine Model

# ML Pipeline

```
smote = SMOTE(random_state=42)
X_train_preprocessed_resampled, y_train_resampled = smote.fit_resample(X_train_preprocessed, y_train)

print(f"Before smote class distribution: {collections.Counter(y_train)}\nAfter: {collections.Counter(y_train_resampled)}")
```

```
Before smote class distribution: Counter({0: 5083526, 1: 6570})
After: Counter({0: 5083526, 1: 5083526})
```

Train test split

SMOTE

Pipeline design

BaseLine Model

# ML Pipeline

```
model = LogisticRegression(class_weight='balanced', random_state=42)
model.fit(X_train_preprocessed_resampled, y_train_resampled)

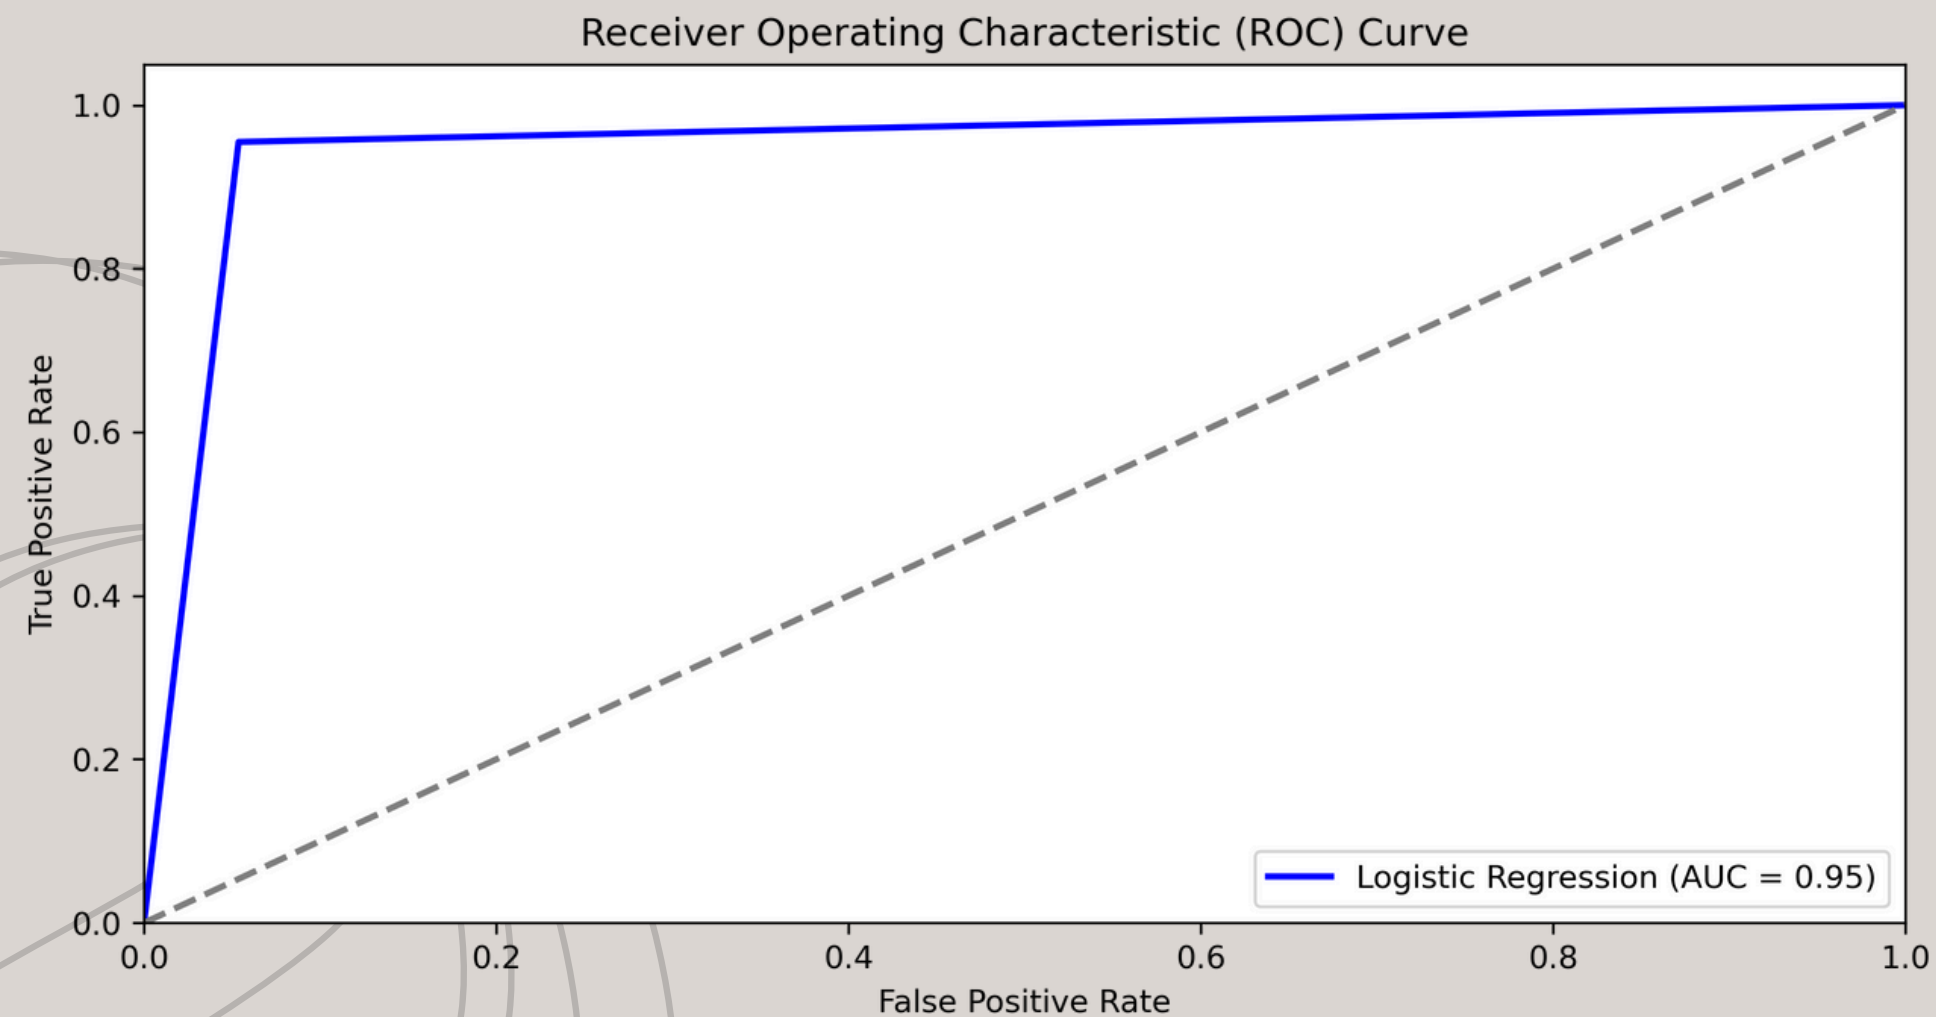
y_pred = model.predict(X_test_preprocessed)
```

Train test split

Pipeline design

SMOTE

BaseLine Model



# ML Models

## Models:

- Logistic Regression, Support Vector Machine Classifier, Random Forest
- Gradient Boosting Classifier
- Feedforward Neural Network
- Kolmogorov Arnolds Networks



# Param search

```
param_grid = [  
    {  
        'classifier': [LogisticRegression(max_iter=500)],  
        'classifier__C': [0.1, 1, 10]  
    },  
    {  
        'classifier': [SVC()],  
        'classifier__C': [0.1, 1, 10],  
        'classifier__kernel': ['linear', 'rbf']  
    },  
    {  
        'classifier': [RandomForestClassifier()],  
        'classifier__n_estimators': [100, 200],  
        'classifier__max_depth': [10, 20],  
        'classifier__min_samples_split': [2, 5, 10],  
        'classifier__min_samples_leaf': [1, 2, 4],  
        'classifier__bootstrap': [True, False]  
    },  
    {  
        'classifier': [GradientBoostingClassifier()],  
        'classifier__n_estimators': [100, 200],  
        'classifier__learning_rate': [0.01, 0.1, 0.2],  
        'classifier__max_depth': [3, 5, 7],  
        'classifier__min_samples_split': [2, 5],  
        'classifier__min_samples_leaf': [1, 2]  
    }  
]
```

```
model_pipeline = Pipeline([  
    ('preprocessor', preprocessor),  
    ('classifier', LogisticRegression())  
])
```

```
sgf=StratifiedGroupKFold(n_splits=5)  
grid_search = GridSearchCV(  
    estimator=model_pipeline,  
    param_grid=param_grid,  
    cv=sgf,  
    scoring='roc_auc',  
    verbose=2,  
    n_jobs=-1 # Use parallel processing  
)
```



# ML Models

Data size: >10m rows

Grid search run took almost 50 hours



To avoid data leakage, on the fly  
preprocessing pipeline is given to each  
stratified kfold

Scoring focus is auc\_roc

Also run gridsearch for each model to  
find best params per model



# Federated Learning



- Decentralized Machine Learning
- Data Privacy Preserved
- Training Across Multiple Devices
- No Raw Data Sharing
- Collaborative Model Updates
- Secure Aggregation
- Ideal for Sensitive Data

# Neural network

```
class Ff1(nn.Module):
    def __init__(self, num_features, embed_dim):
        super(Ff1, self).__init__()
        self.embedding = nn.Embedding(num_features, embed_dim)
        self.fc1 = nn.Linear(embed_dim, 64)
        self.fc2 = nn.Linear(64, 32)
        self.fc3 = nn.Linear(32, 2)

    def forward(self, x):
        x = self.embedding(x).view(-1, self.embed_dim)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

```
class Ff2(nn.Module):
    def __init__(self, input_dim, hidden_dims=[64, 32], dropout=0.2):
        super(Ff2, self).__init__()
        self.layers = nn.ModuleList()
        self.batch_norms = nn.ModuleList()

        self.layers.append(nn.Linear(input_dim, hidden_dims[0]))
        self.batch_norms.append(nn.BatchNorm1d(hidden_dims[0]))

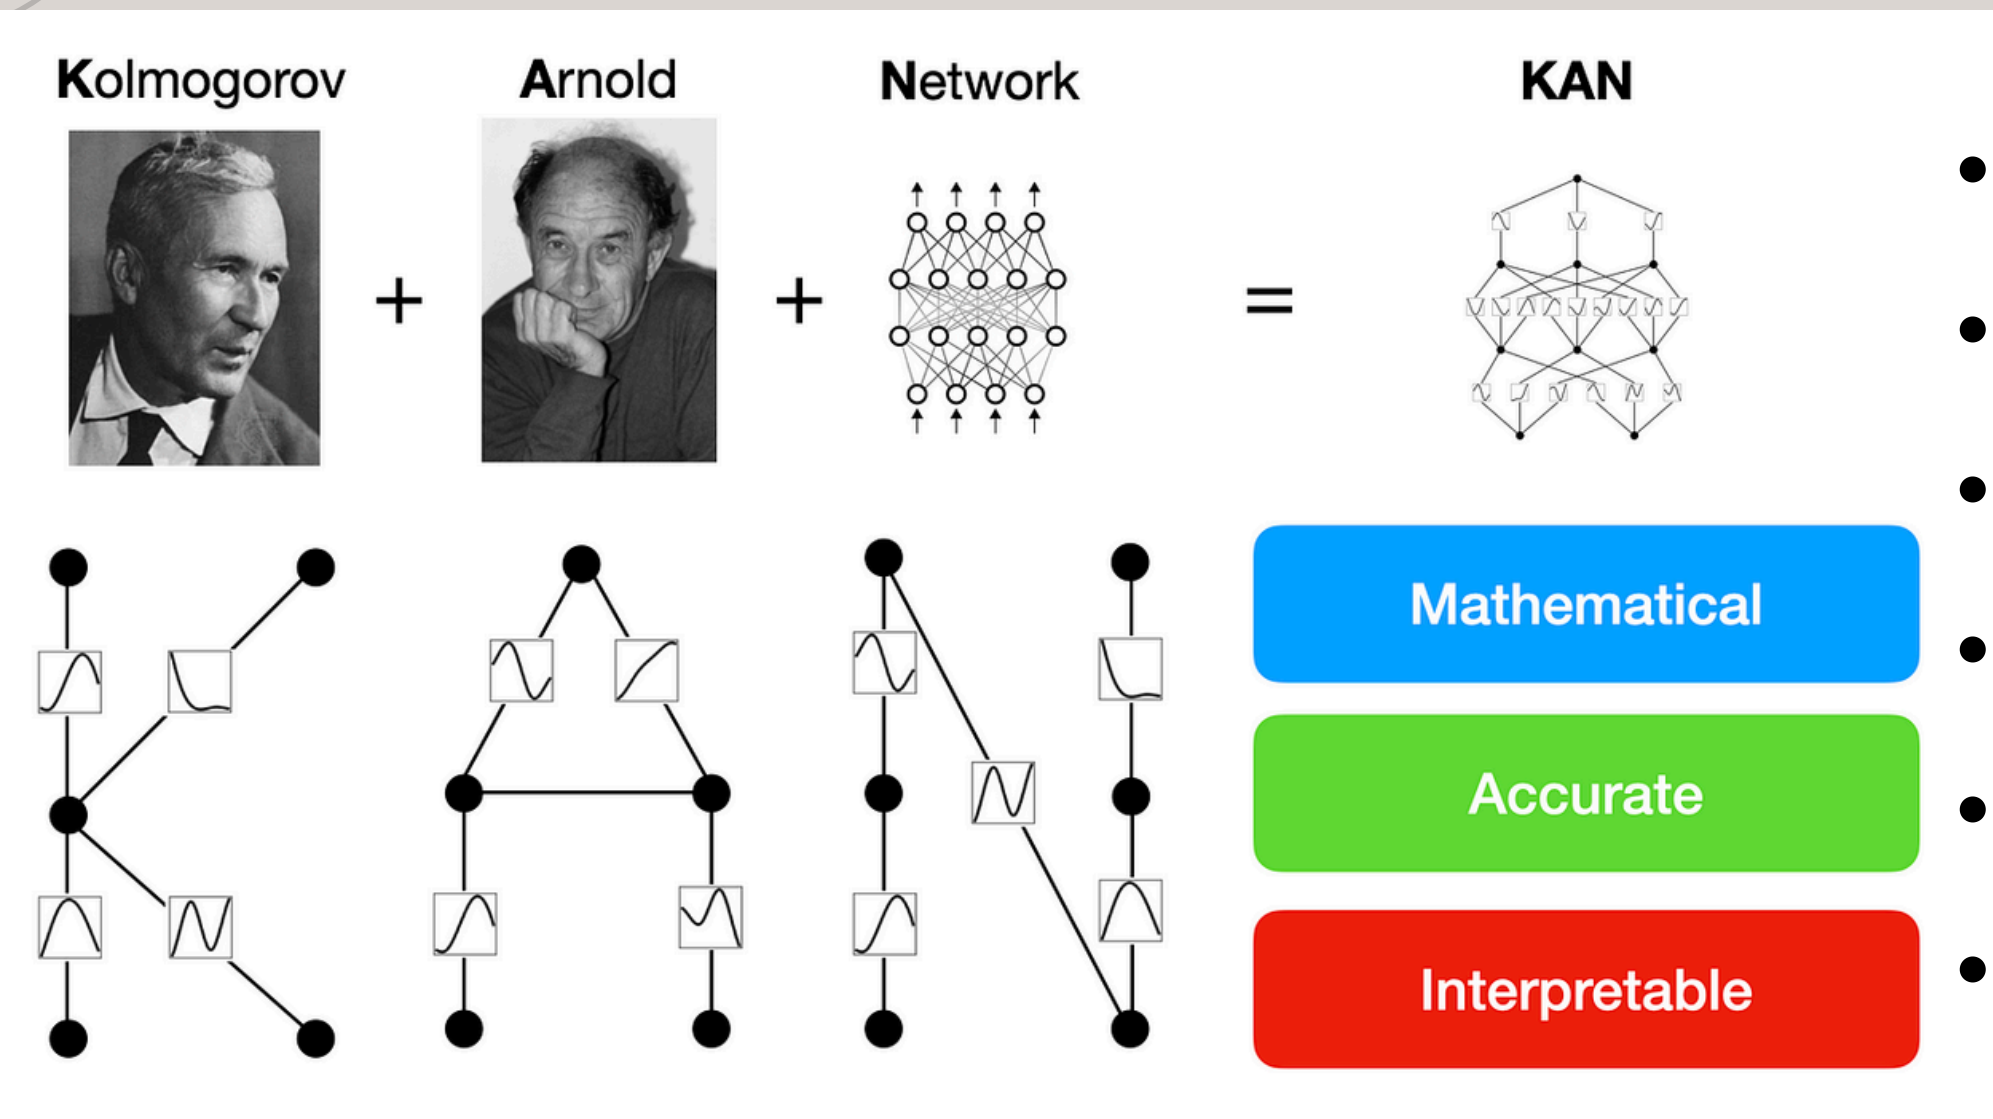
        for i in range(1, len(hidden_dims)):
            self.layers.append(nn.Linear(hidden_dims[i-1], hidden_dims[i]))
            self.batch_norms.append(nn.BatchNorm1d(hidden_dims[i]))

        self.output = nn.Linear(hidden_dims[-1], 2)

        self.dropout = nn.Dropout(dropout)
```

Tested several custom architectures

# Kolmogorov Arnold Networks



- Inspired by Kolmogorov–Arnold Theorem
- Function Approximation Framework
- Decomposes Multivariate Functions
- High Theoretical Expressivity
- Layered Structure of Simple Functions
- Efficient for Complex Tasks

# Model Evaluation



- Precision/Recall
- AUC-ROC
- Confussion matrix

# Model Evaluation

**Logistic Regression - C=10,  
num\_epochs=1000**

The Logistic Regression model  
achieved the following  
performance metrics:

**AUC-ROC: 0.95**

TABLE I  
LOGISTIC REGRESSION RESULTS

Class	P	Recall	F1-Score	Support
0	0.9999	0.9462	0.9724	1,270,881
1	0.0225	0.9550	0.0439	1,643
Accuracy	0.9463			
Macro Avg	0.5112	0.9506	0.5081	1,272,524
Weighted Avg	0.9987	0.9463	0.9712	1,272,524

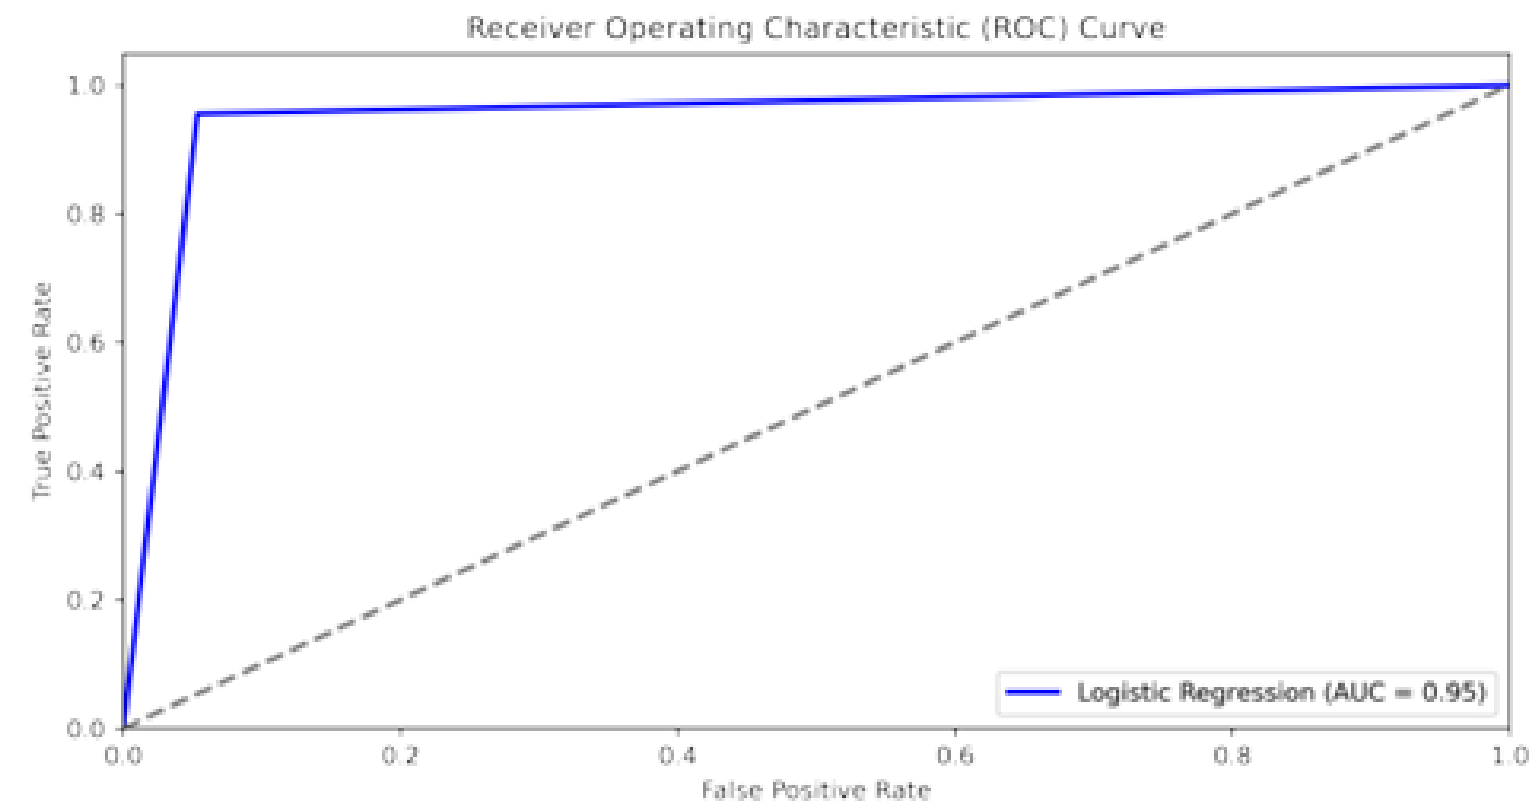


Fig. 7. ROC Curve for Baseline Logistic Regression Model

# Model Evaluation

TABLE III  
RANDOM FOREST RESULTS

Class	Precision	Recall	F1-Score	Support
0	0.9999	0.9835	0.9917	1,270,881
1	0.0723	0.9939	0.1348	1,643
Accuracy	0.9835			
Macro Avg	0.5361	0.9887	0.5632	1,272,524
Weighted Avg	0.9988	0.9835	0.9906	1,272,524

Random Forest with  
n\_estimators=200,  
max\_depth=10, min\_split=2,  
min\_leaf=2, bootstrap=True

Random Forest classifier  
achieved the second-best results,  
with the following metrics:

- AUC: 0.989




Fig. 10. ROC Curve for Random Forest Model



# Model Evaluation

TABLE II  
GRADIENT BOOSTING RESULTS

Class	Precision	Recall	F1-Score	Support
0	0.9999	0.9827	0.9913	1,270,881
1	0.0693	0.9976	0.1295	1,643
Accuracy	0.9827			
Macro Avg	0.5346	0.9901	0.5604	1,272,524
Weighted Avg	0.9988	0.9827	0.9901	1,272,524

Gradient Boosting with  
`n_estimators=200`, `lr=0.01`,  
`max_depth=7`, `min_split=2`,  
`min_leaf=2`

Gradient Boosting provided the best results, achieving superior performance on both training and test sets:

- Train AUC: 0.9989
- Test AUC: 0.9983

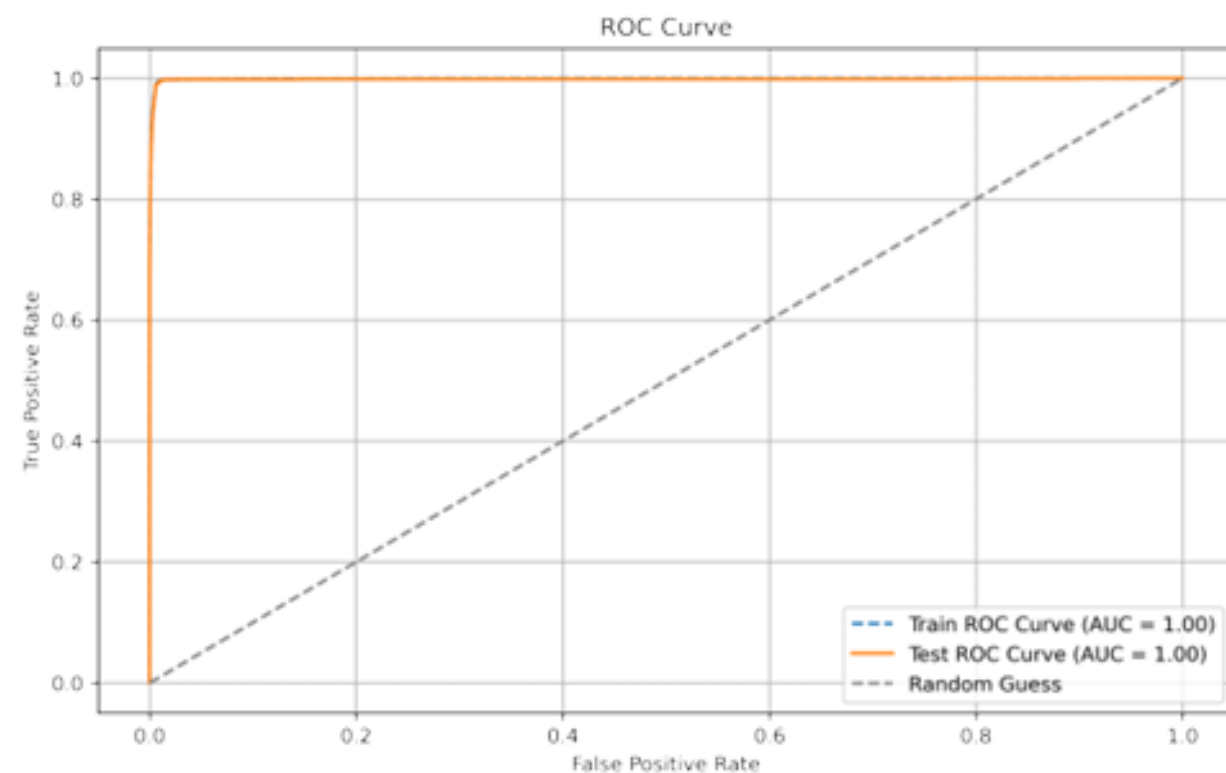


Fig. 8. ROC Curve for Gradient Boosting Model



# Model Evaluation

Using a Neural Feedforward Network within the Federated Learning framework, the model achieved the following performance metrics:

- AUC-ROC: 0.9979
- Balanced Accuracy: 0.9781
- Accuracy: 0.9801

$$\begin{bmatrix} 1220320 & 50963 \\ 4768 & 1265712 \end{bmatrix}$$

NEURAL FEEDFORWARD NETWORK REPORT

Class	Precision	Recall	F1-Score	Support
0	0.9961	0.9601	0.9778	1,270,283
1	0.9613	0.9963	0.9785	1,270,480
Accuracy	0.9801			
Macro Avg	0.9787	0.9782	0.9781	2,540,763
Weighted Avg	0.9787	0.9801	0.9781	2,540,763

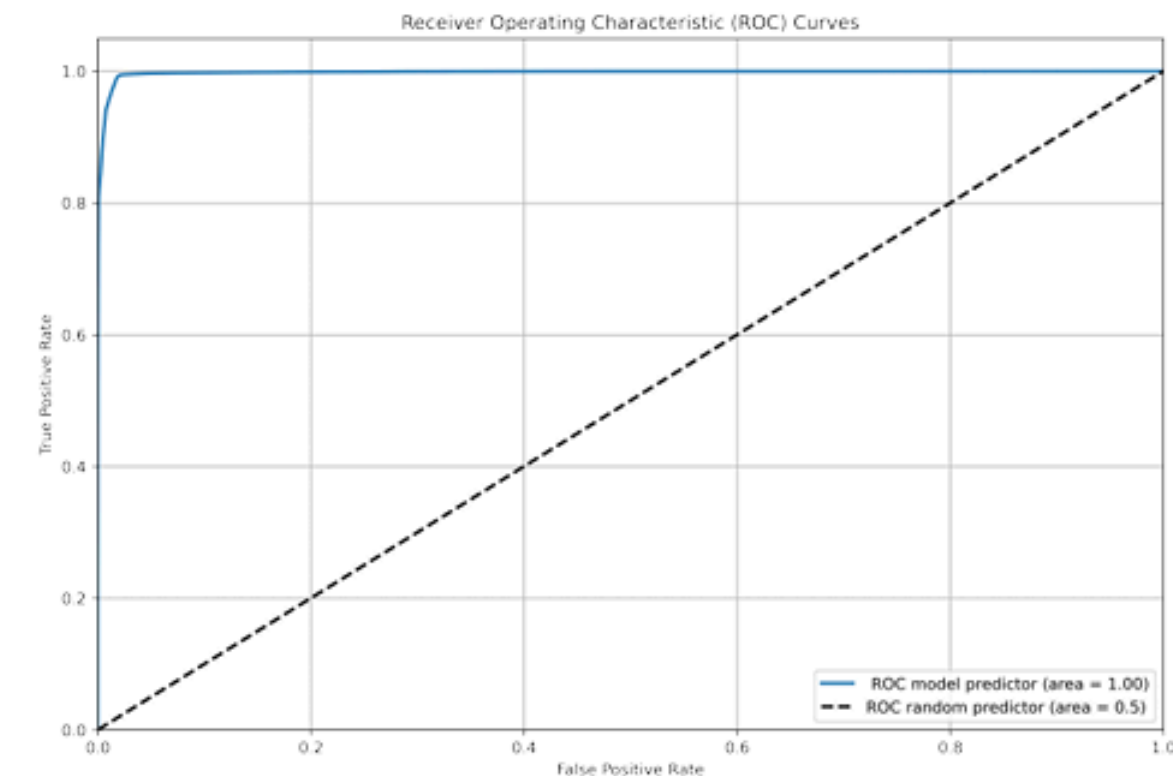


Fig. 11. ROC Curve for Neural Network with Federated Learning

# Model Evaluation

Using Kolmogorov–Arnold Networks (KANs) within the Federated Learning framework, the model achieved the following performance metrics:

- AUC-ROC: 0.9992
- Precision: 0.1233
  - Recall: 0.9970
- Accuracy: 0.9908

1259231	11650
5	1638

TABLE V  
FEDERATED LEARNING WITH KANS CLASSIFICATION REPORT

Class	Precision	Recall	F1-Score	Support
0	0.9999	0.9908	0.9954	1,270,881
1	0.1233	0.9970	0.2194	1,643
Accuracy	0.9908			
Macro Avg	0.5616	0.9939	0.6074	1,272,524
Weighted Avg	0.9989	0.9908	0.9944	1,272,524




Fig. 14. Training Progress for Kolmogorov–Arnold Networks with Federated Learning

# Model Evaluation

```
formula1, formula2 = model.symbolic_formula()[0]  
formula1
```

```
1.61 (− sin (0.55x3 + 9.34) + 0.12 tan (0.57x1 − 9.4) + 0.25 tanh (1.81x2 − 4.89) − 0.28 tanh (2.29x4 − 4.25) + 0.43)2 − 153.19e0.12 sin (2.56x1+3.0)+0.21 tanh (10.0x4−8.8)−0.16|1.59x3−6.98| −  
0.01 sin (4.91 sin (1.43x1 + 1.21) + 0.01 tan (0.75x4 − 3.48) − 1.43 tanh (2.67x2 − 10.0) + 2.27 tanh (2.18x3 − 9.12) + 9.85) +  
63.66 tanh (248.65 (1 − 0.16x1)4 + 0.36 sin (2.83x4 + 1.75) − 1.56 tanh (1.13x3 − 3.4) − 0.23 |3.23x2 − 9.88| + 1.2) + 14.34
```

```
1.61 (− sin (0.55x3 + 9.34) + 0.12 tan (0.57x1 − 9.4) + 0.25 tanh (1.81x2 − 4.89) − 0.28 tanh (2.29x4 − 4.25) + 0.43)2 −  
153.19e0.12 sin (2.56x1+3.0)+0.21 tanh (10.0x4−8.8)−0.16|1.59x3−6.98| −  
0.01 sin (4.91 sin (1.43x1 + 1.21) + 0.01 tan (0.75x4 − 3.48) − 1.43 tanh (2.67x2 − 10.0) + 2.27 tanh (2.18x3 − 9.12) + 9.85) +  
63.66 tanh (248.65 (1 − 0.16x1)4 + 0.36 sin (2.83x4 + 1.75) − 1.56 tanh (1.13x3 − 3.4) − 0.23 |3.23x2 − 9.88| + 1.2) + 14.34
```

# Discussions

TABLE VII  
PERFORMANCE OF PAYSIM MODEL

Metric	Precision (Fraud)	Recall (Fraud)	F1-Score (Fraud)	AUC- ROC
PaySim Model	1.0000	0.0019	0.0039	0.5010

$$\begin{bmatrix} 6,354,407 & 0 \\ 8,197 & 16 \end{bmatrix}$$

Our Solution:  
Significantly outperforms PaySim  
across all key metrics.

- Compare to Paysim Model
- Compare to Kaggle one of the best solutions

Paysim dataset has isFraud and  
isFlaggedFraud

PaySim Model Analysis:

- Perfect precision (1.0000), but extremely low recall (0.0019).
- F1-score is very low (0.0039), indicating poor balance.
- AUC-ROC (0.5010) is close to random guessing.

# Discussions

<https://www.kaggle.com/code/waleedfaheem/credit-card-fraud-detection-auc-0-9#-Model-Building->

- Compare to Paysim Model
- Compare to Kaggle one of the best solutions

TABLE IX

COMPARISON OF OUR SOLUTION AND KAGGLE SOLUTION

Model	Pr	Rec	F1-Score	AUC-ROC
Kaggle Solution	0.9700	0.7700	0.8600	0.9750
Neural Feedforward Network (FL)	0.9613	0.9963	0.9785	0.9979
KANs (FL)	0.1233	0.9970	0.2194	0.9992

**Our Solution:**  
Kaggle model is strong, but our Neural Network and KANs deliver superior performance across key metrics, ensuring robust fraud detection.


Kaggle solution, authored by Waleed Faheem, analysis:

- AUC-ROC: Neural Network (0.9979) and KANs (0.9992) outperform Kaggle solution (0.975).
- Recall: Neural Network (0.9963) and KANs (0.9970) excel vs. Kaggle (0.7700), identifying nearly all fraudulent cases.
- Precision: Kaggle (0.9700) slightly higher than Neural Network (0.9613), but recall and F1 improvements outweigh this.



# Conclusion

- Developed and evaluated various models for fraud detection using the PaySim dataset.
- Federated Learning-based Models
- Achieved superior performance compared to baseline models and existing solutions.
- Implemented a user-friendly Web Platform (Model Zoo)
  - KANs + Federated learning - new research approach not explored yet
  - Allows users to select and utilize different trained models



# Demo

<https://risk-radar.online>





# Thank You

Presented by Suryansh & Ulugbek