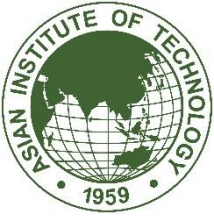


# Convolutional Neural Network

Dr. Mongkol Ekpanyapong

# CNN Example



```
from tensorflow.keras.datasets import mnist
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
test_images.shape
len(test_labels)
from tensorflow import keras
from tensorflow.keras import layers
img_rows = train_images[0].shape[0]
img_cols = train_images[0].shape[1]
model = keras.Sequential([
    layers.Conv2D(64, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(10, activation='softmax')
])

model.summary()
```



A decorative image in the top-left corner consisting of a blue square above a colorful, abstract pattern.


```
model.compile(optimizer="sgd",  
loss="sparse_categorical_crossentropy",  
metrics=["accuracy"])  
train_images = train_images.reshape(train_images.shape[0], img_rows,  
img_cols, 1)  
test_images = test_images.reshape(test_images.shape[0], img_rows,  
img_cols, 1)  
  
train_images = train_images.astype("float32") / 255.0  
  
test_images = test_images.astype("float32") / 255.0  
  
print(train_images.shape)  
  
model.fit(train_images, train_labels, epochs=5, batch_size=128)
```



# Homework


- Modify CNN model in Keras on MNIST data set so that the accuracy is better than 98%

# Layers of a CNN





# Fully Connected Network (FCN)

- Fully connected layers that takes in a 256x256 images and maps to 10 output neural will have  $256 \times 256 \times 10 + 1 = 655,360 + 1$  parameters
- Hence, the FCN model is more complex
- FCN is tended to be more overfit



# FCN/CNN layer Comparison



A decorative image in the top-left corner consisting of a blue square above a colorful, abstract pattern.

# The Limit of Fully Connected


- Higher number of parameters
- We use information from far-away pixel during the prediction
- Non-translation invariant






# The Limit of Fully Connected

- Every input pixel is combined with every other to produce the output results in large number of parameters



# Not-Translation Invariant




# CNN

- A convolution is a weighted sum of the pixel values of the image as the window slides across the whole image
- The difference between a fully connected layer and a convolution layer is that the fully connected layers learn global patterns whereas convolution layers learn local pattern
- Convolution operates over 3D tensor called feature map




# Key Characteristic of CNN

- The patterns they learn are translation invariant (FCN is not)
- They can learn spatial hierarchies of patterns
- Also, it should learn about the filter automatically



# CNN

- Locality and translation invariance



# Convolutional Neural Networks

- CNN has just enough weights to look at a small patch of the image
- The number of parameters are reduced to just  $5 \times 5 + 1 = 25 + 1$  parameters per node




Image from Machine Learning with TensorFlow book

A decorative image in the top-left corner consisting of a blue square above a smaller image of a circuit board.

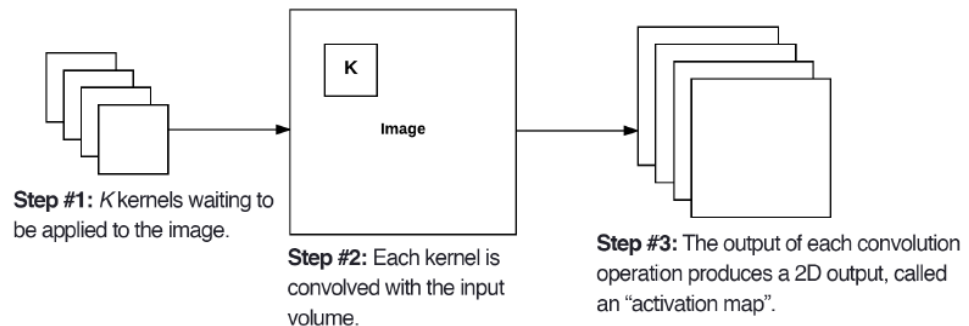
# CNN Layer Types

- Convolutional (CONV)
- Activation (ACT) e.g., RELU or SOFTMAX
- Pooling (POOL)
- Fully-connected (FC)
- Batch Normalization (BN)
- DropOut (DO)



# CONV layers


- Instead of fully connected layer, we can use convolutional layer instead
- Convolutional layer introduces the local connectivity and reduces the number of parameters for training






# Example

- Given a grayscale image of size  $256 \times 256$ , It connects to 25 output neurons, the number of parameters is  $256 \times 256 \times 25 + \text{bias} = 1,638,400 + 1$  for a fully connected layer



# Example with CONV

- Given a grayscale image of size  $256 \times 256$ ,  
It connects to  $5 \times 5$  patch, the number of  
parameters is  
 $5 \times 5 + 1 = 25 + 1$  for a  
convolutional layer




# Backpropagation

- We can consider the weight/parameter of each kernel similar to fully connected layer as shown below in which sigma is the activation function

$$\sigma \left( b + \sum_{l=0}^4 \sum_{m=0}^4 w_{l,m} a_{j+l,k+m} \right)$$


- Hence, the same back propagation can be applied in which now the convolution layer will learn the filters to be used

# Multiple Feature map



# Local Patterns in Image

- Image can be broken into local pattern such as edges, textures




A decorative image in the top-left corner consisting of a blue square above a grid of smaller squares in various colors.

# CNN Hyper-Parameters

- Number of convolution layers
- Convolution window size
- Convolution filter mask (Filter Depth= $K$ )
- Number of stride
- Padding






# Example of random initialized matrices for 32 filter

- The filters to be learnt by CNN



# K filter/Activation Map

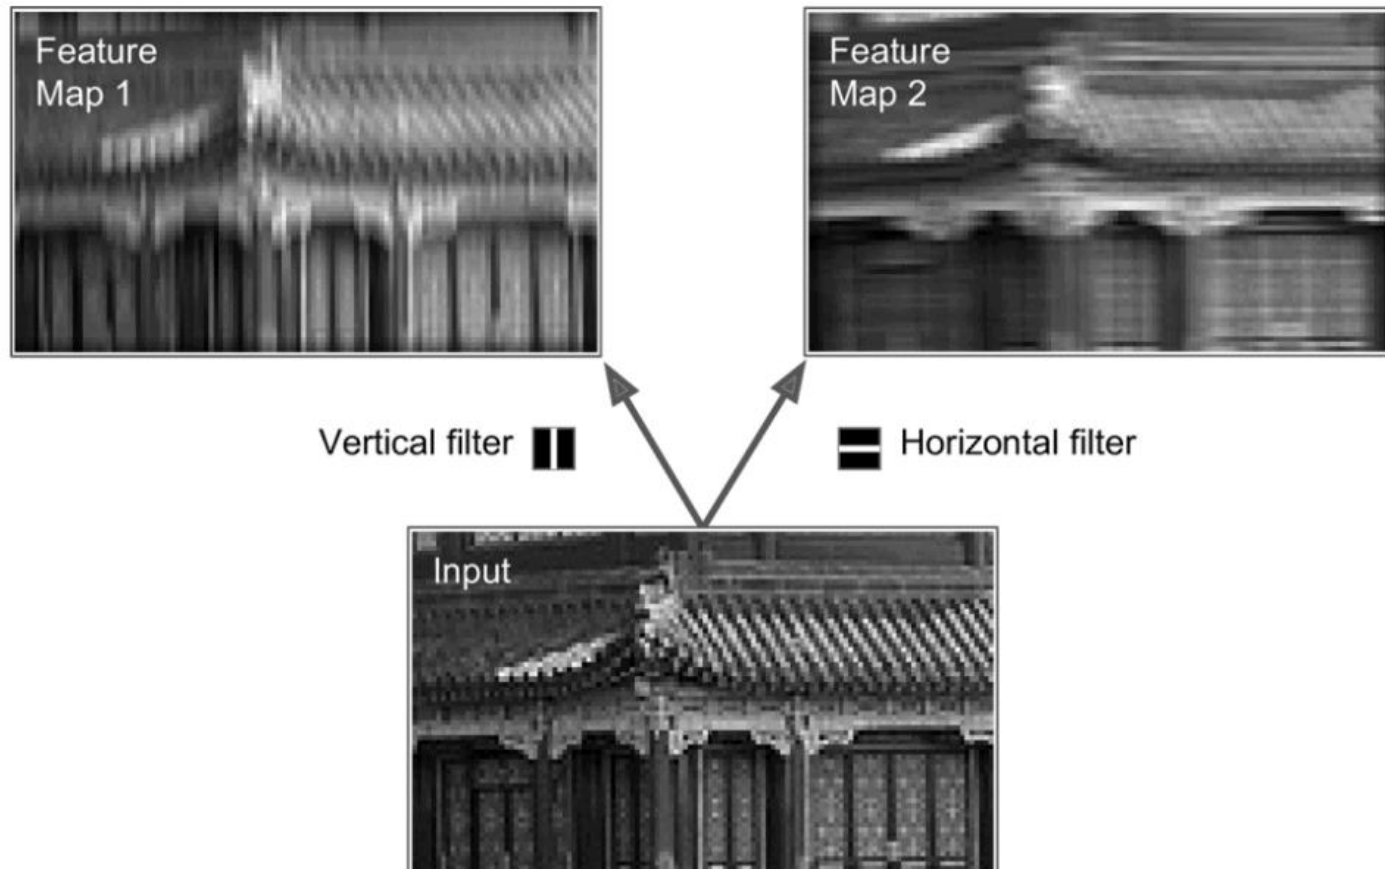
- After we apply K filter, we get the the volume of activation/feature map for the next layer
- Note that the depth can be more than K, as the input can have many channels






# Filter Results

- We call it Feature Map or Activation Map





# Filter Result by applying on a car





# Stride

- The stride is a step of sliding of small matrix over the image (big matrix)
- Example of image (left) and filter (right)

131	162	232	84	91	207
104	<del>-1</del>	<del>109</del>	<del>+1</del>	237	109
243	<del>-2</del>	<del>202</del>	<del>+2</del>	<del>135</del>	<del>126</del>
185	<del>-15</del>	<del>200</del>	<del>+1</del>	61	225
157	124	25	14	102	108
5	155	116	218	232	249

95	242	186	152	39
39	14	220	153	180
5	247	212	54	46
46	77	133	110	74
156	35	74	93	116

0	1	0
1	-4	1
0	1	0

- Output of a convolution with 1x1 stride (left) and 2x2 stride (right)

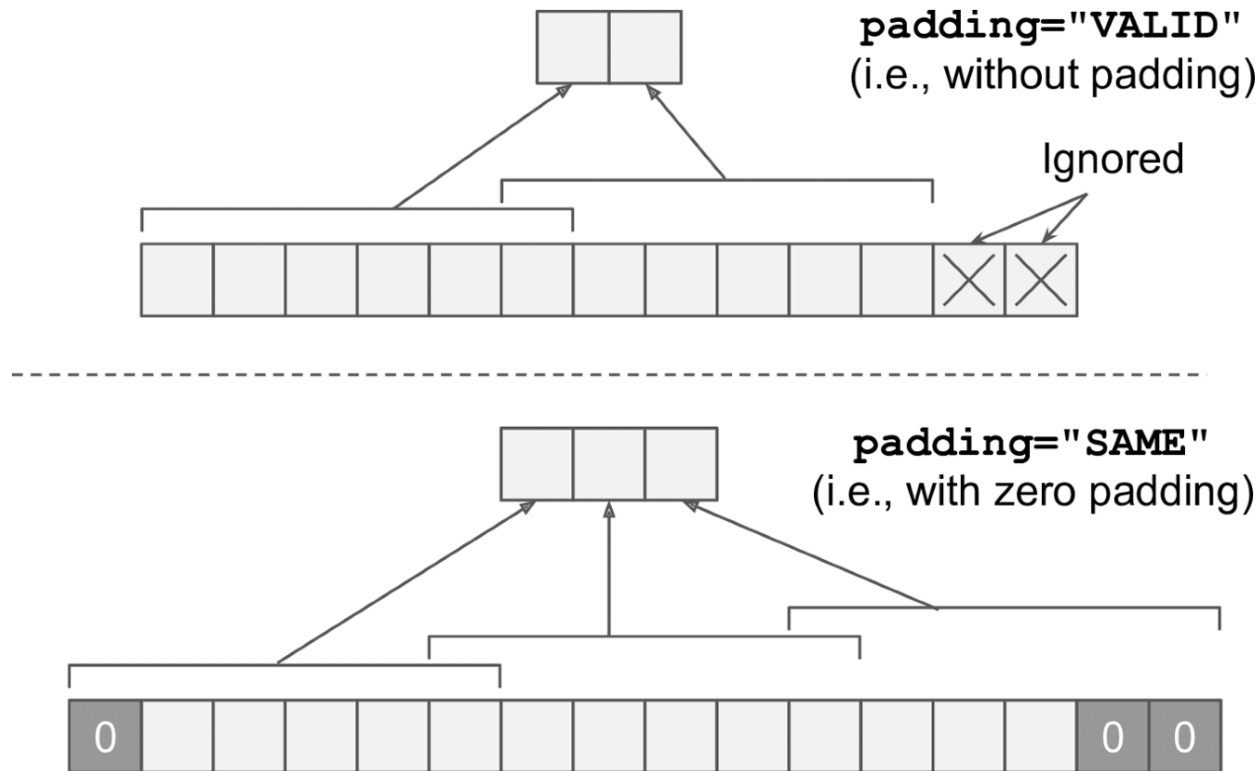
692	-315	-6
-680	-194	305
153	-59	-86

692	-6
153	-86





# Padding Example



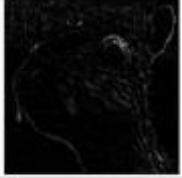






# Convolution Recap

- Convolution Filters or Kernels detect features in images



# Detected features

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	



# What is the Output?

1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel





# Recap

$$(1 \times 0) + (0 \times 1) + (1 \times 0) + (1 \times 1) + (0 \times 0) + (0 \times -1) + (0 \times 0) + (1 \times 1) + (1 \times 0) = 2$$

1x0	0x1	1x0	0	1
1x1	0x0	0x-1	1	1
0x0	1x1	1x0	0	0
1	0	0	1	0
0	0	1	1	0

\*

0	1	0
1	0	-1
0	1	0

=

2		

Input Image

Filter or Kernel

Output or Feature Map



$$(0 \times 0) + (1 \times 1) + (0 \times 0) + (0 \times 1) + (0 \times 0) + (1 \times -1) + (1 \times 0) + (1 \times 1) + (0 \times 0) = 1$$

1	0x0	1x1	0x0	1
1	0x1	0x0	1x-1	1
0	1x0	1x1	0x0	0
1	0	0	1	0
0	0	1	1	0

Input Image

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel

=

2	1	

Output or Feature Map





1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel

=

2	1	-1

Output or Feature Map





1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel

=

2	1	-1
-1		

Output or Feature Map





1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel

=

2	1	-1
-1	1	

Output or Feature Map





1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel

=

2	1	-1
-1	1	3

Output or Feature Map





1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel

=

2	1	-1
-1	1	3
2		

Output or Feature Map





1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel

=

2	1	-1
-1	1	3
2	1	

Output or Feature Map







1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel

=

2	1	-1
-1	1	3
2	1	1

Output or Feature Map

Input 5x5 with filter 3x3, then we get 3x3 output



# Calculating Feature Map Size

Feature Map Size =  $n - f + 1 = m$

Feature Map Size =  $5 - 3 + 1 = 3$

1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

$5 \times 5$   
 $n \times n$

\*

0	1	0
1	0	-1
0	1	0

$3 \times 3$   
 $f \times f$

=

2	1	-1
-1	1	3
2	1	1

$3 \times 3$   
 $m \times m$



# Stride

- It can use to control feature map output
- Larger stride has less overlap

# Stride of 2

1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image

\*


0	1	0
1	0	-1
0	1	0

Filter or Kernel

=

2	

Output or Feature Map



# Shift by 2:

$$\begin{matrix} \xrightarrow{\hspace{2cm}} \\ \begin{array}{|c|c|c|c|c|} \hline 1 & 0 & 1 & 0 & 1 \\ \hline 1 & 0 & 0 & 1 & 1 \\ \hline 0 & 1 & 1 & 0 & 0 \\ \hline 1 & 0 & 0 & 1 & 0 \\ \hline 0 & 0 & 1 & 1 & 0 \\ \hline \end{array} \end{matrix}$$

Input Image

\*

$$\begin{array}{|c|c|c|} \hline 0 & 1 & 0 \\ \hline 1 & 0 & -1 \\ \hline 0 & 1 & 0 \\ \hline \end{array}$$

Filter or Kernel

=

$$\begin{array}{|c|c|} \hline 2 & -1 \\ \hline & \\ \hline \end{array}$$

Output or Feature Map





$$\downarrow$$

1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel

=

2	-1
2	

Output or Feature Map





1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel

=

2	-1
2	1

Output or Feature Map



# Stride = 1

Padding = 0

1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image  
5 x 5

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel  
3 x 3

=

2	-1	1
2	1	0
1	-1	1

$$(n \times n) * (f \times f) = \left(\frac{n + 2p - f}{s} + 1\right) \times \left(\frac{n + 2p - f}{s} + 1\right) = \left(\frac{5 + (2 \times 0) - 3}{1} + 1\right) \times \left(\frac{5 + (2 \times 0) - 3}{1} + 1\right) = 3 \times 3$$





# Stride = 2

Padding = 0

1	0	1	0	1
1	0	0	1	1
0	1	1	0	0
1	0	0	1	0
0	0	1	1	0

Input Image  
5 x 5

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel  
3 x 3

=

2	-1
2	1

$$(n \times n) * (f \times f) = \left(\frac{n + 2p - f}{s} + 1\right) \times \left(\frac{n + 2p - f}{s} + 1\right) = \left(\frac{5 + (2 \times 0) - 3}{2} + 1\right) \times \left(\frac{5 + (2 \times 0) - 3}{2} + 1\right) = 2 \times 2$$



# Padding

Padding = 1

0	0	0	0	0	0	0
0	1	0	1	0	1	0
0	1	0	0	1	1	0
0	0	1	1	0	0	0
0	1	0	0	1	0	0
0	0	0	1	1	0	0
0	0	0	0	0	0	0

Input Image  
7 x 7

\*

0	1	0
1	0	-1
0	1	0

Filter or Kernel  
3 x 3

=










2	-1	1	1	0
2	1	0	1	2
1	-1	1	0	1
0	1	2	1	1
2	0	1	0	2

$$(n \times n) * (f \times f) = \left(\frac{n + 2p - f}{s} + 1\right) \times \left(\frac{n + 2p - f}{s} + 1\right) = \left(\frac{5 + (2 \times 1) - 3}{1} + 1\right) \times \left(\frac{5 + (2 \times 1) - 3}{1} + 1\right) = 5 \times 5$$





We have 3x3 filter and 7x7 image with a stride of 2. How many position which we apply the filter?



# Output Size

- To guarantee, the integer output size, the following equation can be used to check:

$$((W - F + 2P)/S) + 1$$

When  $W$  is the Image size (square)

$F$  is the kernel size (square)

$S$  is stride

$P$  is the padding



# Example

- Alexnet model

The image size is 227x227.

The kernel size is 11x11.

No padding.

Stride is 4.

We obtain number below which is integer

$$((227 - 11 + 2(0))/4) + 1 = 55$$



# Question?

If we have a  $21 \times 21$  filter and a  $251 \times 151$  images with stride of 10. What is the output size?

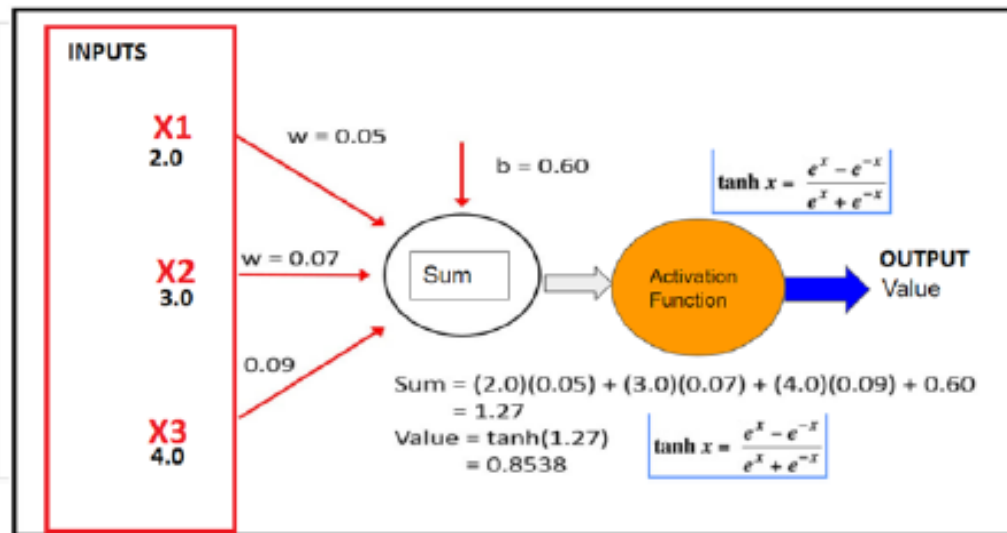


# Answer


$$24 \times 14 = [(251 - 21) / 10 + 1][(151 - 21) / 10 + 1]$$

# Activation Layer


- Activation function is used to introduce non-linearity in the system



- Example of activation function



**Tangent Hyperbolic:**  
Its output always lies between -1 and 1 no matter what inputs are




**Sigmoid--:**  
Its output always lies between 0 and 1 no matter what inputs are



# RELU Operation


- Change all negative values to zero
- Leave all positive value alone

$$f(x) = \max(0, x)$$



# Activation Function

- RELU (Rectifier Linear Unit) and Exponential Linear Unit (ELU)




- The Model:

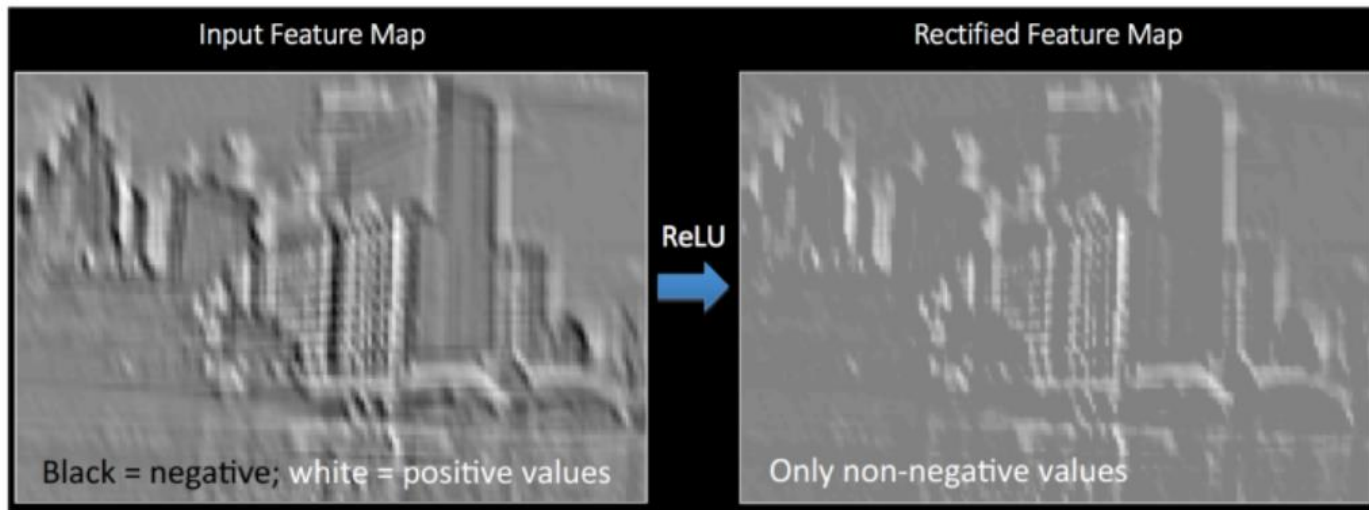
INPUT  $\Rightarrow$  CONV  $\Rightarrow$  RELU  $\Rightarrow$  FC

# Example

- RELU Activation Function




# Output Example



Source - [http://mlss.tuebingen.mpg.de/2015/slides/fergus/Fergus\\_1.pdf](http://mlss.tuebingen.mpg.de/2015/slides/fergus/Fergus_1.pdf)

# Result after Activation Function



# Reduce the Input Size

There are two ways to reduce the input size

- Convolution with stride  $> 1$
- Pooling layer



# Convolution on Grey Scale



1	1	0	0	0
1	1	0	0	0
1	1	0	0	0
1	1	0	0	0
1	1	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1


=

3	3	0
3	3	0
3	3	0






# Convolution on Color Image




Input Image

\*



Filter or Kernel

=



Output or Feature Map





# Output Size


$$(n \times n \times n_c) * (f \times f \times n_c) \Rightarrow (n-f+1) \times (n-f+1) \times n_f$$

$n$  = number of input pixel

$f$  = filter size


$n_c$  = number of channel

$n_f$  = number of filter




Input Image  
5 x 5 x 3

\*



2 Filters or Kernels  
3 x 3 x 3 x 2

=



Output or Feature Map  
3 x 3 x 2

# Pooling

- The pooling is a process of subsample (shrink the image) to reduce the computation
- There are many functions that can be applied such as max pooling, min pooling, mean pooling



# Example

4	123	1	34
56	99	222	253
45	122	165	12
21	187	133	124

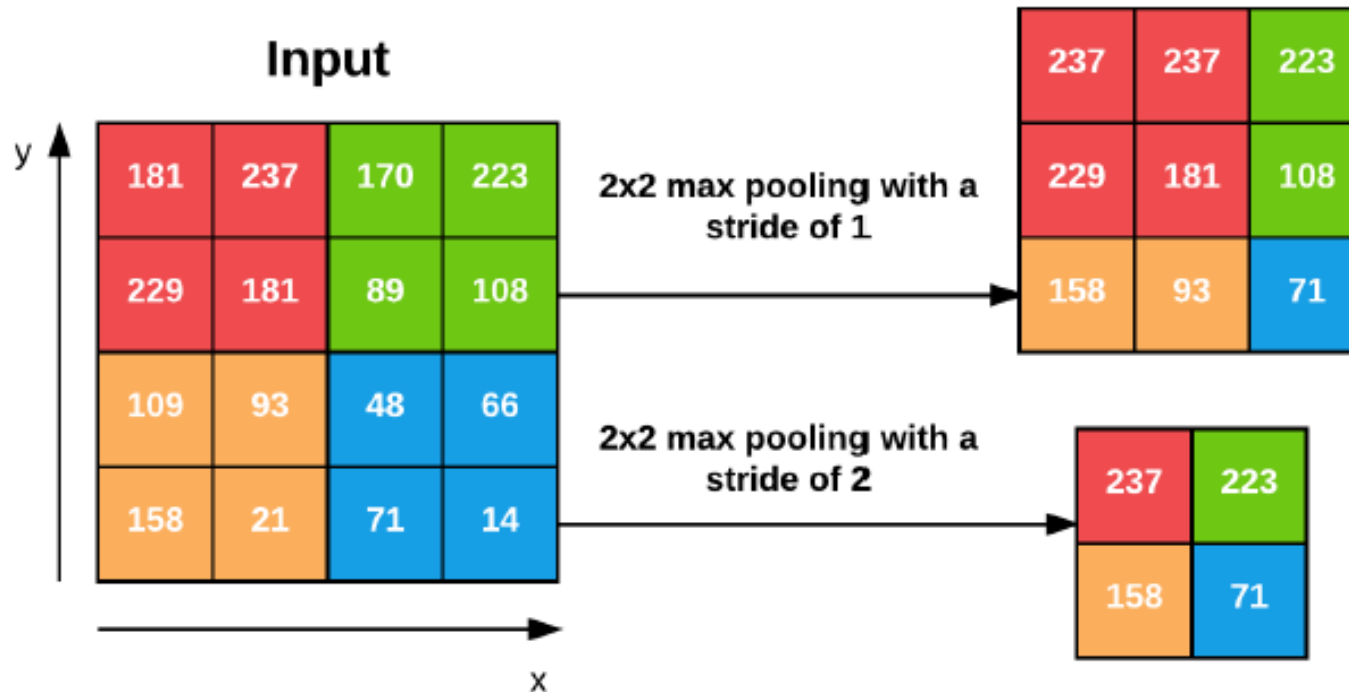
MaxPool Operation



Stride = 2  
Kernel = 2x2


123	167
187	165

# Example




# Pooling Example


- Example of pooling size using a 2x2 kernel with a stride of 2 and no padding



# Pooling Example



# Results after Pooling



# Model with Pooling

- Example of a model with Pooling:

INPUT  $\Rightarrow$  CONV  $\Rightarrow$  RELU  $\Rightarrow$  POOL  $\Rightarrow$  FC


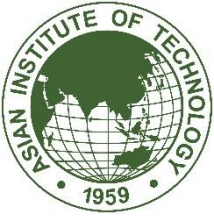


# More on Pooling


- Typically, we use  $2 \times 2$  kernel and stride of 2
- With that, pooling reduces the dimension by a factor of 2 on width and height
- Pooling makes our model more invariant to minor transformations and distortions



# How Max Pooling Achieves Translation Invariance



Shifting Our C down one pixel



# How Pooling Works?

- Neighboring pixels are strongly correlated
- Hence, we can reduce the size of the output by subsampling the filter response without losing information
- A big stride in the pooling layer leads to high information loss
- In practice a stride of 2 was found to be effective




# To POOL or CONV?

- In 2014 paper, striving for simplicity: The ALL Convolutional Net, Springenberg et al., propose to discard the POOL layer entirely and use CONV layers with a larger stride to handle downsampling
- Now, it becomes increasingly common trend to not use POOL


# Another way of Pooling

- We can achieve good translation invariant with non-explode weight by using Convolution Neural Network




# Another way of pooling


- Convolution is done on many filters automatically
- Pooling can also be used between filters



Four convolutional  
kernels predicting  
over the same 2

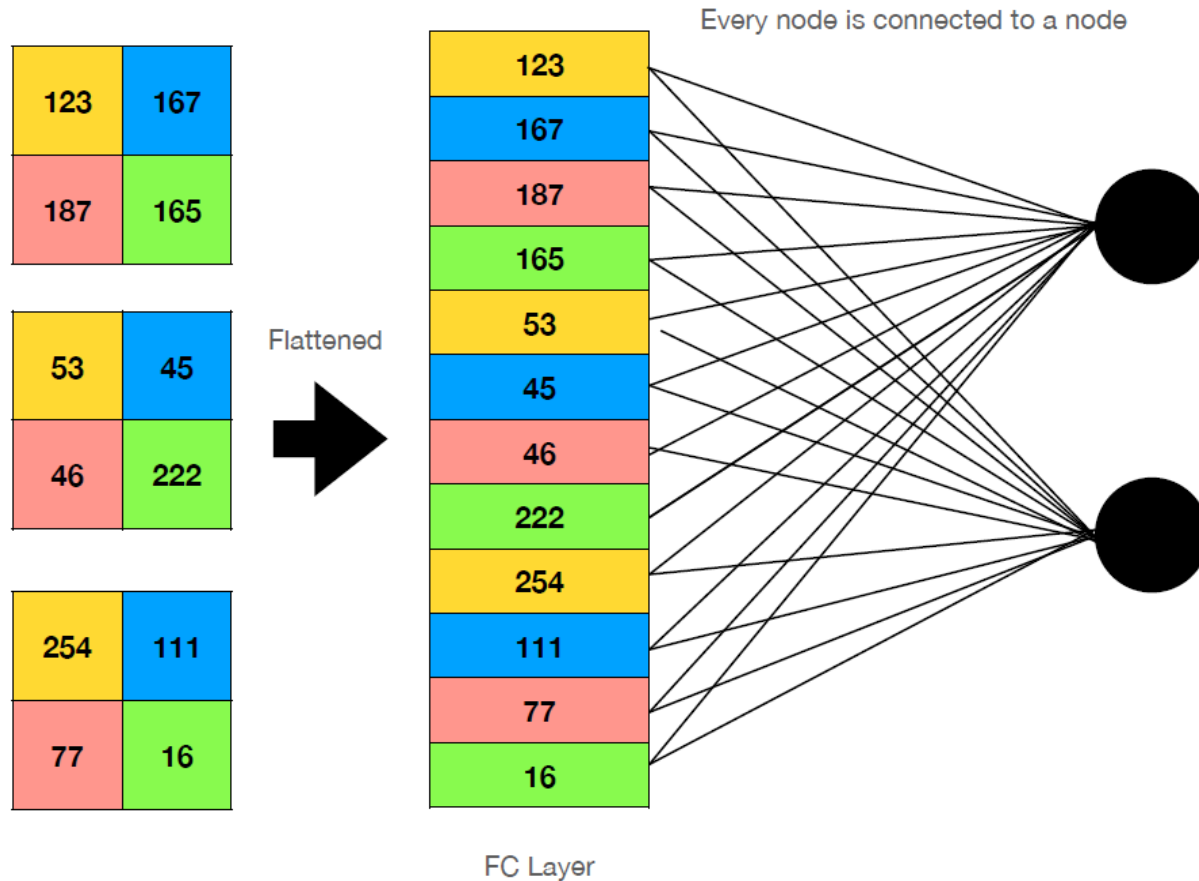


Outputs from each  
of the four kernels in  
each position



The max value of each  
kernel's output forms a  
meaningful representation and  
is passed to the next layer.

# Flattening and Fully Connected Layer



A decorative image in the top-left corner consisting of a blue square above a grid of smaller squares in various colors.

# Fully Connected Layer

- It is common to use one or two FC layers prior to applying the softmax classifier
- There is also a trend to not use FC layer as it is computing intensive
- Model:

INPUT  $\Rightarrow$  CONV  $\Rightarrow$  RELU  $\Rightarrow$  POOL  $\Rightarrow$  FC





# Batch Normalization (BN)

- Recall- the output of CNN is

Batch size x Feature Map Height x Feature Map Width x Channels

- BN calculates the mean and standard deviation of each input variable, to a layer per mini-batch and uses this to perform the standardization

# Batch Normalization

- It is introduced by Ioffe and Szegedy in 2015 paper, Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift to add BN layer
- The idea is to normalize the data where  $x_i$  is mini-batch
- The equation is as follow:

$$\hat{x}_i = \frac{x_i - \mu_\beta}{\sqrt{\sigma_\beta^2 + \epsilon}}$$

when

$$\mu_\beta = \frac{1}{M} \sum_{i=1}^m x_i$$


$$\sigma_\beta^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_\beta)^2$$

A decorative image in the top-left corner consisting of a blue square above a grid of smaller squares in various colors.

# Batch Normalization


- Advantage:
  - Help reduce the number of epochs for training and help for regularization
  - Recommend to put wherever we can
- Drawback:
  - Slow down the system
- Model:

INPUT => CONV => RELU => BN =>  
POOL => FC


A decorative image at the bottom-left corner consisting of three small squares: a red one, a yellow one, and a blue one with a grid pattern.

# Dropout


- The dropout is a regularization technique
- The idea is to keep turn off some neural so that we use many good neural nodes for prediction (not relying only on one)
- It provides multiple redundant nodes



# Dropout Example



(a) Standard Neural Net



(b) After applying dropout.

# Dropout

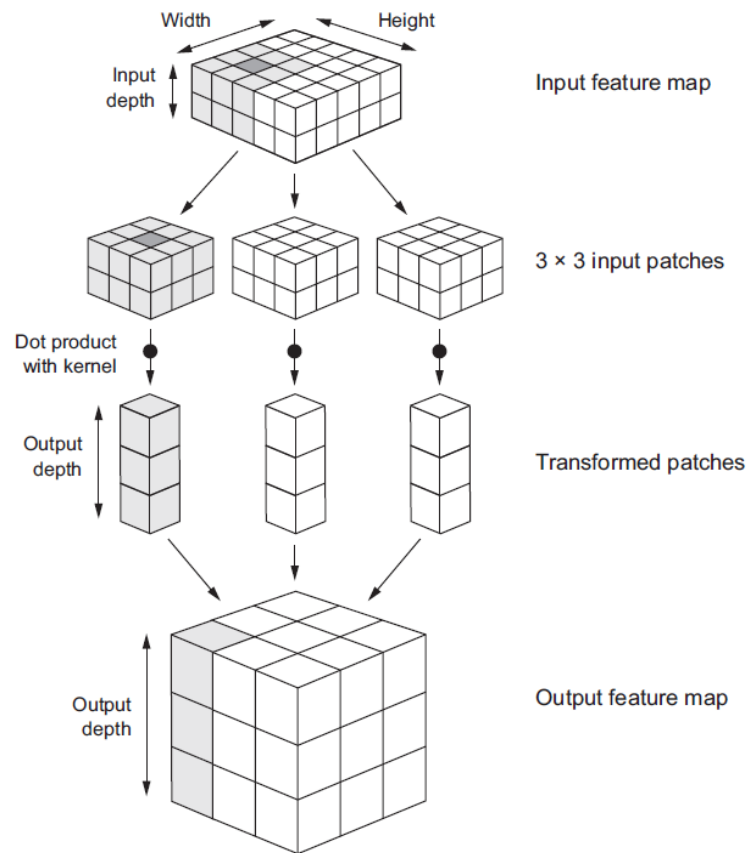
- Note that we randomly add dropout only during the training time, (testing time, we activate back all nodes)

- Model:


INPUT => CONV => RELU => BN =>  
POOL => FC => DO => FC => DO

# Convolution on an image

- How convolution work:




# Complete CNN Architecture







# A Simple CNN



# 4 Layer Deep CNN for MNIST




# Calculate Output Size



Layer	Depth	Width	Height	Filter (w)	Filter (h)
Input	1	28	28		
Conv_1	32	26	26	3	3
Conv_2	64	24	24	3	3
Max Pool	64	12	12	2	2
Flatten	9216	1	1		
Fully Connected	128	1	1		
Output	10	1	1		


# Conv\_1 layer

What is the output size after Conv\_1 layer?



# Conv\_1 layer

What is the output size after the layer?




$$[(n+2p-f)/s + 1] \times [(n+2p-f)/s + 1] = 26 \times 26$$

$$n = 28, f = 3, s = 1, p = 0$$



# Conv\_2 layer

What is the output size after the layer?




$$[(n+2p-f)/s + 1] \times [(n+2p-f)/s + 1] = 24 \times 24$$

$$n = 26, f = 3, s = 1, p = 0$$



# Max Pool layer

What is the output size after the layer?




$$[(n+2p-f)/s + 1] \times [(n+2p-f)/s + 1] = 12 \times 12$$

$$n = 24, f = 0, s = 2, p = 0$$



# Flatten layer

What is the output size after the layer?




$$12 \times 12 \times 64 = 9,216$$



# FC1 layer


What is the output size after the layer?



128


# FC2 layer

What is the output size after the layer?



10





# Calculate the number of parameters







# How many parameters are in 16 Conv. Filters with 3x3 kernels

$$((\text{Height} \times \text{Width} \times \text{Dept}) + \text{bias}) \times \#\text{Filters} =$$

$$((3 \times 3 \times 3) + 1) \times 16 = 448$$




# How many Parameter in this system?



# Conv\_1 layer

How many parameter?




$$((\text{Height} \times \text{Width} \times \text{Depth}) + \text{bias}) \times \# \text{Filters} =$$

$$((3 \times 3 \times 1 + 1) \times 16 = 320$$



# Conv\_2 layer

How many parameter?




$$((\text{Height} \times \text{Width} \times \text{Depth}) + \text{bias}) \times \# \text{Filters} =$$

$$((3 \times 3 \times 32 + 1) \times 64 = 18,496$$



# Max Pool layer

How many parameter?




No trainable parameters





# Flatten layer

How many parameter?




No trainable parameters (9,216 nodes)



# FC1 layer

What is the output size after the layer?




$$(\text{Input node} + \text{bias}) \times \text{Output node} =$$

$$(9,216 + 1) \times 128 = 1,179,776$$

# FC2 layer

What is the output size after the layer?



$$(\text{Input node} + \text{bias}) \times \text{Output node} =$$

$$(128 + 1) \times 10 = 1,290$$



# Total Parameters

Layer	Parameters
Conv_1 + ReLU	320
Conv_2 + ReLU	18494
Max Pool	0
Flatten	0
FC_1	1,179,776
FC_2 (Output)	1,290
Total	1,199,882

A decorative graphic in the top-left corner consisting of a blue square above a grid of smaller squares in various colors.

# Four Important Feature for Deep Learning

- Dataset
- A Loss Function
- A Neural Network Architecture
- An optimization method



# Rules of Thumb

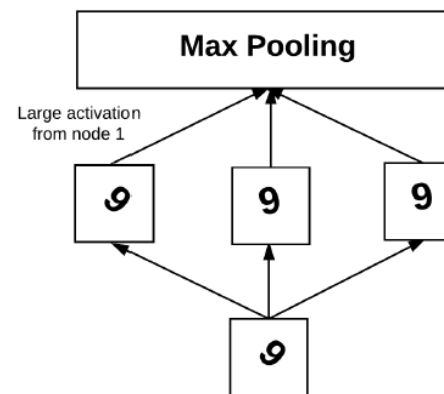

- Common input sizes include  $32 \times 32$ ,  $64 \times 64$ ,  $96 \times 96$ ,  $224 \times 224$ ,  $227 \times 227$  and  $229 \times 229$
- The input layer should be divisible by two multiple times (to use POOL)
- CONV layers should be small size such as  $3 \times 3$ ,  $5 \times 5$ , or  $1 \times 1$
- Large filter can be used in very first CONV such as  $7 \times 7$  and  $11 \times 11$





# Is CNN Translation, rotation, and scaling invariant

- CNN is translation invariant with the help of convolutional layer
- It is not rotation and scaling invariant unless you let the network learn a lot of rotation and scaling samples



A decorative graphic in the top left corner consisting of a blue square above a square image of a circuit board.

# Why CNNs Work So Well For Images?


- ANN requires many parameters
- ANN is not translation-invariant





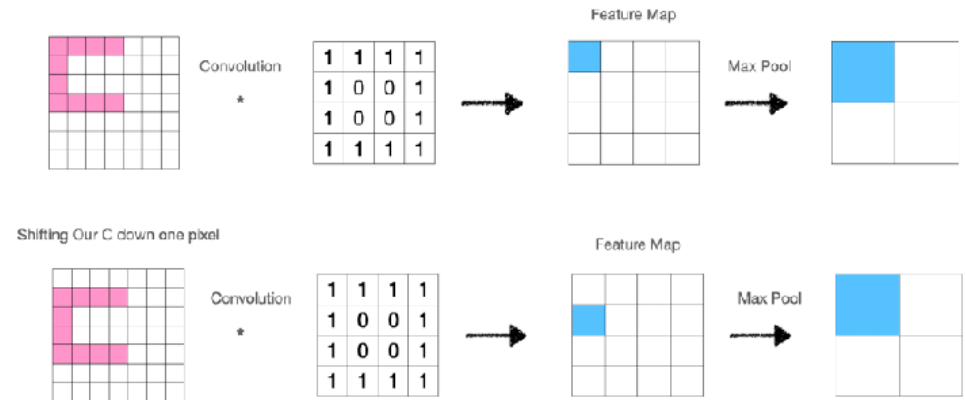
# 3 Layers of ANN on MNIST

- Input  $28 \times 28 = 784$
- For second layer, if we have same number of node as CNN Feature map second layer:  
Feature map second layer:  
 $32 \times 26 \times 26 = 21,631$  nodes
- #trainable parameter =  
 $784 \times 21,631 = 16,958,704$



# Advantages of CNN

- Parameter sharing



- Less number of weight comparing with CNN
- Transition Invariance



A decorative image in the top-left corner consisting of a blue square above a grid of smaller squares in various colors.


# CNN Assumptions

- Low-level features are local
- Features are translational invariant
- High-level features are made up of low-level features



# Rules of Thumb

- We commonly use a stride of  $S=1$ , unless we want to do use CONV instead of POOL
- Zero padding should always be applied
- At a novice, POOL is easier to use. Once, you get enough experience, try to avoid it
- POOL should be used with max pooling with  $2 \times 2$  size and stride  $=2$
- BN and DO should be applied if possible



# Questions?

