

st125457_ulugbek_assignment3

October 6, 2024

1 Machine Learning Assignment 3 - Predicting Car Prices!

This Jupyter notebook is a template for solving the assignment problem, i.e., Chaky company makes some car but he has difficulty setting the price for the car. . Here, I will try to apply the skills I've learned over the past lectures. This notebook contains the following structure: - **1. Setup:** Import block with all necessary imports (also provide some blocks with connection to drive, kaggle, and etc. for future use) - **2. Loading the Data:** Loading, EDA, data cleaning, feature selection, and preprocess the dataset. - **3. Models:** Starter code for basic models to kickstart your experimentation. - **4. Evaluation Metrics:** Tools to evaluate your models using various metrics. - **5. Inference and Conclusion:** Testing the best model and generating Report.

Let's start!

1.1 Some notes:

The typical workflow of data science project is following:

1. Problem Definition

- Objective: Clearly define the problem you're trying to solve. Understand the business or research context.
Tasks:

- Identify the key objectives and success metrics.
- Understand the constraints and resources available.
- Formulate hypotheses or research questions.

2. Data Collection

- Objective: Gather the necessary data from various sources, which could be internal databases or external data providers.
Tasks:

- Identify data sources and acquire the data.
- Integrate data from multiple sources if needed.
- Ensure data privacy and compliance with regulations (e.g., GDPR).

3. Data Exploration and Analysis (Exploratory Data Analysis - EDA)

- Objective: Understand the data, its patterns, and any potential issues through visualization and statistical analysis.
Tasks:

- Summarize the data using descriptive statistics.
- Visualize distributions, correlations, and trends.
- Identify patterns, outliers, and potential relationships between features.
- Formulate additional hypotheses based on the data.

4. Data Preprocessing

- Objective: Clean and prepare the data for modeling.
Tasks:

- - Handle missing values (imputation or removal).
 - - Handle outliers.
 - - Encode categorical variables.
 - - Normalize or standardize numerical features.
 - - Split the data into training, validation, and test sets.
5. Feature Engineering
- Objective: Create new features or modify existing ones to improve model performance.
- Tasks:
- - Create new features from existing data (e.g., interaction terms, polynomial features).
 - - Apply feature scaling (normalization or standardization).
 - - Transform features to handle skewness (e.g., log transformations).
 - - Reduce dimensionality if necessary (e.g., PCA).
6. Model Selection
- Objective: Choose the appropriate machine learning models for the problem.
- Tasks:
- - Compare different algorithms (e.g., linear models, decision trees, ensemble methods, neural networks).
 - - Consider baseline models for comparison.
 - - Choose models based on the problem type (e.g., classification, regression).
7. Model Training
- Objective: Train the chosen models on the preprocessed data.
- Tasks:
- - Train the models using the training dataset.
 - - Perform hyperparameter tuning (e.g., using grid search or random search).
 - - Use cross-validation to evaluate model performance.
8. Model Evaluation
- Objective: Assess the model's performance using relevant metrics and ensure it meets the project requirements.
- Tasks:
- - Evaluate model performance on the validation dataset.
 - - Use appropriate metrics (e.g., accuracy, precision, recall, F1-score, RMSE).
 - - Analyze model errors and refine the model if necessary.
9. Model Deployment
- Objective: Integrate the model into a production environment where it can be used to make predictions.
- Tasks:
- - Deploy the model as a service (e.g., REST API, microservice).
 - - Ensure scalability and monitor the model's performance in production.
 - - Handle model retraining as needed (e.g., with new data).
10. Monitoring and Maintenance
- Objective: Continuously monitor the model's performance and maintain its accuracy over time.
- Tasks:
- - Track model performance using key metrics.
 - - Monitor for data drift and update the model if necessary.
 - - Address any issues in production and ensure model reliability.
11. Documentation and Reporting
- Objective: Document the entire process and communicate the results to stakeholders.
- Tasks:
- - Prepare detailed reports and visualizations.
 - - Document the data, model, and processes.
 - - Share insights and actionable recommendations with stakeholders.

12. Iteration and Optimization

- Objective: Refine the project by iterating over the steps to improve results.

Tasks:

- - Revisit earlier steps based on feedback and new insights.
- - Optimize the model and the workflow for better performance.

2 1. Setup

The following libraries are required to run this notebook. If you are running this on Colab it should be all smooth sailing. If you are running it locally please make sure you have all of these installed.

```
[ ]: # Import section, basically importing everything what I need later + default
      ↪ imports
import os
import random

%matplotlib inline
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

import mlflow

# For sklearn imports I will import them in model sections for better
      ↪ explanation purposes
from sklearn.metrics import classification_report

from sklearn.metrics import mean_squared_error, r2_score
```

2.1 2.1 DataLoading

First thing we need to do is load in the data. We will be looking at the cars dataset (shared for this assignment [cars](#)). This dataset is tabular and contains information regarding car details(year, brand, mileage, and etc.) and we need to predict the price of the car(regression).

```
[ ]: # Loading the data
train_csv_path = '../dataset/cars.csv'
df = pd.read_csv(train_csv_path)
```

2.1.1 Data Preprocessing and Label Encoding

We need to represent categorical data into numerical form via encoding. This step should be done before EDA

```
[ ]: # Let's observe what are columns and their data types
df.dtypes
```

```
[ ]: name          object
     year          int64
     selling_price  int64
     km_driven      int64
     fuel           object
     seller_type    object
     transmission   object
     owner          object
     mileage        object
     engine         object
     max_power      object
     torque         object
     seats          float64
     dtype: object
```

The task is following (note: I will deal with nan values on the fly alongs tasks, I am thinking since its regression task our predictions is approximate (it is okay to change with mean/median based on distribution), but for classification tasks I think it is better to drop such rows. Therefore, I will keep them.):

1. Feature owner - map First owner to 1, ..., Test drive car to 5
2. Feature fuel - remove all rows with CNG and LPG because CNG and LPG use a different mileage
3. Feature mileage - remove "kmpl" and convert to float
4. Feature engine - remove "CC" and convert to numerical
5. Feature max power - same as engine
6. Feature brand - take first word and remove other
7. Drop feature torque
8. Test Drive cars are expensive, so delete all samples

```
[ ]: # task 1 - Feature owner - map First owner to 1, ..., Test drive car to 5
df_copy = df.copy()

# First Owner          5289
# Second Owner         2105
# Third Owner          555
# Fourth & Above Owner  174
# Test Drive Car 5

# Better to use one-hot encoding, but as per hw instructions doing mapping.
owner_map = {
    'owner': {
        "First Owner": 1,
        "Second Owner": 2,
        "Third Owner": 3,
        "Fourth & Above Owner": 4,
        "Test Drive Car": 5,
    }
}
```

```
df_copy.replace(owner_map, inplace=True)

# task8 - Test Drive cars are expensive, so delete all samples
df_copy = df_copy[df_copy.owner != 5]
print(df_copy.owner.value_counts())

# doing in such a sandwich way for testing purposes on the fly
df = df_copy.copy()
```

```
owner
1    5289
2    2105
3     555
4     174
Name: count, dtype: int64
```

```
C:\Users\eraco\AppData\Local\Temp\ipykernel_22248\1370062162.py:21:
FutureWarning: Downcasting behavior in `replace` is deprecated and will be
removed in a future version. To retain the old behavior, explicitly call
`result.infer_objects(copy=False)`. To opt-in to the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
  df_copy.replace(owner_map, inplace=True)
```

```
[ ]: # Encoder for binary categorical values
from sklearn.preprocessing import LabelEncoder

# defining encoder
le = LabelEncoder()
```

```
[ ]: # task2 Feature fuel - remove all rows with CNG and LPG
df_copy = df.copy()

# df_copy['fuel'].value_counts()
# CNG      57
# LPG      38

print(df_copy.shape)
df_copy = df_copy[~df_copy.fuel.isin(['CNG', 'LPG'])]

print(df_copy.shape)
df_copy.fuel.value_counts()

# And also let's encode it
df_copy.fuel = le.fit_transform(df_copy.fuel)
print(df_copy.fuel.value_counts())
```

```
df = df_copy.copy()
```

```
(8123, 13)
```

```
(8028, 13)
```

```
fuel
```

```
0    4401
```

```
1    3627
```

```
Name: count, dtype: int64
```

```
[ ]: # task6 - Feature brand - take first word and remove other
# same approach

df_copy = df.copy()

# Changing name to brand
df_copy.rename(columns = {'name': 'brand'}, inplace=True)
df_copy.brand = df_copy.brand.str.split().str[0]

print(df_copy.brand.isna().sum())

# Doing mapping

# Bad choice, I will proceed with one-hot encoding (though too much values)
# brand_name_map = {'brand': {v:k for k, v in zip(range(1, 33),
#         ['Maruti', 'Skoda', 'Honda', 'Hyundai', 'Toyota', 'Ford', 'Renault',
#         'Mahindra', 'Tata', 'Chevrolet', 'Fiat', 'Datsun', 'Jeep',
#         'Mercedes-Benz', 'Mitsubishi', 'Audi', 'Volkswagen', 'BMW',
#         'Nissan', 'Lexus', 'Jaguar', 'Land', 'MG', 'Volvo', 'Daewoo',
#         'Kia', 'Force', 'Ambassador', 'Ashok', 'Isuzu', 'Opel', 'Peugeot'])
# }
# }

# I will proceed with grouped one-hot encoding
group_map = {
    'Economy': ['Maruti', 'Tata', 'Hyundai', 'Datsun', 'Renault', 'Ford',
    ↪ 'Chevrolet', 'Fiat'],
    'Midrange': ['Honda', 'Toyota', 'Mahindra', 'Nissan', 'Skoda',
    ↪ 'Mitsubishi', 'Kia', 'MG'],
    'Luxury': ['Audi', 'BMW', 'Mercedes-Benz', 'Volvo', 'Jaguar', 'Lexus',
    ↪ 'Jeep', 'Land'],
    'Others': ['Daewoo', 'Ambassador', 'Ashok', 'Isuzu', 'Opel', 'Peugeot',
    ↪ 'Force']
}

# local mapper - later maybe need to define this in backend code
brand_to_group = {brand: group for group, brands in group_map.items() for brand
    ↪ in brands}
```

```

# mapping cars to its groups
df_copy.brand = df_copy.brand.map(brand_to_group)

# creating columns of brand grouping
df_encoded = pd.get_dummies(df_copy, columns=['brand'], drop_first=True)

df_encoded.head()

df = df_encoded.copy()

```

0

```

[ ]: # Transmission feature has 2 classes only, so use LabelEncoder
df_copy = df.copy()

df_copy.transmission = le.fit_transform(df_copy.transmission)
print(df_copy.transmission.value_counts())

df = df_copy.copy()

```

```

transmission
1    6982
0    1046
Name: count, dtype: int64

```

```

[ ]: # seller_type feature has 3 classes: individual, dealer, trustmark dealer ->
      ↪ use one-hot encoding
df_copy = df.copy()

# one-hot encoding, drop_first=True to drop one not required column
df_copy = pd.get_dummies(df_copy, columns=['seller_type'], drop_first=True)
df_copy.head()

df = df_copy.copy()

```

```

[ ]: # task3 - Feature mileage - remove "kmpl" and convert to float
      # Hint: use df_copy.mileage.str.split()

df.mileage = df.mileage.str.split().str[0].astype(float)

```

```

[ ]: # task4 - Feature engine - remove "CC" and convert to numerical
      # Same as task3

df.engine = df.engine.str.split().str[0].astype(float)

```

```

[ ]: # task5 - Feature max power - same as engine

```

```
df.max_power = df.max_power.str.split().str[0].astype(float)
```

```
[ ]: # task7 - dropping torque column
      # so that it would not have impact on EDA - even though its bad practice
```

```
df.drop(columns=['torque'], inplace=True)
```

```
[ ]: # Checking if everything is fine
      # But probably, it would be better to keep torque and transfer for numerical_
      ↪ form for the EDA basis
      # I will test it in next iteration
df.dtypes
```

```
[ ]: year                int64
      selling_price       int64
      km_driven           int64
      fuel                int32
      transmission        int32
      owner               int64
      mileage             float64
      engine              float64
      max_power           float64
      seats              float64
      brand_Luxury         bool
      brand_Midrange       bool
      brand_Others         bool
      seller_type_Individual bool
      seller_type_Trustmark Dealer bool
      dtype: object
```

2.1.2 Now we can proceed with EDA

3 2.2 Exploratory Data Analysis (EDA)

DataFrame columns:

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	name	8128 non-null	object
1	year	8128 non-null	int64
2	selling_price	8128 non-null	int64
3	km_driven	8128 non-null	int64
4	fuel	8128 non-null	object
5	seller_type	8128 non-null	object
6	transmission	8128 non-null	object
7	owner	8128 non-null	object
8	mileage	7907 non-null	object
9	engine	7907 non-null	object


```

10 max_power      7913 non-null    object
11 torque         7906 non-null    object
12 seats          7907 non-null    float64

```

3.0.1 General Notes about EDA:

value_counts(): Frequency counts

outliers: the value that is considerably higher or lower from rest of the data

Value at 75% is Q3 and value at 25% is Q1 -> Q stands for "quartile"

Outlier are smaller than $Q1 - 1.5(Q3-Q1)$ and bigger than $Q3 + 1.5(Q3-Q1)$. $(Q3-Q1) = IQR$

IQR stands for "interquartile range"

We will use describe() method. Describe method includes:

count: number of entries

mean: average of entries

std: standard deviation

min: minimum entry

25%: first quantile

50%: median or second quantile

75%: third quantile

max: maximum entry

```
[ ]: # Let's see all columns
df.columns
```

```
[ ]: Index(['year', 'selling_price', 'km_driven', 'fuel', 'transmission', 'owner',
          'mileage', 'engine', 'max_power', 'seats', 'brand_Luxury',
          'brand_Midrange', 'brand_Others', 'seller_type_Individual',
          'seller_type_Trustmark Dealer'],
          dtype='object')
```

```
[ ]: # Some basic info about each column
# We see there are null values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 8028 entries, 0 to 8127
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	year	8028 non-null	int64
1	selling_price	8028 non-null	int64
2	km_driven	8028 non-null	int64
3	fuel	8028 non-null	int32
4	transmission	8028 non-null	int32
5	owner	8028 non-null	int64
6	mileage	7814 non-null	float64
7	engine	7814 non-null	float64

8	max_power	7820 non-null	float64
9	seats	7814 non-null	float64
10	brand_Luxury	8028 non-null	bool
11	brand_Midrange	8028 non-null	bool
12	brand_Others	8028 non-null	bool
13	seller_type_Individual	8028 non-null	bool
14	seller_type_Trustmark Dealer	8028 non-null	bool

dtypes: bool(5), float64(4), int32(2), int64(4)
memory usage: 666.4 KB


3.1 Plotting:

```
[ ]: # # lets drop for now not number columns
# df_only_nums = df[['year', 'selling_price', 'km_driven', 'seats']]
# #only numbers

# previously was keeping categorical as categorical, but now we can plot all
↳ features
# print(df.corr())

# #correlation map
f, ax = plt.subplots(figsize=(15, 15))
sns.heatmap(df.corr(), annot=True, linewidths=.5, fmt= '.1f',ax=ax,
↳ cmap="coolwarm")

# So the more red color, the more correlation
plt.show()
```



```
[ ]: # Observing first 5 data
df.head(5)

# Observing last 5 data
#df.tail()
```

```
[ ]:   year  selling_price  km_driven  fuel  transmission  owner  mileage  engine \
0  2014         450000    145500    0           1         1     23.40   1248.0
1  2014         370000    120000    0           1         2     21.14   1498.0
2  2006         158000    140000    1           1         3     17.70   1497.0
3  2010         225000    127000    0           1         1     23.00   1396.0
```

4	2007	130000	120000	1	1	1	16.10	1298.0
---	------	--------	--------	---	---	---	-------	--------


	max_power	seats	brand_Luxury	brand_Midrange	brand_Others	\
0	74.00	5.0	False	False	False	
1	103.52	5.0	False	True	False	
2	78.00	5.0	False	True	False	
3	90.00	5.0	False	False	False	
4	88.20	5.0	False	False	False	

	seller_type_Individual	seller_type_Trustmark Dealer
0	True	False
1	True	False
2	True	False
3	True	False
4	True	False

Let's try to plot some line, scatter and histogram plots. To choose between, there are some differences in plots: - Line plot is better when x axis is time. - Box plots: visualize basic statistics like outliers, min/max or quantiles - Scatter is better when there is correlation between two variables - Histogram is better when we need to see distribution of numerical data. - Customization: Colors, labels, thickness of line, title, opacity, grid, figsize, ticks of axis and linestyle

```
[ ]: # Line plot
# It might be seen there is no correlation between features,
# but basically I am just exploring type of plots


# Line plot is better when x axis is time
df['selling_price'].plot(kind = 'line', color = 'g', label = 'Price',
linewidth=1, alpha = 0.5, grid = True, linestyle = ':')
df['km_driven'].plot(color = 'r', label = 'Driven', linewidth=1, alpha = 0.5, grid=True,
linestyle = '-.')
plt.legend(loc='upper right') # legend = puts label into plot
plt.xlabel('x axis') # label = name of label
plt.ylabel('y axis')
plt.title('Line Plot') # title = title of plot
plt.show()
```



```
[ ]: # Scatter plot
      # Scatter is better when there is correlation between two variables


      df.plot(kind='scatter', x='selling_price', y='km_driven',alpha = 0.5,color = 'red')
      plt.xlabel('Price')           # label = name of label
      plt.ylabel('Driven')
      plt.title('Scatter Plot')
```

```
[ ]: Text(0.5, 1.0, 'Scatter Plot')
```



```
[ ]: # Histogram
      # bins = number of bar in figure
      # Histogram is better when we need to see distribution of numerical data.

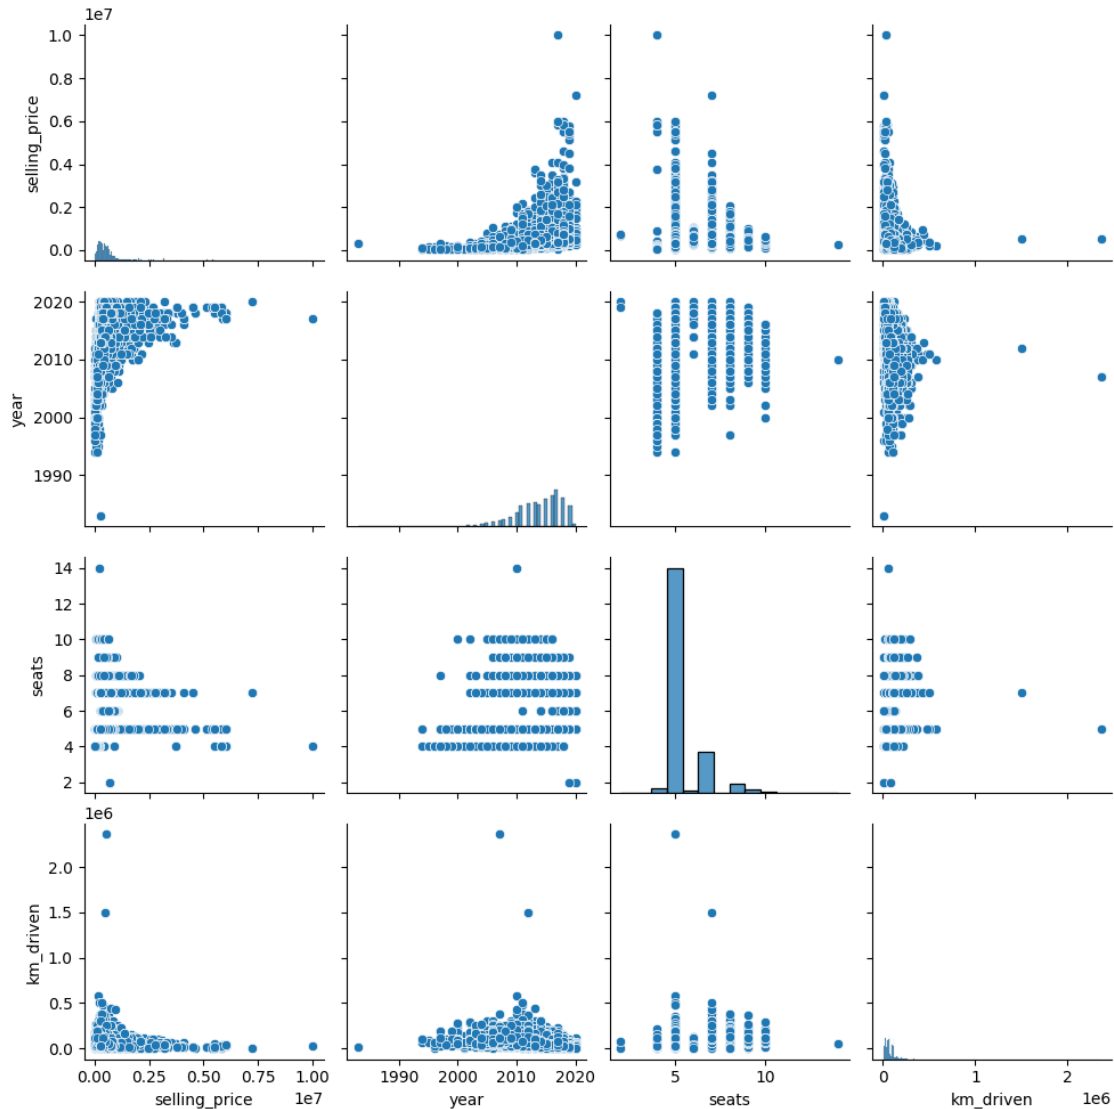
      df['selling_price'].plot(kind = 'hist',bins = 50,figsize = (4,4))
      plt.show()
```



```
[ ]: sns.pairplot(df[['selling_price', 'year', 'seats', 'km_driven']])
```

```
# From the result we can see that we need to normalize features
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x2252c9d3af0>
```



4 2.3 Feature Engineering

[08/18/2024] 1st attempt: I think there is no need to create new features, I will try with existing ones (created this section for future use - will require this)

5 2.4 Feature Selection

[08/18/2024] 1st attempt: I am thinking taking all features except the ones that we need to drop

[08/20/2024] 2nd attempt: I will choose 5 features that are most important: max_power, engine, transmission, fuel_type, and seats

```
[ ]: # Just to remind what are columns
df.columns
```



```
[ ]: Index(['year', 'selling_price', 'km_driven', 'fuel', 'transmission', 'owner',
          'mileage', 'engine', 'max_power', 'seats', 'brand_Luxury',
          'brand_Midrange', 'brand_Others', 'seller_type_Individual',
          'seller_type_Trustmark Dealer'],
          dtype='object')
```

5.0.1 Outliers

I want to handle them before proceeding to training ['year', 'km_driven', 'mileage', 'engine', 'max_power'] - will chose those only

[08/21/2024] 3rd attempt - trying with outliers before splitting the dataset

```
[ ]: # To see all outliers

def outlier_count(col, data = df):
    # calculate your 25% quatile and 75% quatile
    q75, q25 = np.percentile(data[col], [75, 25])

    # calculate your inter quatile
    iqr = q75 - q25

    # min_val and max_val
    min_val = q25 - (iqr*1.5)
    max_val = q75 + (iqr*1.5)

    # count number of outliers, which are the data that are less than min_val
    # or more than max_val calculated above
    outlier_count = len(np.where((data[col] > max_val) | (data[col] <
    min_val))[0])

    # calculate the percentage of the outliers
    outlier_percent = round(outlier_count/len(data[col])*100, 2)

    if(outlier_count > 0):
        print("\n"+15*'- ' + col + 15*'- '+ "\n")
        print('Number of outliers: {}'.format(outlier_count))
        print('Percent of data that is outlier: {}'.format(outlier_percent))

[ ]: # Printing outliers per column
for col in ['year', 'km_driven', 'mileage', 'engine', 'max_power']:
    outlier_count(col)
```

-----year-----

Number of outliers: 78

Percent of data that is outlier: 0.97%

-----km_driven-----

Number of outliers: 168

Percent of data that is outlier: 2.09%

Let's not remove them, but cap them to a fixed value (5th or 95th percentile) - reduce impact of extreme values

```
[ ]: # Capping outliers
def cap_outliers(df, column):
    lower_limit = df[column].quantile(0.05)
    upper_limit = df[column].quantile(0.95)
    df[column] = np.where(df[column] < lower_limit, lower_limit, df[column])
    df[column] = np.where(df[column] > upper_limit, upper_limit, df[column])
    return df

# applying for 'year' and 'km_driven' since there are only two outliers
# for chosen set of features
df = cap_outliers(df, 'year')
df = cap_outliers(df, 'km_driven')
```

```
[ ]: # Printing outliers per column
for col in ['year', 'km_driven', 'mileage', 'engine', 'max_power']:
    outlier_count(col)
```

```
[ ]: # Same approach as in label encoding
df_copy = df.copy()

# shape (m,)
y = df_copy['selling_price']
# df_copy = df_copy.drop(columns=['selling_price'])
print(y.shape)
assert len(y.shape) == 1

# Taking shape (m, n)
X = df_copy[['year', 'km_driven', 'mileage', 'engine', 'max_power']]
print(X.shape)
assert len(X.shape) == 2
```

(8028,)

(8028, 5)

Converting the regression problem into classification problem:

```
[ ]: test = y.copy()
```

```
[ ]: q25 = y.quantile(0.25)
q50 = y.quantile(0.50)
q75 = y.quantile(0.75)
```

```
def regr_to_class(val):
    return 0 if val < q25 else 1 if q25 <= val < q50 else 2 if q50 <= val < q75
    ↪ else 3

y = y.apply(regr_to_class)
```

```
[ ]: q25, q50, q75
```

```
[ ]: (260000.0, 450000.0, 680000.0)
```

```
[ ]: y.value_counts()
```

```
[ ]: selling_price
2    2113
3    2013
0    2003
1    1899
Name: count, dtype: int64
```

```
[ ]: from sklearn.model_selection import train_test_split

# Splitting the dataset, will proceed with processing it
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.13,
    ↪ random_state=42)
```

6 2.5 Preprocessing

```
[ ]: # Let's see what is the train dataset size
X_train.shape
```

```
[ ]: (6984, 5)
```

```
[ ]: # Same for test
X_test.shape
```

```
[ ]: (1044, 5)
```

6.0.1 NULL values

```
[ ]: # Let's observe all null values in training set (did not deal with them - to
    ↪ avoid data leakage)
X_train.isna().sum()
```

```
[ ]: year          0
     km_driven     0
     mileage      187
```

```
engine      187
max_power    181
dtype: int64
```

```
[ ]: # Same for the testing dataset
X_test.isna().sum()
```

```
[ ]: year      0
km_driven    0
mileage      27
engine       27
max_power    27
dtype: int64
```

```
[ ]: # Removing null values for mileage
print(X_train.mileage.mean(), X_train.mileage.median())

# sns.distplot(X_train, x=X_train.mileage)

# Interchanging nan values with mean - the distribution is normal
X_train.mileage.fillna(X_train.mileage.mean(), inplace=True)
X_test.mileage.fillna(X_train.mileage.mean(), inplace=True)
```

```
19.38204354862439 19.3
```

<ipython-input-37-4dc187eca875>:7: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
X_train.mileage.fillna(X_train.mileage.mean(), inplace=True)
```

<ipython-input-37-4dc187eca875>:8: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
X_test.mileage.fillna(X_train.mileage.mean(), inplace=True)
```

```
[ ]: # df_copy.engine.isna().sum() # 214
print(X_train.engine.mean(), X_train.engine.median())

# sns.distplot(X_train, x=X_train['engine'])

# Interchanging nan values with median - the distribution is skewed
X_train.engine.fillna(X_train.engine.mean(), inplace=True)
X_test.engine.fillna(X_train.engine.mean(), inplace=True)
```

```
1463.756068853906 1248.0
```

<ipython-input-38-9d9e8b96bc3f>:7: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
X_train.engine.fillna(X_train.engine.mean(), inplace=True)
```

<ipython-input-38-9d9e8b96bc3f>:8: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
X_test.engine.fillna(X_train.engine.mean(), inplace=True)
```

```
[ ]: # df_copy.max_power.isna().sum() # 208
print(X_train.max_power.mean(), X_train.max_power.median())

# sns.distplot(X_train, x=X_train.max_power) # distribution is skewed a little

# Interchanging nan values with median - the distribution is skewed
```

```
X_train.max_power.fillna(X_train.max_power.mean(), inplace=True)
X_test.max_power.fillna(X_train.max_power.mean(), inplace=True)
```

```
91.74543877701014 82.85
```

<ipython-input-39-9d3173a3f2c6>:7: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
X_train.max_power.fillna(X_train.max_power.mean(), inplace=True)
<ipython-input-39-9d3173a3f2c6>:8: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
```

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
X_test.max_power.fillna(X_train.max_power.mean(), inplace=True)
```

```
[ ]: # # And we want to remove all null values from seats feature - 214 rows
      # print(X_train.seats.mean(), X_train.seats.median())

      # # sns.distplot(X_train, x=X_train.seats) # distribution is skewed a little

      # # Interchanging nan values with median - the distribution is skewed
      # X_train.seats.fillna(X_train.seats.mean(), inplace=True)
      # X_test.seats.fillna(X_train.seats.mean(), inplace=True)
```

```
[ ]: # Now verify if everything is fine
      X_train.isna().sum()
```

```
[ ]: year          0
      km_driven     0
      mileage       0
      engine        0
```

```
max_power    0
dtype: int64
```

```
[ ]: # Same for test set
X_test.isna().sum()
```

```
[ ]: year        0
km_driven      0
mileage        0
engine         0
max_power      0
dtype: int64
```

```
[ ]: # Just to be sure
y_train.isna().sum()
```

```
[ ]: 0
```

```
[ ]: # Now we can proceed
y_test.isna().sum()
```

```
[ ]: 0
```

6.1 Scaling

```
[ ]: # Observing what need to be scaled
X_train.head()
```

```
[ ]:      year  km_driven  mileage  engine  max_power
4419  2016.0    68089.0    19.16  2494.0    157.70
6103  2011.0    81500.0    14.84  2143.0    167.62
7893  2011.0   140000.0    14.40  1598.0    103.60
7427  2016.0   120000.0    13.58  2499.0     72.40
1448  2018.0    30000.0    18.60  1197.0     81.83
```

```
[ ]: # We need to scale all numerics whose difference is large
from sklearn.preprocessing import StandardScaler

# After observing above, we can proceed with the following columns
col_names = ['year', 'km_driven', 'mileage', 'engine', 'max_power']

# Defining Scaler
sc = StandardScaler()

# Scaling is performed
X_train[col_names] = sc.fit_transform(X_train[col_names])
X_test[col_names] = sc.transform(X_test[col_names])
```

```
[ ]: # Let's see if its fine  
X_train.head()
```

```
[ ]:  
      year  km_driven  mileage  engine  max_power  
4419  0.559461    0.028728 -0.056110  2.064080    1.885161  
6103 -0.806911    0.364115 -1.147758  1.360856    2.168701  
7893 -0.806911    1.827105 -1.258944  0.268956    0.338836  
7427  0.559461    1.326937 -1.466156  2.074098   -0.552945  
1448  1.106010   -0.923816 -0.197620 -0.534442   -0.283410
```

```
[ ]: # Same for test set  
X_test.head()
```

```
[ ]:  
      year  km_driven  mileage  engine  max_power  
5948 -0.806911   -0.798774 -1.107327 -0.191846   -0.049889  
6039  0.559461    1.076853  0.072766  0.236900    0.984805  
3069  0.286187   -0.173565  0.285031 -0.933136   -0.706149  
6531  0.012912    0.326602 -0.094014 -0.005522    0.477462  
322   1.106010   -0.948824  1.166941 -0.556481   -0.508070
```

```
[ ]: # Same for selling price, we want to do np.log transformation  
# y_train = np.log(y_train)  
# y_train  
  
# We dont need to log the actual target set - so leaving it  
# y_test = np.log(y_test)  
  
X_train = X_train.to_numpy()  
y_train = y_train.to_numpy()  
  
X_test = X_test.to_numpy()  
y_test = y_test.to_numpy()
```

7 3. Modeling

Configuring MLFlow workplace

```
[ ]: # Task 3  
mlflow_url = 'https://mlflow.ml.brain.cs.ait.ac.th/'  
mlflow.set_tracking_uri(mlflow_url)  
  
os.environ["MLFLOW_TRACKING_USERNAME"] = "admin"  
os.environ["MLFLOW_TRACKING_PASSWORD"] = "password"  
  
os.environ["LOGNAME"] = "st125457-ulugbek"  
  
mlflow.set_experiment(experiment_name="st125457-a3-ulugbek-experiment")
```



```
import logging
```

```
# Setting only logging warning messages, later it would make output neat  
logging.getLogger("mlflow").setLevel(logging.WARNING)
```

```
[ ]: import json
```

```
# Task 2
```

```
class RidgePenalty:
```

```
    def __init__(self, l):  
        self.l = l
```

```
    def __call__(self, theta): #__call__ allows us to call class as method  
        return self.l * np.sum(np.square(theta))
```

```
    def derivation(self, theta):  
        return self.l * 2 * theta
```

```
class LogisticRegression:
```

```
    def __init__(self, k, n, method='batch', batch_size=512, alpha=0.001,  
↪max_iter=5000, regularization=True, l=0.1):  
        self.k = k  
        self.n = n  
        self.method = method  
        self.alpha = alpha  
        self.max_iter = max_iter  
        self.batch_size = batch_size
```

```
        self.l = l
```

```
        self.reg_flag = regularization  
        self.regularization = RidgePenalty(self.l) if self.reg_flag else None
```

```
        self.W = None
```

```
    def fit(self, X, Y):  
        if self.method not in ["batch", "minibatch", "sto"]:  
            raise ValueError('Method must be one of the followings: "batch",  
↪"minibatch" or "sto".')
```

```
        # Creating new set of weights each time fit function is called is bad  
↪idea
```

```

        # probably we do not want to do it that way -> initialization during
        ↪init? - for kfold, introduce flag?
        np.random.seed(42)
        self.W = np.random.rand(self.n, self.k)

        self.losses = []

        Y = np.eye(self.k)[Y]
        # with mlflow.start_run(run_name=f"{type(self).__name__}", nested=True):
        #     params = {"method": self.method, "lr": self.alpha, "reg": ↪
        ↪type(self).__name__, "regularization": self.reg_flag}
        #     mlflow.log_params(params=params)
        self.val_loss_old = np.infty

        if self.method == "batch":
            for i in range(self.max_iter):
                loss, grad = self.gradient(X, Y)
                self.losses.append(loss)

                self.W = self.W - self.alpha * grad

                # if i % 10 == 0:
                #     print(f"Loss at iteration {i}", loss)

                # if i == 30:
                #     self.alpha = 0.0009025
            if i % 125 == 0:
                self.learning_rate_decay(i)

                # print(f'Learning rate: ', self.alpha)

            mlflow.log_metric(key="train_loss", value=loss, step=i)

        elif self.method == "minibatch":
            for i in range(self.max_iter):
                ix = np.random.randint(0, X.shape[0]) #<----with replacement
                batch_X = X[ix:ix+self.batch_size]
                batch_Y = Y[ix:ix+self.batch_size]
                loss, grad = self.gradient(batch_X, batch_Y)
                self.losses.append(loss)
                self.W = self.W - self.alpha * grad
                if i % 10 == 0:
                    # print(f"Loss at iteration {i}", loss)

                    if i % 90 == 0:
                        self.learning_rate_decay(i)
                    # print(f'Learning rate: ', self.alpha)

```

```

        mlflow.log_metric(key="train_loss", value=loss, step=i)

    elif self.method == "sto":
        for i in range(self.max_iter):
            idx = np.random.randint(X.shape[0])
            X_train = X[idx, :].reshape(1, -1)
            Y_train = Y[idx]

            loss, grad = self.gradient(X_train, Y_train)
            self.losses.append(loss)

            self.W = self.W - self.alpha * grad

            # if i % 500 == 0:
            #     print(f"Loss at iteration {i}", loss)

            mlflow.log_metric(key="train_loss", value=loss, step=i)

        return self

    # if np.allclose(loss, self.val_loss_old):
    #     print(f"Break - Loss at iteration {i}", loss)

    # self.val_loss_old = loss

    # print(f"Time taken: {time.time() - start_time}")

def learning_rate_decay(self, epoch, min_lr=1e-8):
    new_alpha = self.alpha * (0.95 ** (epoch // 10))
    # new_alpha = self.alpha * 0.1
    self.alpha = max(new_alpha, min_lr)

def predict(self, X_test):
    return np.argmax(self.h_theta(X_test, self.W), axis=1)

def score(self, X, Y):
    """Score method required by GridSearchCV."""
    Y_pred = self.predict(X)
    accuracy = np.mean(Y_pred == Y)
    return accuracy

def gradient(self, X, Y):
    m = X.shape[0]
    h = self.h_theta(X, self.W)

    loss = - np.sum(Y * np.log(h)) / m

```

```

        if self.reg_flag:
            loss += self.regularization(self.W) / (m) # Divide by 2*m to
↪normalize

        error = h - Y

        grad = self.softmax_grad(X, error)

        if self.reg_flag:
            grad += (self.regularization.derivation(self.W) / m)

        return loss, grad

    def softmax(self, theta_t_x):
        return np.exp(theta_t_x) / np.sum(np.exp(theta_t_x), axis=1,
↪keepdims=True)

    def softmax_grad(self, X, error):
        return X.T @ error

    def h_theta(self, X, W):
        return self.softmax(X @ W)

    # For performing GridSearch
    def get_params(self, deep=True):
        """Get the parameters for GridSearchCV."""
        return {
            'k': self.k,
            'n': self.n,
            'method': self.method,
            'batch_size': self.batch_size,
            'alpha': self.alpha,
            'max_iter': self.max_iter,
            'regularization': self.reg_flag,
            'l': self.l
        }

    # For performing GridSearch
    def set_params(self, **params):
        """Set the parameters for GridSearchCV."""
        for key, value in params.items():
            setattr(self, key, value)
        return self

    def plot(self):
        plt.plot(np.arange(len(self.losses)), self.losses, label = "Train
↪Losses")

```

```

plt.title("Losses")
plt.xlabel("epoch")
plt.ylabel("losses")
plt.legend()

# Task 1
@staticmethod
def accuracy(y_true, y_pred):
    correct_predictions = np.sum(y_true == y_pred)
    total_predictions = len(y_true)

    return correct_predictions / total_predictions

@staticmethod
def precision(y_true, y_pred, class_label):
    TP = np.sum((y_true == class_label) & (y_pred == class_label))
    FP = np.sum((y_true != class_label) & (y_pred == class_label))

    return TP / (TP + FP) if TP + FP > 0 else 0.0

@staticmethod
def recall(y_true, y_pred, class_label):
    TP = np.sum((y_true == class_label) & (y_pred == class_label))
    FN = np.sum((y_true == class_label) & (y_pred != class_label))

    return TP / (TP + FN) if TP + FN > 0 else 0.0

@staticmethod
def f1_score(y_true, y_pred, class_label):
    P = LogisticRegression.precision(y_true, y_pred, class_label)
    R = LogisticRegression.recall(y_true, y_pred, class_label)

    return 2 * P * R / (P + R) if (P + R) > 0 else 0.0

@staticmethod
def macro_precision(y_true, y_pred):
    classes = np.unique(y_true)
    precisions = [LogisticRegression.precision(y_true, y_pred, class_label)
    ↪for class_label in classes]

    return np.mean(precisions)

@staticmethod
def macro_recall(y_true, y_pred):
    classes = np.unique(y_true)
    recalls = [LogisticRegression.recall(y_true, y_pred, class_label) for
    ↪class_label in classes]

```

```

        return np.mean(recalls)

    @staticmethod
    def macro_f1(y_true, y_pred):
        classes = np.unique(y_true)
        f1_scores = [LogisticRegression.f1_score(y_true, y_pred, class_label)
        ↪for class_label in classes]

        return np.mean(f1_scores)

    @staticmethod
    def weighted_precision(y_true, y_pred):
        classes = np.unique(y_true)
        total_samples = len(y_true)

        weights = [(np.sum(y_true == class_label) / total_samples) for
        ↪class_label in classes]
        precisions = [LogisticRegression.precision(y_true, y_pred, class_label)
        ↪for class_label in classes]

        return np.sum([precision * w_coef for precision, w_coef in
        ↪zip(precisions, weights)])

    @staticmethod
    def weighted_recall(y_true, y_pred):
        classes = np.unique(y_true)
        total_samples = len(y_true)

        weights = [(np.sum(y_true == class_label) / total_samples) for
        ↪class_label in classes]
        recalls = [LogisticRegression.recall(y_true, y_pred, class_label) for
        ↪class_label in classes]

        return np.sum([recall * w_coef for recall, w_coef in zip(recalls,
        ↪weights)])

    @staticmethod
    def weighted_f1(y_true, y_pred):
        classes = np.unique(y_true)
        total_samples = len(y_true)

        weights = [(np.sum(y_true == class_label) / total_samples) for
        ↪class_label in classes]
        f1_scores = [LogisticRegression.f1_score(y_true, y_pred, class_label)
        ↪for class_label in classes]

```

```

        return np.sum([f1_score * w_coef for f1_score, w_coef in zip(f1_scores,
↪weights)])

    @staticmethod
    def classification_report(y_true, y_pred):
        classes = np.unique(y_true)
        report = []
        header = f"{'Class':<16}{'Precision':<12}{'Recall':<12}{'F1-Score':
↪<12}{'Support':<10}"
        report.append(header)
        report.append("=" * len(header))

        for class_label in classes:
            precision = LogisticRegression.precision(y_true, y_pred,
↪class_label)
            recall = LogisticRegression.recall(y_true, y_pred, class_label)
            f1 = LogisticRegression.f1_score(y_true, y_pred, class_label)
            support = np.sum(y_true == class_label)

            report.append(
                f"{class_label:<16}{precision:<12.4f}{recall:<12.4f}{f1:<12.
↪4f}{support:<10}"
            )

        report.append("=" * len(header))
        macro_precision = LogisticRegression.macro_precision(y_true, y_pred)
        macro_recall = LogisticRegression.macro_recall(y_true, y_pred)
        macro_f1 = LogisticRegression.macro_f1(y_true, y_pred)

        weighted_precision = LogisticRegression.weighted_precision(y_true,
↪y_pred)
        weighted_recall = LogisticRegression.weighted_recall(y_true, y_pred)
        weighted_f1 = LogisticRegression.weighted_f1(y_true, y_pred)

        accuracy = LogisticRegression.accuracy(y_true, y_pred)
        total_support = len(y_true)

        report.append(
            f"{'Accuracy':<16}{'':<12}{'':<12}{accuracy:<12.4f}{total_support:
↪<10}"
        )

        report.append(
            f"{'Macro Avg':<16}{macro_precision:<12.4f}{macro_recall:<12.
↪4f}{macro_f1:<12.4f}{total_support:<10}"
        )

```

```

    )
    report.append(
        f"{'Weighted Avg':<16}{weighted_precision:<12.4f}{weighted_recall:
↪<12.4f}{weighted_f1:<12.4f}{total_support:<10}"
    )

    return "\n".join(report)

# For understanding which feature is important
def plot_feature_importance(self, feature_names=None):
    if not hasattr(self, 'W'):
        raise ValueError("Model coefficients are not available. Fit the
↪model first.")

    # Coefficients
    coefficients = self._coef()
    importance = np.abs(coefficients)

    # Assign default names if feature_names are not provided
    if feature_names is None:
        feature_names = [f"Feature {i}" for i in range(1, len(coefficients)
↪+ 1)]

    # Sort features by importance
    mask = np.argsort(importance)[::-1]
    sorted_importance = importance[mask]
    sorted_feature_names = np.array(feature_names)[mask]

    sorted_feature_names = sorted_feature_names.tolist() if
↪isinstance(sorted_feature_names, np.ndarray) else sorted_feature_names
    sorted_importance = sorted_importance.tolist() if
↪isinstance(sorted_importance, np.ndarray) else sorted_importance

    # Plotting
    plt.figure(figsize=(10, 6))
    plt.barh(sorted_feature_names, sorted_importance, color='skyblue')
    plt.xlabel('Coefficient Magnitude (Absolute)')
    plt.title('Feature Importance based on Coefficients')
    plt.gca().invert_yaxis() # To display the most important feature at
↪the top
    plt.show()

    def _coef(self):
        return self.W

```



```

def save(self, file_path):
    with open(file_path, 'w') as f:
        json.dump({'coefficients': self.W.tolist()}, f)

@classmethod
def load(cls, file_path):
    with open(file_path, 'r') as f:
        data = json.load(f)
    return cls(theta=np.array(data['coefficients']))

```

```

[ ]: from sklearn.model_selection import GridSearchCV
k = 4

# Performing Gridsearch
model = LogisticRegression(k=k, n=X_train.shape[1])

param_grid = {
    'alpha': [0.001, 0.01],      # Learning rates
    'l': [0.01, 0.1, 1.0],       # Ridge regularization strengths
    'batch_size': [32, 64],      # Batch sizes
    'max_iter': [200],
    'method': ['batch', 'minibatch', 'sto']
}

grid_search = GridSearchCV(estimator=model, param_grid=param_grid, cv=3,
    ↳scoring='accuracy')
grid_search.fit(X_train, y_train)
best_model = grid_search.best_estimator_
best_params = grid_search.best_params_

```

```

-----
AttributeError                                Traceback (most recent call last)
<ipython-input-54-f2177fc18a87> in <cell line: 19>()
    17 best_model = grid_search.best_estimator_
    18 best_params = grid_search.best_params_
----> 19 grid_search.plot()

AttributeError: 'GridSearchCV' object has no attribute 'plot'

```

- What does support in the classification report means?

Support refers to the number of actual occurrences of each class in the true labels. It helps to understand the distribution of classes in your dataset, which is particularly useful when dealing with imbalanced datasets. It doesn't affect metrics like precision, recall, or F1-score directly, but it provides context about how many instances of each class were present when calculating these metrics.

```
[ ]: # k, n, method, batch_size=32, alpha=0.001, max_iter=5000, regularization=True,
      ↪l=0.1):
methods = ['sto', 'minibatch', 'batch']
regularization = [False, True]

mlflow_params = []
k = len(np.unique(y_train))

for method in methods:
    for reg in regularization:
        mlflow_prm = f"LogisticRegression_{k=}_{method=}_{reg=}"
        mlflow_params.append(mlflow_prm)

        with mlflow.start_run(run_name=f"{mlflow_prm}", nested=True):
            params = {"method": method, "reg": "LogisticRegression",
            ↪"regularization": reg}
            mlflow.log_params(params=params)
            print(params)
            model = LogisticRegression(k, X_train.shape[1], method=method,
            ↪regularization=reg, max_iter=300, alpha=0.001)

            model.fit(X_train, y_train)

            signature = mlflow.models.infer_signature(X_train, model.
            ↪predict(X_train))
            mlflow.sklearn.log_model(model, artifact_path='model',
            ↪signature=signature)

            yhat = model.predict(X_test)
            print(LogisticRegression.classification_report(y_test, yhat))
            print("Report: ", classification_report(y_test, yhat, digits=4))
```

```
{'method': 'sto', 'reg': 'LogisticRegression', 'regularization': False}
```

Class	Precision	Recall	F1-Score	Support
0	0.4299	0.3446	0.3825	267
1	0.2840	0.2584	0.2706	267
2	0.2808	0.2840	0.2824	257
3	0.0856	0.1107	0.0966	253
Accuracy			0.2510	1044
Macro Avg	0.2701	0.2494	0.2580	1044
Weighted Avg	0.2724	0.2510	0.2600	1044
Report:	precision	recall	f1-score	support
0	0.4299	0.3446	0.3825	267
1	0.2840	0.2584	0.2706	267

2	0.2808	0.2840	0.2824	257
3	0.0856	0.1107	0.0966	253
accuracy			0.2510	1044
macro avg	0.2701	0.2494	0.2580	1044
weighted avg	0.2724	0.2510	0.2600	1044

```
{'method': 'sto', 'reg': 'LogisticRegression', 'regularization': True}
```

Class	Precision	Recall	F1-Score	Support
0	0.4326	0.3483	0.3859	267
1	0.2869	0.2622	0.2740	267
2	0.2797	0.2840	0.2819	257
3	0.0864	0.1107	0.0971	253

Accuracy			0.2529	1044
Macro Avg	0.2714	0.2513	0.2597	1044
Weighted Avg	0.2738	0.2529	0.2617	1044
Report:	precision	recall	f1-score	support

0	0.4326	0.3483	0.3859	267
1	0.2869	0.2622	0.2740	267
2	0.2797	0.2840	0.2819	257
3	0.0864	0.1107	0.0971	253

accuracy			0.2529	1044
macro avg	0.2714	0.2513	0.2597	1044
weighted avg	0.2738	0.2529	0.2617	1044

```
{'method': 'minibatch', 'reg': 'LogisticRegression', 'regularization': False}
```

Class	Precision	Recall	F1-Score	Support
0	0.5896	0.9363	0.7236	267
1	0.6623	0.1910	0.2965	267
2	0.5934	0.4202	0.4920	257
3	0.6094	0.8696	0.7166	253

Accuracy			0.6025	1044
Macro Avg	0.6137	0.6043	0.5572	1044
Weighted Avg	0.6139	0.6025	0.5557	1044
Report:	precision	recall	f1-score	support

0	0.5896	0.9363	0.7236	267
1	0.6623	0.1910	0.2965	267
2	0.5934	0.4202	0.4920	257
3	0.6094	0.8696	0.7166	253

accuracy			0.6025	1044
----------	--	--	--------	------

macro avg	0.6137	0.6043	0.5572	1044
weighted avg	0.6139	0.6025	0.5557	1044

```
{'method': 'minibatch', 'reg': 'LogisticRegression', 'regularization': True}
```

Class	Precision	Recall	F1-Score	Support	
0	0.5896	0.9363	0.7236	267	
1	0.6623	0.1910	0.2965	267	
2	0.5934	0.4202	0.4920	257	
3	0.6094	0.8696	0.7166	253	
=====					
Accuracy			0.6025	1044	
Macro Avg	0.6137	0.6043	0.5572	1044	
Weighted Avg	0.6139	0.6025	0.5557	1044	
Report:		precision	recall	f1-score	support
	0	0.5896	0.9363	0.7236	267
	1	0.6623	0.1910	0.2965	267
	2	0.5934	0.4202	0.4920	257
	3	0.6094	0.8696	0.7166	253

accuracy			0.6025	1044
macro avg	0.6137	0.6043	0.5572	1044
weighted avg	0.6139	0.6025	0.5557	1044

```
{'method': 'batch', 'reg': 'LogisticRegression', 'regularization': False}
```

Class	Precision	Recall	F1-Score	Support	
0	0.5939	0.9476	0.7302	267	
1	0.6500	0.1948	0.2997	267	
2	0.5819	0.4008	0.4747	257	
3	0.6094	0.8696	0.7166	253	
=====					
Accuracy			0.6015	1044	
Macro Avg	0.6088	0.6032	0.5553	1044	
Weighted Avg	0.6091	0.6015	0.5539	1044	
Report:		precision	recall	f1-score	support
	0	0.5939	0.9476	0.7302	267
	1	0.6500	0.1948	0.2997	267
	2	0.5819	0.4008	0.4747	257
	3	0.6094	0.8696	0.7166	253

accuracy			0.6015	1044
macro avg	0.6088	0.6032	0.5553	1044
weighted avg	0.6091	0.6015	0.5539	1044

```
{'method': 'batch', 'reg': 'LogisticRegression', 'regularization': True}
```

Class	Precision	Recall	F1-Score	Support
0	0.5939	0.9476	0.7302	267
1	0.6500	0.1948	0.2997	267
2	0.5819	0.4008	0.4747	257
3	0.6094	0.8696	0.7166	253

Accuracy			0.6015	1044
Macro Avg	0.6088	0.6032	0.5553	1044
Weighted Avg	0.6091	0.6015	0.5539	1044

Report:	precision	recall	f1-score	support
0	0.5939	0.9476	0.7302	267
1	0.6500	0.1948	0.2997	267
2	0.5819	0.4008	0.4747	257
3	0.6094	0.8696	0.7166	253

accuracy			0.6015	1044
macro avg	0.6088	0.6032	0.5553	1044
weighted avg	0.6091	0.6015	0.5539	1044

```
[ ]: # So best model has the following params: {'method': 'minibatch', 'reg':
    ↪ 'LogisticRegression', 'regularization': True}
k = 4
method = 'minibatch'
reg = True
mlflow_prm = f"LogisticRegression_{k=}_{method=}_{reg=}"
with mlflow.start_run(run_name=f"{mlflow_prm}", nested=True):
    params = {"method": method, "reg": "LogisticRegression", "regularization":
    ↪ reg}
    mlflow.log_params(params=params)
    print(params)
    model = LogisticRegression(k, X_train.shape[1], method=method,
    ↪ regularization=reg, max_iter=300, alpha=0.001)

    model.fit(X_train, y_train)

    signature = mlflow.models.infer_signature(X_train, model.predict(X_train))
    mlflow.sklearn.log_model(model, artifact_path='model', signature=signature)

    yhat = model.predict(X_test)
    print(LogisticRegression.classification_report(y_test, yhat))
    print("Report: ", classification_report(y_test, yhat, digits=4))
```

```
{'method': 'minibatch', 'reg': 'LogisticRegression', 'regularization': True}
```

Class	Precision	Recall	F1-Score	Support
-------	-----------	--------	----------	---------

0	0.5896	0.9363	0.7236	267
1	0.6623	0.1910	0.2965	267
2	0.5934	0.4202	0.4920	257
3	0.6094	0.8696	0.7166	253

=====


Accuracy			0.6025	1044
Macro Avg	0.6137	0.6043	0.5572	1044
Weighted Avg	0.6139	0.6025	0.5557	1044

Report:

	precision	recall	f1-score	support
0	0.5896	0.9363	0.7236	267
1	0.6623	0.1910	0.2965	267
2	0.5934	0.4202	0.4920	257
3	0.6094	0.8696	0.7166	253

accuracy			0.6025	1044
macro avg	0.6137	0.6043	0.5572	1044
weighted avg	0.6139	0.6025	0.5557	1044

```
[ ]: model.plot()
```



```
[ ]: # Saving the model
import joblib

# save the model to disk
model_name = 'Logistic_regression'
filename = f'{model_name}_model.pkl'
joblib.dump(model, filename)

model.save(f'{model_name}_weights.json')
# with open('polynomial_regression_model.pkl', 'wb') as file:
#     pickle.dump(model, file)
```

8 Report - Conclusion

So regarding the findings, let's break it down into parts: 1. I have been using the same features 2. The loss function's result is not the best, but since we are mapping range of values (wide range) into single class, I suppose this result is fine. 3. Problem with dataset has arranged, since we mapped regression problem into classification - slight errors in prediction now cause to fall down into another class if the values of regression problem is at the border of the class. 4. Possible understanding is choosing more features and feature engineering which will increase the results. 5. Implemented also the possibility to run GridSearch on the model - was running like 10 hours based on the params provided above - and it still was running - infinite runner :) 6. More details regarding scores can be found on ML Flow. 7. Also implemented learning rate decay