

classwork

October 7, 2024

```
[ ]: import torch
import torch.nn as nn

from torchvision import datasets, transforms
from torch.utils.data import DataLoader
from torch.utils.data import TensorDataset
```

```
[ ]: torch.__version__
```

```
[ ]: '2.4.0+cu118'
```

```
[ ]: torch.manual_seed(42)
```

```
[ ]: <torch._C.Generator at 0x281abee4710>
```

```
[ ]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
device
```

```
[ ]: device(type='cuda', index=0)
```

```
[ ]: map_courses = {0: 'DSAI', 1: 'IM', 2: 'WEM', 3: 'GIS', 4: 'RGLS', 5: 'CS'}
```

```
[ ]: import numpy as np

x_train = np.array([[400, 0, 28, 10],
                    [100, 1, 26, 8],
                    [200, 0, 29, 2],
                    [700, 2, 23, 5],
                    [900, 3, 27, 9],
                    [800, 4, 40, 5],
                    [500, 5, 29, 7],
                    [600, 0, 35, 10]], dtype='float32')

y_train = np.array([3.9, 3.7, 2.4, 2.9, 3.9, 3.0, 3.0, 3.1], dtype='float32')

inputs = torch.from_numpy(x_train)
targets = torch.from_numpy(y_train)
print(inputs.size())
```

```
print(targets.size())
```

```
torch.Size([8, 4])  
torch.Size([8])
```

```
[ ]: train_ds = TensorDataset(inputs, targets)  
      train_ds[0:3]
```

```
[ ]: (tensor([[400.,  0., 28., 10.],  
             [100.,  1., 26.,  8.],  
             [200.,  0., 29.,  2.]]),  
      tensor([3.9000, 3.7000, 2.4000]))
```

```
[ ]: batch_size = 3  
      train_dl = DataLoader(train_ds, batch_size, shuffle=True)
```

```
[ ]: for xb, yb in train_dl:  
      print(xb)  
      print(yb)  
      break
```

```
tensor([[400.,  0., 28., 10.],  
        [800.,  4., 40.,  5.],  
        [100.,  1., 26.,  8.]])  
tensor([3.9000, 3.0000, 3.7000])
```

```
[ ]: class Custom(nn.Module):  
      def __init__(self):  
          super().__init__()  
          self.fc = nn.Linear(4, 1)  
  
      def forward(self, x):  
          x = self.fc(x)  
          return x
```

```
[ ]: model = Custom()  
      print(model.fc.weight)  
      print(model.fc.weight.size()) # (out_features, in_features)  
      print(model.fc.bias)  
      print(model.fc.bias.size())
```

```
Parameter containing:  
tensor([[ 0.0472, -0.4938,  0.4516, -0.4247]], requires_grad=True)  
torch.Size([1, 4])  
Parameter containing:  
tensor([0.3860], requires_grad=True)  
torch.Size([1])
```

```
[ ]: list(model.parameters())
```

```
[ ]: [Parameter containing:
      tensor([[ 0.0472, -0.4938,  0.4516, -0.4247]], requires_grad=True),
      Parameter containing:
      tensor([0.3860], requires_grad=True)]
```

```
[ ]: print(sum(p.numel() for p in model.parameters() if p.requires_grad))
```

```
5
```

```
[ ]: preds = model(inputs)
     preds
```

```
[ ]: tensor([[27.6588],
           [12.9539],
           [22.0699],
           [40.6945],
           [49.7463],
           [52.1024],
           [31.6345],
           [40.2580]], grad_fn=<AddmmBackward0>)
```

```
[ ]: criterion_mse = nn.MSELoss()
     criterion_softmax_cross_entropy_loss = nn.CrossEntropyLoss()
```

```
[ ]: mse = criterion_mse(preds, targets)
     print(mse)
     print(mse.item())
```

```
tensor(1147.2864, grad_fn=<MseLossBackward0>)
1147.286376953125
```

```
d:\DataScience\Anaconda3\envs\dl4cv\lib\site-
packages\torch\nn\modules\loss.py:538: UserWarning: Using a target size
(torch.Size([8])) that is different to the input size (torch.Size([8, 1])). This
will likely lead to incorrect results due to broadcasting. Please ensure they
have the same size.
```

```
    return F.mse_loss(input, target, reduction=self.reduction)
```

```
[ ]: opt = torch.optim.Adam(model.parameters(), lr=0.001)
```

```
[ ]: import sys

     # Utility function to train the model
     def fit(num_epochs, model, loss_fn, opt, train_dl):

         # Repeat for given number of epochs
         for epoch in range(num_epochs):

             # Train with batches of data
```

```

for xb,yb in train_dl:

    xb.to(device) #move them to gpu if possible, if not, it will be cpu
    yb.to(device)

    # 1. Predict
    pred = model(xb)

    # 2. Calculate loss
    loss = loss_fn(pred, yb)

    # 3. Calculate gradient
    opt.zero_grad() #if not, the gradients will accumulate
    loss.backward()

    # Print out the gradients.
    #print ('dL/dw: ', model.weight.grad)
    #print ('dL/db: ', model.bias.grad)

    # 4. Update parameters using gradients
    opt.step()

    # Print the progress
    if (epoch+1) % 10 == 0:
        sys.stdout.write("\rEpoch [{}/{}], Loss: {:.4f}".format(epoch+1,
↪num_epochs, loss.item()))

```

```

[ ]: #train for 100 epochs
fit(100, model, criterion_mse, opt, train_dl)

```

Epoch [100/100], Loss: 0.0628

```

[ ]: preds = model(inputs)
loss = criterion_mse(preds, targets)
print(loss.item())

```

0.35663479566574097

```

[ ]: preds = model(inputs)
preds

```

```

[ ]: tensor([[3.3230],
           [3.1469],
           [2.6553],
           [2.6472],
           [3.1514],
           [3.3107],
           [3.1696],

```

```
[3.5440]], grad_fn=<AddmmBackward0>)
```

```
[ ]: from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()
```

```
model.fit(x_train, y_train)
```

```
preds2 = model.predict(x_train)
```

```
preds2
```

```
[ ]: array([3.696323 , 3.4681077, 2.4174857, 3.0752082, 3.63602 , 2.733765 ,  
          3.347672 , 3.5254207], dtype=float32)
```

```
[ ]: from sklearn.metrics import mean_squared_error, r2_score
```

```
print('MSE: ', mean_squared_error(preds.detach().numpy(), preds2))
```

```
print('R2_score: ', r2_score(preds.detach().numpy(), preds2))
```

```
MSE: 0.1352608
```

```
R2_score: -0.5435817138557375
```

We are getting Linear Regression model better because of number of training samples. Less training sample - classical ml is better. Deep learning benefits from a lot of data. Our training sample is just 8 examples, thats why Linear regression results are better

MSE: 0.135 R2_score is really bad, it is worser than mean ($r2_score(\text{mean}) = 0$)